

Anil Poonai

Github link for files and code: https://github.com/DevonARP/DeepLearning_A2

For question 1 I wrote out all my notes and work before writing them here, I'll add my work/notes at the end of each part for this question

Problem 1: Part A

$$L(w_1, w_2) = .5(aw_1^2 + bw_2^2)$$

$$\nabla L(w_1, w_2) = aw_1 + bw_2$$

$$\nabla^2 L(w_1, w_2) = a + b$$

I'm making $a = 1$ and $b = 2$ to keep the double derivate positive and keep the graph concave up so it has a minimum value.

$$\nabla^2 L(w_1, w_2) = a + b = 1 + 2 = 3$$

Now, we go back to the first derivative

$$\nabla L(w_1, w_2) = aw_1 + bw_2 = w_1 + 2w_2$$

I can use $w_1 = -2$ and $w_2 = 1$, this gives the minimum value of 0.

References:

<https://study.com/learn/lesson/how-to-find-the-maximum-value-of-a-function.html#:~:text=We%20will%20set%20the%20first,will%20be%20a%20minimum%20value>

https://www.ocf.berkeley.edu/~reinholz/ed/07fa_m155/lectures/second_derivative.pdf

a. $L(w_1, w_2) = .5(aw_1^2 + bw_2^2)$
 $\nabla L(w_1, w_2) = aw_1 + bw_2$
 $\nabla \nabla L(w_1, w_2) = a + b$
 $\quad \quad \quad = a + b$
making $a=1$ & $b=2$
to keep it concave
up & have a
min value
 $= 1 + 2$
 $\rightarrow \quad \quad \quad = w_1 + 2w_2$
 $\quad \quad \quad w_1 = -2 \quad w_2 = 1$

Problem 1: Part B

P_i is a dropout layer, I'll be using N for the dropout rate

Have to grab the derivative of L with respect to the corresponding weight

$W_i(t+1) = w_i(t) - N(\partial L / \partial w_i)$ This is being rearranged to follow the dropout rate argument formula

$$W_1 = w_1(t) - N a w_1(t) = w_1(t)(1 - N a) = p_1 w_1(t)$$

$$W_2 = w_2(t) - N b w_2(t) = w_2(t)(1 - N b) = p_2 w_2(t)$$

$$P_1 = 1 - N a$$

$$P_2 = 1 - N b$$

References:

<https://towardsdatascience.com/simplified-math-behind-dropout-in-deep-learning-6d50f3f47275#:~:text=In%20Keras%2C%20the%20dropout%20rate,can%20adversely%20affect%20the%20training.>

6. p is a dropout layer
using N for dropout rate

$$w_i(t+1) = w_i(t) - N \frac{\partial L}{\partial w_i}$$

\hookrightarrow following dropout rate argument formula $\hookrightarrow N a w_1$
 $N b w_2$

$$w_1 = w_1(t) - N a w_1(t)$$
$$w_1(t)(1 - N a) = p_1 w_1(t)$$
$$w_2 = w_2(t) - N b w_2(t)$$
$$w_2(t)(1 - N b) = p_2 w_2(t)$$
$$p_1 = 1 - N a$$
$$p_2 = 1 - N b$$

Problem 1: Part C

It converges when the gradient of the cost function becomes 0.

In this case both p_1 and p_2 need to be 0

We can rearrange that to be $1 - Na$ and $1 - Nb$ equal 0

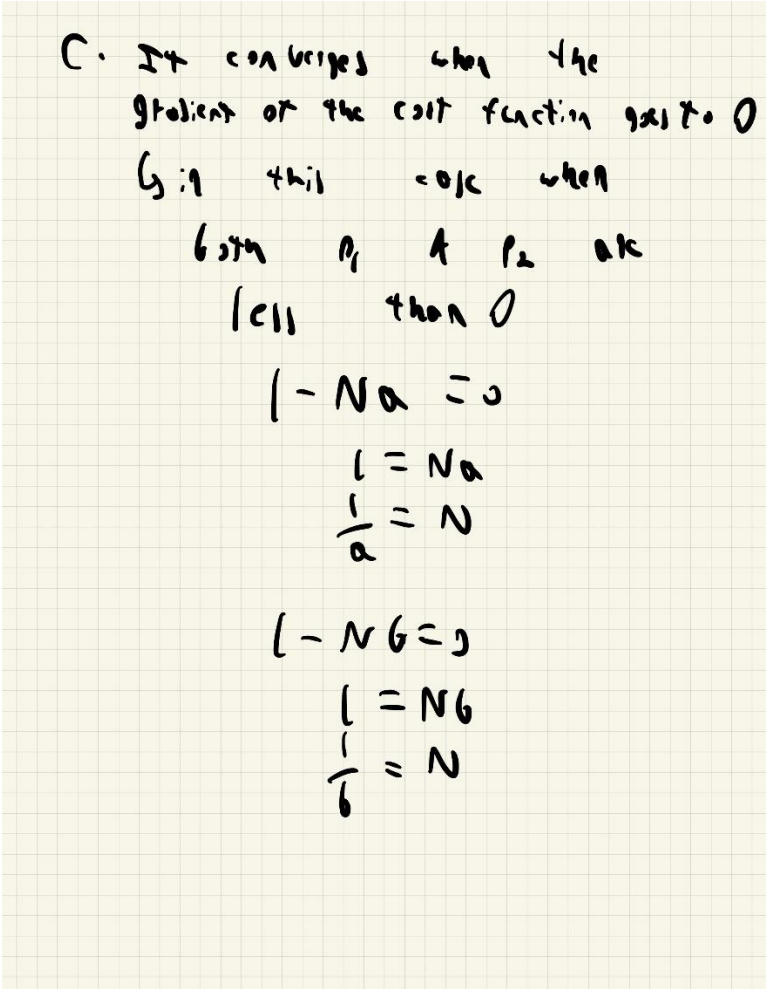
Then we can make it so that Na and Nb both equal 1

And we can end off with N equals $1/a$ and $1/b$.

References:

<https://www.cs.umd.edu/~djacobs/CMSC426/GradientDescent.pdf>

<https://www.cs.ubc.ca/~schmidtm/Courses/540-W18/L4.pdf>



C. It converges when the
gradient of the cost function goes to 0
by this case when
both p_1 & p_2 are
less than 0

$$1 - Na = 0$$
$$1 = Na$$
$$\frac{1}{a} = N$$

$$1 - Nb = 0$$
$$1 = Nb$$
$$\frac{1}{b} = N$$

Problem 1: Part D

When either a/b or b/a is very large, in the first case, a/b , then a would be the larger number and N for w_1 would be comparatively smaller than w_2 , and therefore w_1 would have to have more updates than w_2 . The opposite is true for the second case, b would be the larger number and N for w_2 would be comparatively smaller than w_1 , leaving w_2 to have more updates. The cases of a/b and b/a being very small is just the same case as the b/a and a/b being very large respectively.

Problem 2: Part A

I would use the Sobel filter, which uses two kernels, one for the horizontal edges and one for the vertical edges.

Horizontal edges

-1	0	1
-2	0	2
1	0	1

Vertical edges

-1	-2	-1
0	0	0
1	2	1

This works by looking for string changes in the image. The higher the sum of the numbers after the convolution, the more likely there is an edge there and the positive or negative sign indicates the direction of the edge. The output from both filters are then combined to see all the edges detected.

References:

https://www.projectrhea.org/rhea/index.php/An_Implementation_of_Sobel_Edge_Detection

<https://automaticaddison.com/how-the-sobel-operator-works/>

https://www.cs.auckland.ac.nz/compsci373s1c/PatricesLectures/Edge%20detection-Sobel_2up.pdf

Problem 2: Part B

I'm going to use a Box Blur Kernel because I end up mentioning the Gaussian Blur Kernel in Part D.

1	1	1
1	1	1
1	1	1

This works by giving each pixel the same weight and adding them all up then dividing by 9 in this case at the end, this makes it so that the output is a relative value to the other output points after the kernel as it adds up all of the points in a region and averages it out for every region.

References:

<https://medium.com/hackernoon/cv-for-busy-developers-convolutions-5c984f216e8c#:~:text=The%20convolution%20of%20a%20Gaussian,the%20kernel%20values%20is%2016.>

Problem 2: Part C

I'm combining some concepts from regular sharpening and edge detection for this.

0	-1	0
0	2	0
0	-1	0

We're focusing on the horizontal sharpening by grabbing the middle row of the region, this focuses on the center primarily and the middle row right after, with no focus on anything else as those values on the filter are zero.

References:

<https://medium.com/@boelsmaxence/introduction-to-image-processing-filters-179607f9824a>

[https://en.wikipedia.org/wiki/Kernel_\(image_processing\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))

Problem 2: Part D

I'm going to use a Gaussian Blur Kernel, this can also be used to blur an image.

1	2	1
2	4	2
1	2	1

This works by giving the pixel near the center of the kernel more weight than the ones on the edges, this helps mute noise as the box blur would treat all the points with the same weight.

References:

<https://medium.com/hackernoon/cv-for-busy-developers-convolutions-5c984f216e8c#:~:text=The%20convolution%20of%20a%20Gaussian,the%20kernel%20values%20is%2016.>

Problem 3: Part A

Jaccard similarity is supposed to find how similar 2 sets of data are, so it has to range between 0, no correlation at all, and 1, they are the exact same set. It can't be out of that range since the extremes would be either 0 and 1 for no relation and 100% relation, any number between 0 and 1 would represent a partial match with the higher the value being the higher match percentage. This can also be described as the IOU being the division of the overlap of the predicted and ground truth, so it can only be between the ranges 0 and 1.

$IOU = |A \cap B| / |A \cup B|$ #A is the predicted and B is the ground truth

B counts for all true positives, false negatives, and false positives while A counts for true positives

Problem 3: Part B

IOU isn't differentiable inherently as it ranges from 0 to 1, which means it can't differentiate between distances from how similar sets are as you would need all real numbers available. It also isn't differentiable in the case mentioned because the parameters have no influence on its gradient, so in the case of the top left to bottom right corners, the gradients are just going to be 0. Also, using the equation mentioned above indicates that the gradients would just be 0 everywhere as well, the parameters don't play a role in IOU calculation, it also is because the IOU metric itself isn't continuous natively.

Problem 4:

Code and answers will be below in an attached pdf to this document

Problem 5: Part A

Code and answers will be below in an attached pdf to this document

Problem 5: Part B

Code will be below in an attached pdf to this document

The backbone model is slightly better regarding loss after training as it has a lower loss and loss classifier but when looking at the example image it actually picks up 1 less person than the finetuning model but it is more confident in its predictions. This also leads to the IOU metric being an indicator, with the higher values being given towards the finetuning model, meaning it actually ends up matching more objects with that model. I'm not saying which is better as that's depends on what outcome is wanted but I do like the results from the finetuning a model a bit more, the IOU metric helps a lot.

Finetuning model

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1793 (0.1891) loss_classifier: 0.0253 (0.0269)
loss_box_reg: 0.0378 (0.0417) loss_mask: 0.1127 (0.1169) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0029 (0.0030) time: 0.5977 data: 0.0096 max mem: 3778

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.839

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.979

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.931

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.355

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.659

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.860

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.424

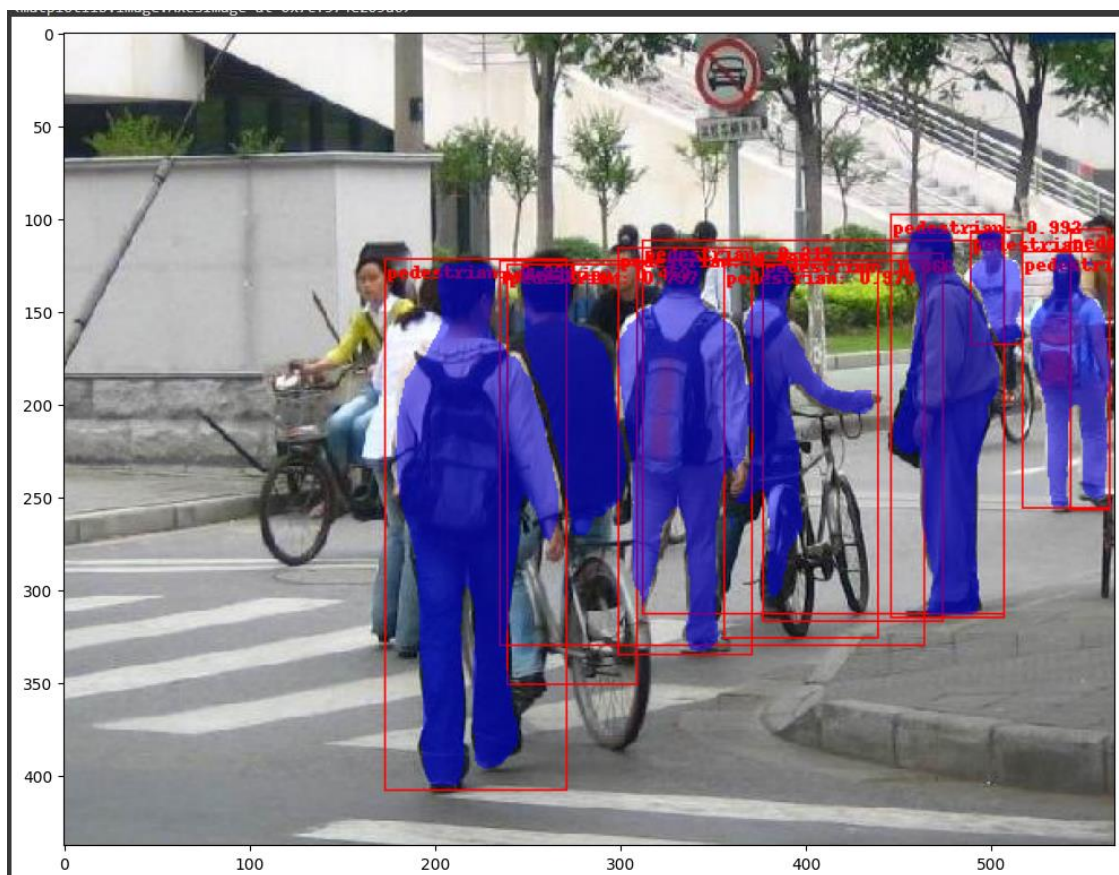
Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.877

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.877

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.467

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.825

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.894



```

pred_labels
[
  'pedestrian: 0.995',
  'pedestrian: 0.994',
  'pedestrian: 0.992',
  'pedestrian: 0.988',
  'pedestrian: 0.979',
  'pedestrian: 0.797',
  'pedestrian: 0.453',
  'pedestrian: 0.215',
  'pedestrian: 0.089',
  'pedestrian: 0.066',
  'pedestrian: 0.055'
]

[14] pred_boxes
tensor([
  [299, 116, 371, 335],
  [173, 122, 271, 408],
  [446, 98, 507, 315],
  [517, 118, 564, 256],
  [356, 125, 439, 326],
  [239, 125, 309, 351],
  [235, 123, 464, 330],
  [312, 112, 507, 313],
  [489, 107, 517, 168],
  [377, 119, 474, 317],
  [543, 106, 564, 257]], device='cuda:0')

```


Backbone Model

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1621 (0.1808) loss_classifier: 0.0219 (0.0247)
loss_box_reg: 0.0326 (0.0377) loss_mask: 0.1089 (0.1143) loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0033 (0.0029) time: 0.5792 data: 0.0088 max mem: 3780

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.806

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.975

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.919

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.348

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.700

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.830

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.354

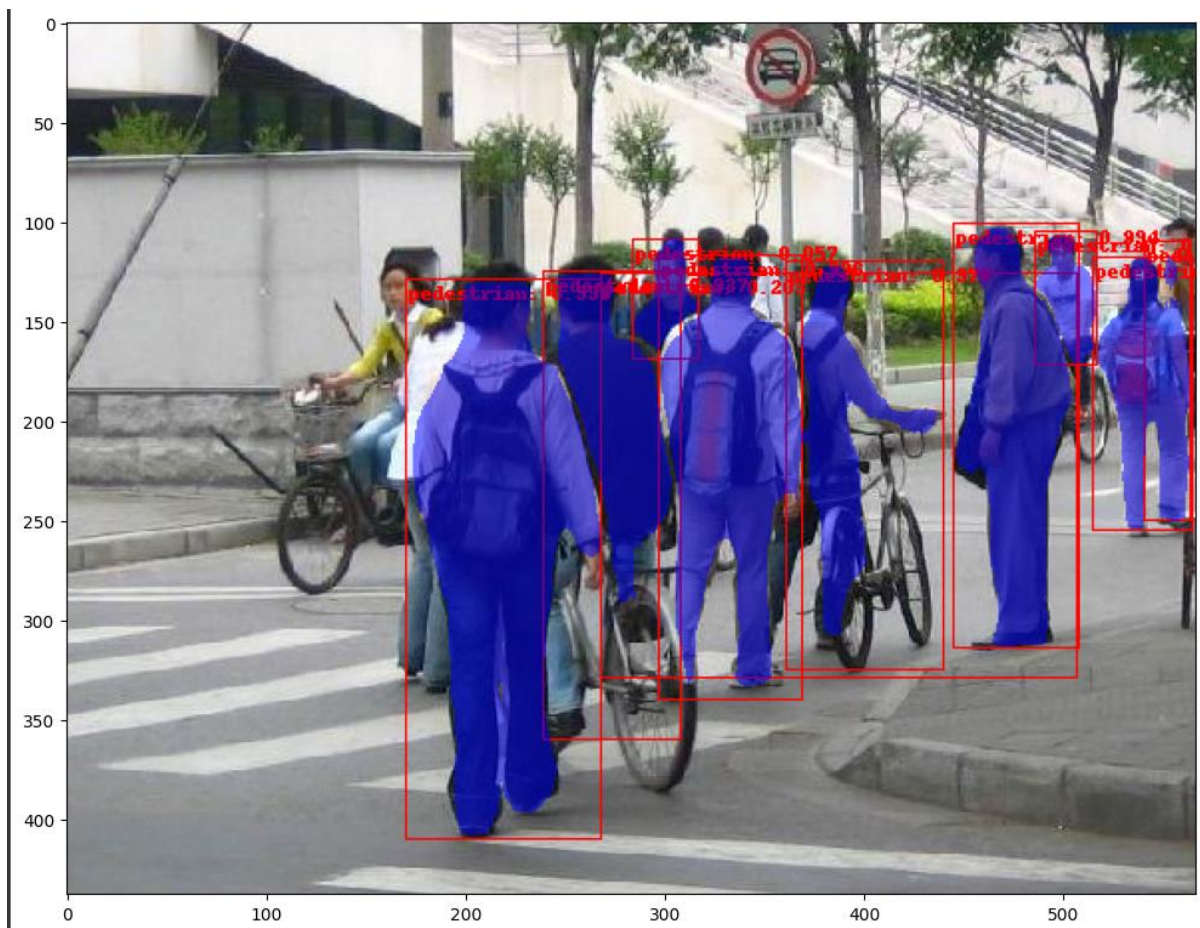
Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.853

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.853

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.467

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.775

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.872



pred_labels

```
[ 'pedestrian: 0.996',
  'pedestrian: 0.996',
  'pedestrian: 0.994',
  'pedestrian: 0.987',
  'pedestrian: 0.976',
  'pedestrian: 0.937',
  'pedestrian: 0.203',
  'pedestrian: 0.181',
  'pedestrian: 0.119',
  'pedestrian: 0.057']
```

[14] pred_boxes

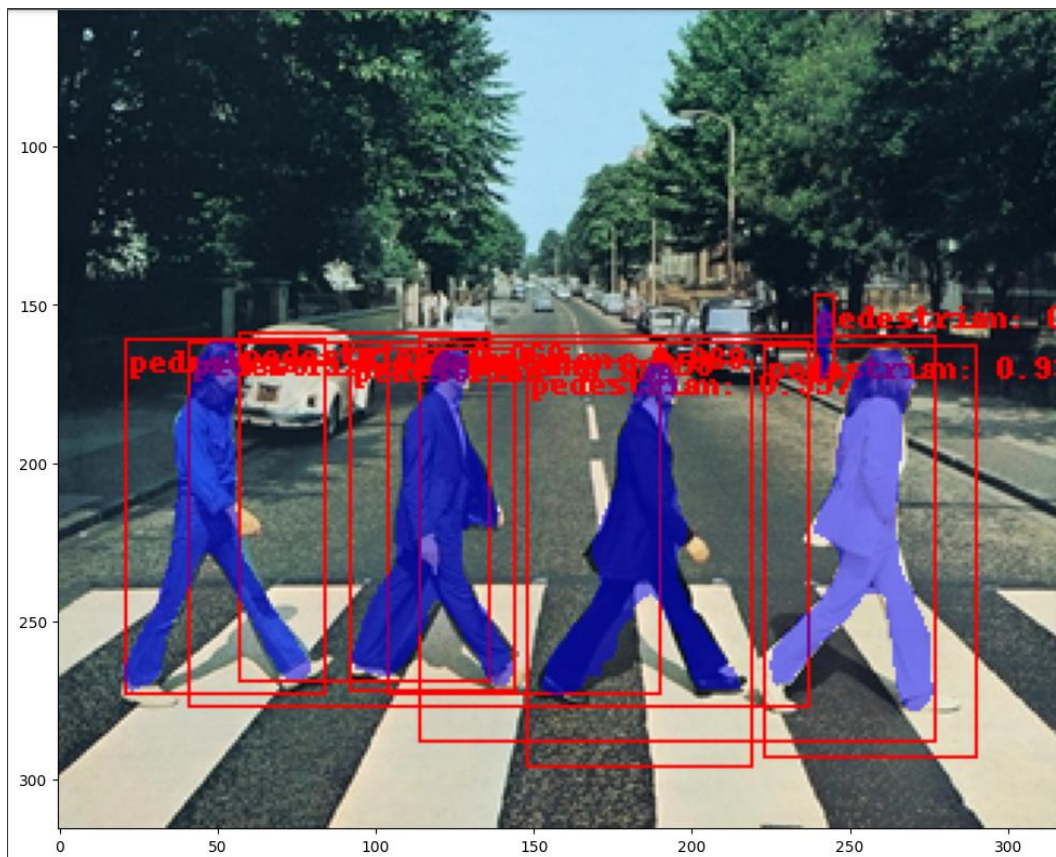
```
tensor([[170, 129, 268, 410],
        [297, 117, 369, 340],
        [445, 101, 508, 314],
        [515, 118, 565, 255],
        [361, 120, 440, 325],
        [239, 125, 308, 360],
        [268, 126, 507, 329],
        [486, 105, 517, 172],
        [541, 110, 564, 250],
        [284, 109, 317, 169]], device='cuda:0')
```

Problem 5: Part C

Code will be below in an attached pdf to this document

There are 5 people in the photograph, 4 are in the center and easy to spot but there's a fifth on the right side a bit in the back, the finetuning model picked up on that and with low confidence but the backbone model only recognized 4 people, being the 4 Beatles members in the middle of the photograph. The finetuning model also picked up some errors, it had bounding boxes around multiple Beatles members as possible objects but gave it a really low confidence score. So the finetuning model performed better here as it caught the 4 members with high confidence, a person in the back with fairly low confidence and had 4 other bounding boxes with multiple people in them but gave those boxes a really low confidence score, I would calculate those as False Positives. The backbone model performed worst not only because it missed a person entirely but because the confidence it had about the 4 Beatles members in the center of the image was low, it gave George Harrison a confidence score of less than .5 but the rest as above .9. All that being said is it important to point out that the finetuning model does have a higher error rate. Again, I'm not saying which is better but I still like the finetuning model a bit more, it might have more errors but it is more confident about detecting objects and can pick up on hard to see objects.

Finetuning Model




```
[10] pred_labels
```

```
['pedestrian: 0.987',  
 'pedestrian: 0.985',  
 'pedestrian: 0.957',  
 'pedestrian: 0.939',  
 'pedestrian: 0.103',  
 'pedestrian: 0.093',  
 'pedestrian: 0.089',  
 'pedestrian: 0.060',  
 'pedestrian: 0.059']
```

```
[11] pred_boxes
```

```
⇒ tensor([[ 21, 161,  84, 273],  
          [ 92, 164, 144, 272],  
          [148, 168, 219, 296],  
          [223, 163, 290, 293],  
          [ 41, 162, 237, 277],  
          [239, 147, 245, 174],  
          [114, 160, 277, 288],  
          [ 57, 159, 136, 269],  
          [104, 162, 190, 273]], device='cuda:0')
```

Backbone Model



pred_labels

['pedestrian: 0.977',
'pedestrian: 0.975',
'pedestrian: 0.933',
'pedestrian: 0.496']

[11] pred_boxes

```
tensor([[ 92, 161, 145, 269],  
        [ 21, 157,  81, 270],  
        [147, 168, 215, 286],  
        [224, 160, 293, 287]], device='cuda:0')
```


Question 4

AlexNet

In this problem, you are asked to train a deep convolutional neural network to perform image classification. In fact, this is a slight variation of a network called *AlexNet*. This is a landmark model in deep learning, and arguably kickstarted the current (and ongoing, and massive) wave of innovation in modern AI when its results were first presented in 2012. AlexNet was the first real-world demonstration of a *deep* classifier that was trained end-to-end on data and that outperformed all other ML models thus far.

We will train AlexNet using the [CIFAR10](#) dataset, which consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. The classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.

A lot of the code you will need is already provided in this notebook; all you need to do is to fill in the missing pieces, and interpret your results.

Warning : AlexNet takes a good amount of time to train (~1 minute per epoch on Google Colab). So please budget enough time to do this homework.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import _LRScheduler
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
```

```

from sklearn import manifold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time

SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

```

Loading and Preparing the Data

Our dataset is made up of color images but three color channels (red, green and blue), compared to MNIST's black and white images with a single color channel. To normalize our data we need to calculate the means and standard deviations for each of the color channels independently, and normalize them.

```

ROOT = '.data'
train_data = datasets.CIFAR10(root = ROOT,
                              train = True,
                              download = True)

```

Files already downloaded and verified

Compute means and standard deviations along the R,G,B channel

```

means = train_data.data.mean(axis = (0,1,2)) / 255
stds = train_data.data.std(axis = (0,1,2)) / 255

```

Next, we will do data augmentation. For each training image we will randomly rotate it (by up to 5 degrees), flip/mirror with probability 0.5, shift by +/-1 pixel. Finally we will normalize each color channel using the means/stds we calculated above.

```

train_transforms = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(32, padding = 2),
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                        std = stds)
])

```

```

    ])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                          std = stds)
])

```

Next, we'll load the dataset along with the transforms defined above.

We will also create a validation set with 10% of the training samples. The validation set will be used to monitor loss along different epochs, and we will pick the model along the optimization path that performed the best, and report final test accuracy numbers using this model.

```

train_data = datasets.CIFAR10(ROOT,
                               train = True,
                               download = True,
                               transform = train_transforms)

test_data = datasets.CIFAR10(ROOT,
                              train = False,
                              download = True,
                              transform = test_transforms)

Files already downloaded and verified
Files already downloaded and verified

VALID_RATIO = 0.9

n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples

train_data, valid_data = data.random_split(train_data,
                                           [n_train_examples,
                                           n_valid_examples])

valid_data = copy.deepcopy(valid_data)
valid_data.dataset.transform = test_transforms

```

Now, we'll create a function to plot some of the images in our dataset to see what they actually look like.

Note that by default PyTorch handles images that are arranged [channel, height, width], but matplotlib expects images to be [height, width, channel], hence we need to permute the dimensions of our images before plotting them.

```

def plot_images(images, labels, classes, normalize = False):
    n_images = len(images)

```



```

rows = int(np.sqrt(n_images))
cols = int(np.sqrt(n_images))

fig = plt.figure(figsize = (10, 10))

for i in range(rows*cols):

    ax = fig.add_subplot(rows, cols, i+1)

    image = images[i]

    if normalize:
        image_min = image.min()
        image_max = image.max()
        image.clamp_(min = image_min, max = image_max)
        image.add_(-image_min).div_(image_max - image_min + 1e-5)

    ax.imshow(image.permute(1, 2, 0).cpu().numpy())
    ax.set_title(classes[labels[i]])
    ax.axis('off')

```

One point here: `matplotlib` is expecting the values of every pixel to be between $[0, 1]$, however our normalization will cause them to be outside this range. By default `matplotlib` will then clip these values into the $[0, 1]$ range. This clipping causes all of the images to look a bit weird - all of the colors are oversaturated. The solution is to normalize each image between $[0, 1]$.

```

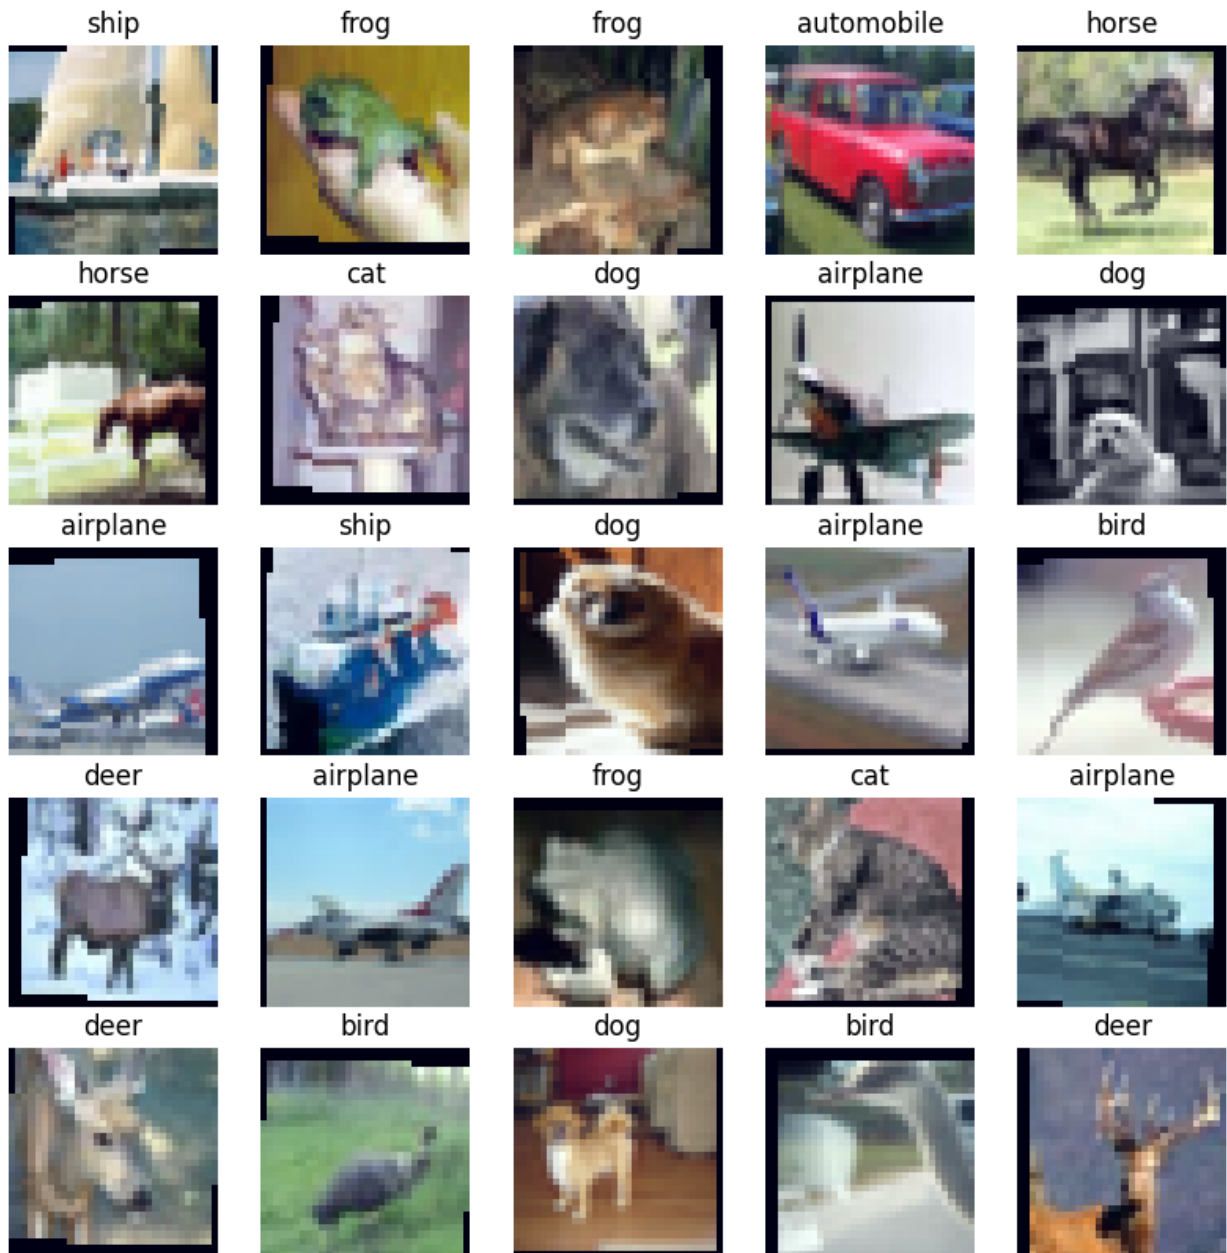
N_IMAGES = 25

images, labels = zip(*[(image, label) for image, label in
                        [train_data[i] for i in range(N_IMAGES)]])

classes = test_data.classes

plot_images(images, labels, classes, normalize = True)

```



We'll be normalizing our images by default from now on, so we'll write a function that does it for us which we can use whenever we need to renormalize an image.

```
def normalize_image(image):
    image_min = image.min()
    image_max = image.max()
    image.clamp_(min = image_min, max = image_max)
    image.add_(-image_min).div_(image_max - image_min + 1e-5)
    return image
```

The final bit of the data processing is creating the iterators. We will use a large. Generally, a larger batch size means that our model trains faster but is a bit more susceptible to overfitting.

```
# Q1: Create data loaders for train_data, valid_data, test_data
# Use batch size 256

BATCH_SIZE = 256

train_iterator = torch.utils.data.DataLoader(train_data,
batch_size=BATCH_SIZE, shuffle=True)#Added the iterable wraps for the
datasets

valid_iterator = torch.utils.data.DataLoader(valid_data,
batch_size=BATCH_SIZE, shuffle=True)

test_iterator = torch.utils.data.DataLoader(test_data,
batch_size=BATCH_SIZE, shuffle=True)
```

Defining the Model

Next up is defining the model.

AlexNet will have the following architecture:

- There are 5 2D convolutional layers (which serve as *feature extractors*), followed by 3 linear layers (which serve as the *classifier*).
- All layers (except the last one) have ReLU activations. (Use `inplace=True` while defining your ReLUs.)
- All convolutional filter sizes have kernel size 3 x 3 and padding 1.
- Convolutional layer 1 has stride 2. All others have the default stride (1).
- Convolutional layers 1,2, and 5 are followed by a 2D maxpool of size 2.
- Linear layers 1 and 2 are preceded by Dropouts with Bernoulli parameter 0.5.
- For the convolutional layers, the number of channels is set as follows. We start with 3 channels and then proceed like this:
 - 3 → 64 → 192 → 384 → 256 → 256

In the end, if everything is correct you should get a feature map of size $2 \times 2 \times 256 = 1024$.

- For the linear layers, the feature sizes are as follows:
 - 1024 → 4096 → 4096 → 10.

(The 10, of course, is because 10 is the number of classes in CIFAR-10).

```
class AlexNet(nn.Module):
    def __init__(self, output_dim):
```

```

super().__init__()

self.features = nn.Sequential(
    # Define according to the steps described above
    nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
#Added layers for AlexNet
    nn.ReLU(inplace=True),
    nn.MaxPool2d(kernel_size=2),
    nn.Conv2d(64, 192, kernel_size=3, stride=1, padding=1),
    nn.ReLU(inplace=True),
    nn.MaxPool2d(kernel_size=2),
    nn.Conv2d(192, 384, kernel_size=3, stride=1, padding=1),
    nn.ReLU(inplace=True),
    nn.Conv2d(384, 256, kernel_size=3, stride=1, padding=1),
    nn.ReLU(inplace=True),
    nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),
    nn.ReLU(inplace=True),
    nn.MaxPool2d(kernel_size=2),
)

self.classifier = nn.Sequential(
    # define according to the steps described above
    nn.Dropout(p=0.5),
    nn.Linear(1024, 4096),
    nn.ReLU(inplace=True),
    nn.Dropout(p=0.5),
    nn.Linear(4096, 4096),
    nn.ReLU(inplace=True),
    nn.Linear(4096, 10),
)

def forward(self, x):
    x = self.features(x)
    h = x.view(x.shape[0], -1)
    x = self.classifier(h)
    return x, h

```

We'll create an instance of our model with the desired amount of classes.

```

OUTPUT_DIM = 10
model = AlexNet(OUTPUT_DIM)

```

Training the Model

We first initialize parameters in PyTorch by creating a function that takes in a PyTorch module, checking what type of module it is, and then using the `nn.init` methods to actually initialize the parameters.

For convolutional layers we will initialize using the *Kaiming Normal* scheme, also known as *He Normal*. For the linear layers we initialize using the *Xavier Normal* scheme, also known as *Glorot Normal*. For both types of layer we initialize the bias terms to zeros.

```
def initialize_parameters(m):
    if isinstance(m, nn.Conv2d):
        nn.init.kaiming_normal_(m.weight.data, nonlinearity = 'relu')
        nn.init.constant_(m.bias.data, 0)
    elif isinstance(m, nn.Linear):
        nn.init.xavier_normal_(m.weight.data, gain =
nn.init.calculate_gain('relu'))
        nn.init.constant_(m.bias.data, 0)
```

We apply the initialization by using the model's `apply` method. If your definitions above are correct you should get the printed output as below.

```
model.apply(initialize_parameters)

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=1024, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
```

```

        (6): Linear(in_features=4096, out_features=10, bias=True)
    )
)

```

We then define the loss function we want to use, the device we'll use and place our model and criterion on to our device.

```

optimizer = optim.Adam(model.parameters(), lr = 1e-3)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
criterion = nn.CrossEntropyLoss()

model = model.to(device)
criterion = criterion.to(device)

# This is formatted as code

```

We define a function to calculate accuracy...

```

def calculate_accuracy(y_pred, y):
    top_pred = y_pred.argmax(1, keepdim = True)
    correct = top_pred.eq(y.view_as(top_pred)).sum()
    acc = correct.float() / y.shape[0]
    return acc

```

As we are using dropout we need to make sure to "turn it on" when training by using `model.train()`.

```

def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in iterator:

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred, _ = model(x)

        loss = criterion(y_pred, y)

        acc = calculate_accuracy(y_pred, y)

        loss.backward()

```

```

optimizer.step()

epoch_loss += loss.item()
epoch_acc += acc.item()

return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

We also define an evaluation loop, making sure to "turn off" dropout with `model.eval()`.

```

def evaluate(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for (x, y) in iterator:

            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

Next, we define a function to tell us how long an epoch takes.

```

def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs

```

Then, finally, we train our model.

Train it for 25 epochs (using the train dataset). At the end of each epoch, compute the validation loss and keep track of the best model. You might find the command `torch.save` helpful.

At the end you should expect to see validation losses of ~76% accuracy.

```
# Q3: train your model here for 25 epochs.  
# Print out training and validation loss/accuracy of the model after  
# each epoch  
# Keep track of the model that achieved best validation loss thus far.  
import time
```

```
EPOCHS = 25  
model_copy = copy.deepcopy(model)  
accuracy_ref = -1 #Lowest the accuracy can go is 0, so this lets the  
accuracy reference and new model be updated on the first epoch  
# Fill training code here  
for epoch in range(EPOCHS):  
    start_time = time.time()  
    loss, accuracy = train(model, train_iterator, optimizer, criterion,  
device) #Training data  
    print(f'Epoch {epoch}, Train Loss {loss}, Train Accuracy  
{accuracy}')    valid_loss, valid_accuracy = evaluate(model, valid_iterator,  
criterion, device) #Validation data  
    print(f'Epoch {epoch}, Valid Loss {valid_loss}, Valid Accuracy  
{valid_accuracy}')    if valid_accuracy > accuracy_ref: #Copies the new model if the  
validation accuracy is better than the previous copy  
        accuracy_ref = valid_accuracy  
        model_copy = copy.deepcopy(model)  
        print(f'Model Copied')  
    end_time = time.time()  
    min, sec = epoch_time(start_time, end_time)  
    print(f'Epoch took {min} minutes and {sec} seconds')
```

```
Epoch 0, Train Loss 2.3859293684363365, Train Accuracy  
0.21515447443181818  
Epoch 0, Valid Loss 1.6082023859024048, Valid Accuracy  
0.40342371314764025  
Model Copied  
Epoch took 2 minutes and 32 seconds  
Epoch 1, Train Loss 1.5127338414842433, Train Accuracy  
0.43745649859986524  
Epoch 1, Valid Loss 1.3785551190376282, Valid Accuracy  
0.48405330926179885  
Model Copied  
Epoch took 2 minutes and 30 seconds  
Epoch 2, Train Loss 1.3497779613191432, Train Accuracy  
0.5120134942910888  
Epoch 2, Valid Loss 1.2114853024482728, Valid Accuracy  
0.5644990801811218  
Model Copied  
Epoch took 2 minutes and 20 seconds  
Epoch 3, Train Loss 1.252753977071155, Train Accuracy  
0.5508149858902801
```


Epoch 3, Valid Loss 1.1492135107517243, Valid Accuracy 0.5816061586141587
Model Copied
Epoch took 2 minutes and 21 seconds
Epoch 4, Train Loss 1.1764032918621192, Train Accuracy 0.5811496803706343
Epoch 4, Valid Loss 1.1288484692573548, Valid Accuracy 0.6064223349094391
Model Copied
Epoch took 2 minutes and 27 seconds
Epoch 5, Train Loss 1.1137279759753833, Train Accuracy 0.607915483076464
Epoch 5, Valid Loss 1.041766142845154, Valid Accuracy 0.6327435672283173
Model Copied
Epoch took 2 minutes and 25 seconds
Epoch 6, Train Loss 1.0533894764428788, Train Accuracy 0.6306143467399207
Epoch 6, Valid Loss 1.0123582810163498, Valid Accuracy 0.6456916362047196
Model Copied
Epoch took 2 minutes and 34 seconds
Epoch 7, Train Loss 1.016191769729961, Train Accuracy 0.6427503549917177
Epoch 7, Valid Loss 0.9771516680717468, Valid Accuracy 0.656031709909439
Model Copied
Epoch took 2 minutes and 51 seconds
Epoch 8, Train Loss 0.9675288193605163, Train Accuracy 0.6617116477679122
Epoch 8, Valid Loss 0.9208255469799042, Valid Accuracy 0.6822150737047196
Model Copied
Epoch took 3 minutes and 7 seconds
Epoch 9, Train Loss 0.9362695372917436, Train Accuracy 0.673189808360555
Epoch 9, Valid Loss 0.9574713259935379, Valid Accuracy 0.6697265625
Epoch took 2 minutes and 47 seconds
Epoch 10, Train Loss 0.915364300662821, Train Accuracy 0.67897372150963
Epoch 10, Valid Loss 0.8721305072307587, Valid Accuracy 0.6986787676811218
Model Copied
Epoch took 2 minutes and 43 seconds
Epoch 11, Train Loss 0.8793735033409162, Train Accuracy 0.6945791904899207
Epoch 11, Valid Loss 0.8667667210102081, Valid Accuracy 0.6991842836141586
Model Copied

Epoch took 2 minutes and 40 seconds
Epoch 12, Train Loss 0.8457419658926401, Train Accuracy 0.7068705609576269
Epoch 12, Valid Loss 0.8426105201244354, Valid Accuracy 0.7082375913858414
Model Copied
Epoch took 2 minutes and 43 seconds
Epoch 13, Train Loss 0.8229072757742621, Train Accuracy 0.7161807529628277
Epoch 13, Valid Loss 0.8680290371179581, Valid Accuracy 0.7025850176811218
Epoch took 2 minutes and 53 seconds
Epoch 14, Train Loss 0.8056459250775251, Train Accuracy 0.7222514203326269
Epoch 14, Valid Loss 0.7988647848367691, Valid Accuracy 0.7249540448188782
Model Copied
Epoch took 2 minutes and 49 seconds
Epoch 15, Train Loss 0.7857820635492151, Train Accuracy 0.7266770242290064
Epoch 15, Valid Loss 0.7764606267213822, Valid Accuracy 0.7301700353622437
Model Copied
Epoch took 2 minutes and 44 seconds
Epoch 16, Train Loss 0.7603419257158582, Train Accuracy 0.73777432536537
Epoch 16, Valid Loss 0.7669012516736984, Valid Accuracy 0.7389131426811218
Model Copied
Epoch took 2 minutes and 50 seconds
Epoch 17, Train Loss 0.7526960955424742, Train Accuracy 0.74021573161537
Epoch 17, Valid Loss 0.7977280527353287, Valid Accuracy 0.7365923702716828
Epoch took 2 minutes and 35 seconds
Epoch 18, Train Loss 0.7299478788944808, Train Accuracy 0.74827237224037
Epoch 18, Valid Loss 0.7806057006120681, Valid Accuracy 0.7405560672283172
Model Copied
Epoch took 2 minutes and 42 seconds
Epoch 19, Train Loss 0.7180966392836787, Train Accuracy 0.7520321377299048
Epoch 19, Valid Loss 0.7532464444637299, Valid Accuracy 0.7397977948188782
Epoch took 2 minutes and 41 seconds
Epoch 20, Train Loss 0.695985344323245, Train Accuracy 0.7610040838745508
Epoch 20, Valid Loss 0.7738403081893921, Valid Accuracy

```
0.7301470577716828
Epoch took 2 minutes and 45 seconds
Epoch 21, Train Loss 0.687265569852157, Train Accuracy
0.7641637074676427
Epoch 21, Valid Loss 0.7499783217906952, Valid Accuracy
0.7522977948188782
Model Copied
Epoch took 2 minutes and 44 seconds
Epoch 22, Train Loss 0.6752578406171366, Train Accuracy
0.7682998935607347
Epoch 22, Valid Loss 0.7266848772764206, Valid Accuracy
0.7558938413858414
Model Copied
Epoch took 2 minutes and 48 seconds
Epoch 23, Train Loss 0.6537438587031581, Train Accuracy
0.7748322087255392
Epoch 23, Valid Loss 0.7215611875057221, Valid Accuracy
0.7567210465669632
Model Copied
Epoch took 2 minutes and 48 seconds
Epoch 24, Train Loss 0.6516660617833788, Train Accuracy
0.7765660512853753
Epoch 24, Valid Loss 0.6996587306261063, Valid Accuracy
0.7605583637952804
Model Copied
Epoch took 2 minutes and 49 seconds
```

Evaluating the model

We then load the parameters of our model that achieved the best validation loss. You should expect to see ~75% accuracy of this model on the test dataset.

Finally, plot the confusion matrix of this model and comment on any interesting patterns you can observe there. For example, which two classes are confused the most?

```
# Q4: Load the best performing model, evaluate it on the test dataset,
and print test accuracy.
```

```
# Also, print out the confusion matrox.
```

```
def get_predictions(model, iterator, device):
```

```
    model.eval()
    with torch.no_grad():
        labels = []
        probs = []
```

```
# Q4: Fill code here.
```

```

        for (x, y) in iterator:
            x, y = x.to(device), y.to(device)
            predicted_output = model(x) #Prediction on test data
            labels.append(y)
            probs.append(predicted_output[0]) #Ignore the second
tensor

        labels = torch.cat(labels, dim = 0) #Converts all of the
tensors into 1 big tensor
        probs = torch.cat(probs, dim = 0)

        return labels, probs

labels, probs = get_predictions(model_copy, test_iterator, device)
#New model is used

pred_labels = torch.argmax(probs, 1) #Gets index or largest
possibility

print(torch.eq(labels, pred_labels)) #Matches the prediction and label
tensor values

tensor([False,  True,  True,  ...,  True,  True, False])

torch.eq(labels, pred_labels).sum() #Since it's boolean, I can just
sum it all up to get the amount of matches

tensor(7629)

torch.eq(labels, pred_labels).sum()/len(torch.eq(labels, pred_labels))
#Percentage that was predicted correctly at 76%

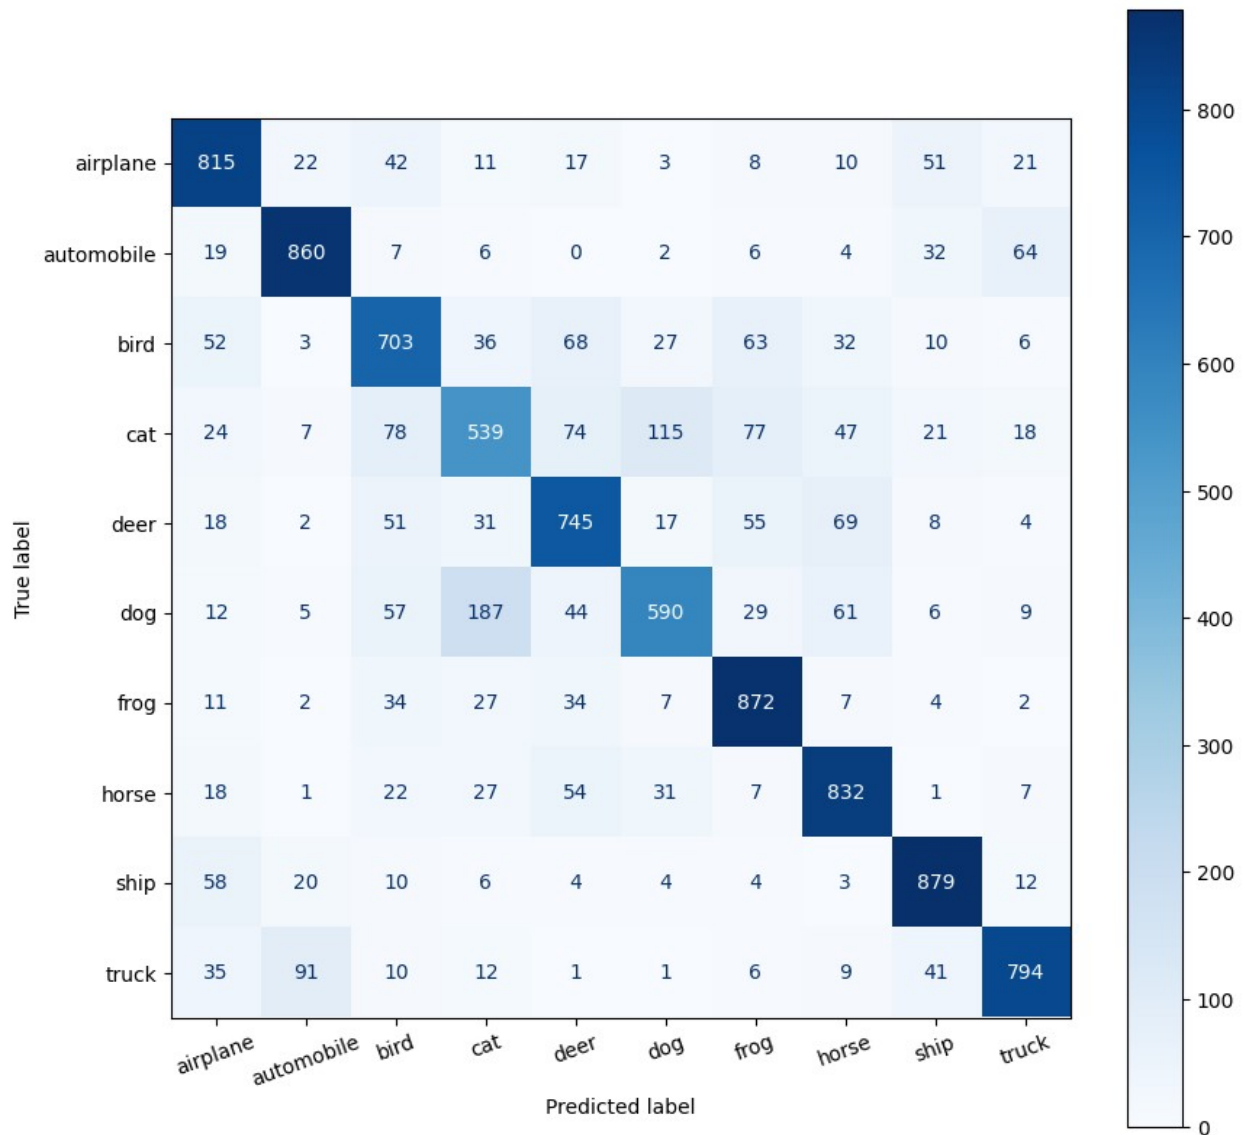
tensor(0.7629)

def plot_confusion_matrix(labels, pred_labels, classes):

    fig = plt.figure(figsize = (10, 10));
    ax = fig.add_subplot(1, 1, 1);
    cm = confusion_matrix(labels, pred_labels);
    cm = ConfusionMatrixDisplay(cm, display_labels = classes);
    cm.plot(values_format = 'd', cmap = 'Blues', ax = ax)
    plt.xticks(rotation = 20)

plot_confusion_matrix(labels, pred_labels, classes)

```



Conclusion

That's it! As a side project (this is not for credit and won't be graded), feel free to play around with different design choices that you made while building this network.

- Whether or not to normalize the color channels in the input.
- The learning rate parameter in Adam.
- The batch size.
- The number of training epochs.
- (and if you are feeling brave -- the AlexNet architecture itself.)

This is for Question 5 Part A

```
# For tips on running notebooks in Google Colab, see
# https://pytorch.org/tutorials/beginner/colab
%matplotlib inline

from google.colab import drive
drive.mount('/content/drive') #Updated the spots where they needed the
locations to change, also downloaded the PennFudan dataset and
tutorial source for the test photo at the end
# Dataset: https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
# Test Photo:
https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425fa
a00c06bcf502/_static/img/tv_tutorial/tv_image05.png
#Source Code:
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

Mounted at /content/drive
```

TorchVision Object Detection Finetuning Tutorial

.. tip::

To get the most of this tutorial, we suggest using this [Colab Version](https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_fineting_instance_segmentation.ipynb). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model on the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection and instance segmentation model on a custom dataset.

.. note ::

This tutorial works only with torchvision version `>=0.16` or `nightly`. If you're using `torchvision<=0.15`, please follow [this tutorial instead](https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425faa00c06bcf502/intermediate_source/torchvision_tutorial.rst).

Defining the Dataset

The reference scripts for training object detection, instance segmentation and person keypoint detection allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return a tuple:

- `image`: `:class:torchvision.tv_tensors.Image` of shape `[3, H, W]`, a pure tensor, or a PIL Image of size `(H, W)`
- `target`: a dict containing the following fields
 - `boxes`, `:class:torchvision.tv_tensors.BoundingBoxes` of shape `[N, 4]`: the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels`, integer `:class:torch.Tensor` of shape `[N]`: the label for each bounding box. `0` represents always the background class.
 - `image_id`, int: an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area`, float `:class:torch.Tensor` of shape `[N]`: the area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd`, uint8 `:class:torch.Tensor` of shape `[N]`: instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks`, `:class:torchvision.tv_tensors.Mask` of shape `[N, H, W]`: the segmentation masks for each one of the objects

If your dataset is compliant with above requirements then it will work for both training and evaluation codes from the reference script. Evaluation code will use scripts from `pycocotools` which can be installed with `pip install pycocotools`.

.. note :: For Windows, please install `pycocotools` from [gautamchitnis_](https://github.com/gautamchitnis/cocoapi) with command

```
pip install
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-
master#subdirectory=PythonAPI
```

One note on the `labels`. The model considers class `0` as background. If your dataset does not contain the background class, you should not have `0` in your `labels`. For example, assuming you have just two classes, *cat* and *dog*, you can define `1` (not `0`) to represent *cats* and `2` to represent *dogs*. So, for instance, if one of the images has both classes, your `labels` tensor should look like `[1, 2]`.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratios), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

Writing a custom dataset for PennFudan

Let's write a dataset for the PennFudan dataset. After [downloading and extracting the zip file](#), we have the following folder structure:

::

```
PennFudanPed/ PedMasks/ FudanPed00001_mask.png FudanPed00002_mask.png
FudanPed00003_mask.png FudanPed00004_mask.png ... PNGImages/ FudanPed00001.png
FudanPed00002.png FudanPed00003.png FudanPed00004.png
```

Here is one example of a pair of images and segmentation masks

So each image has a corresponding segmentation mask, where each color correspond to a different instance. Let's write a :class:torch.utils.data.Dataset class for this dataset. In the code below, we are wrapping images, bounding boxes and masks into torchvision.TVTensor classes so that we will be able to apply torchvision built-in transformations ([new Transforms API](#)) for the given object detection and segmentation task. Namely, image tensors will be wrapped by :class:torchvision.tv_tensors.Image, bounding boxes into :class:torchvision.tv_tensors.BoundingBoxes and masks into :class:torchvision.tv_tensors.Mask. As torchvision.TVTensor are :class:torch.Tensor subclasses, wrapped objects are also tensors and inherit the plain :class:torch.Tensor API. For more information about torchvision tv_tensors see [this documentation](#).

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
```



```

        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,)), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,)), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

That's all for the dataset. Now let's define a model that can perform predictions on this dataset.

Defining your model

In this tutorial, we will be using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.

Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

There are two common situations where one might want to modify one of the available models in TorchVision Model Zoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

1 - Finetuning from a pretrained model

Let's suppose that you want to start from a model pre-trained on COCO and want to finetune it for your particular classes. Here is a possible way of doing it:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

# load a model pre-trained on COCO
model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)

# replace the classifier with a new one, that has
# num_classes which is user-defined
num_classes = 2 # 1 class (person) + background
# get number of input features for the classifier
in_features = model.roi_heads.box_predictor.cls_score.in_features
# replace the pre-trained head with a new one
model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

Downloading:
"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth" to
/root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth
100%|██████████| 160M/160M [00:01<00:00, 141MB/s]
```

2 - Modifying the model to add a different backbone

```
import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator

# load a pre-trained model for classification and return
# only the features
backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features
# ``FasterRCNN`` needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(
    sizes=((32, 64, 128, 256, 512)),
    aspect_ratios=((0.5, 1.0, 2.0)),
)

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# ``OrderedDict[Tensor]``, and in ``featmap_names`` you can choose
# which
# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(
    featmap_names=['0'],
    output_size=7,
    sampling_ratio=2
)

# put the pieces together inside a Faster-RCNN model
model = FasterRCNN(
    backbone,
    num_classes=2,
    rpn_anchor_generator=anchor_generator,
    box_roi_pool=roi_pooler
)

Downloading: "https://download.pytorch.org/models/mobilenet_v2-
7ebf99e0.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-
7ebf99e0.pth
100%|██████████| 13.6M/13.6M [00:00<00:00, 32.3MB/s]
```

Object detection and instance segmentation model for PennFudan Dataset

In our case, we want to finetune from a pre-trained model, given that our dataset is very small, so we will be following approach number 1.

Here we want to also compute the instance segmentation masks, so we will be using Mask R-CNN:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256
    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(
        in_features_mask,
        hidden_layer,
        num_classes
    )

    return model
```

That's it, this will make `model` be ready to be trained and evaluated on your custom dataset.

Putting everything together

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py` and `references/detection/utils.py`. Just download everything under `references/detection` to your folder and use them here. On Linux if you have `wget`, you can download them using below commands:

```

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API
<https://pytorch.org/vision/stable/transforms.html>`_
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#
# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
        transforms.append(T.ToDtype(torch.float, scale=True))
        transforms.append(T.ToPureTensor())
    return T.Compose(transforms)

# Testing ``forward()`` method (Optional)
# -----
#
# Before iterating over the dataset, it's good to see what the model
# expects during training and inference time on sample data.
import utils

model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
data_loader = torch.utils.data.DataLoader(

```

```

        dataset,
        batch_size=2,
        shuffle=True,
        num_workers=4,
        collate_fn=utils.collate_fn
    )

# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
print(output)

# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions
print(predictions[0])

/usr/local/lib/python3.10/dist-packages/torch/utils/data/
dataloader.py:557: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

{'loss_classifier': tensor(0.1144, grad_fn=<NllLossBackward0>),
'loss_box_reg': tensor(0.0404, grad_fn=<DivBackward0>),
'loss_objectness': tensor(0.0071,
grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), 'loss_rpn_box_reg':
tensor(0.0054, grad_fn=<DivBackward0>)}
{'boxes': tensor([], size=(0, 4), grad_fn=<StackBackward0>), 'labels':
tensor([], dtype=torch.int64), 'scores': tensor([],
grad_fn=<IndexBackward0>)}

```

Let's now write the main function which performs the training and the validation:

```

from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',

```

```

get_transform(train=True))
dataset_test = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 5

```

```

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

Downloading:

```

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth" to
/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth
100%|██████████| 170M/170M [00:01<00:00, 140MB/s]

```

```

Epoch: [0] [ 0/60] eta: 0:08:32 lr: 0.000090 loss: 2.9040 (2.9040)
loss_classifier: 0.8063 (0.8063) loss_box_reg: 0.2864 (0.2864)
loss_mask: 1.7528 (1.7528) loss_objectness: 0.0557 (0.0557)
loss_rpn_box_reg: 0.0028 (0.0028) time: 8.5405 data: 0.7895 max
mem: 2148
Epoch: [0] [10/60] eta: 0:01:11 lr: 0.000936 loss: 1.4138 (1.8127)
loss_classifier: 0.5713 (0.5281) loss_box_reg: 0.2840 (0.2832)
loss_mask: 0.5955 (0.9743) loss_objectness: 0.0220 (0.0239)
loss_rpn_box_reg: 0.0028 (0.0032) time: 1.4238 data: 0.0855 max
mem: 3041
Epoch: [0] [20/60] eta: 0:00:40 lr: 0.001783 loss: 0.8679 (1.2905)
loss_classifier: 0.2435 (0.3770) loss_box_reg: 0.2475 (0.2576)
loss_mask: 0.3593 (0.6307) loss_objectness: 0.0164 (0.0204)
loss_rpn_box_reg: 0.0033 (0.0048) time: 0.6335 data: 0.0115 max
mem: 3041
Epoch: [0] [30/60] eta: 0:00:26 lr: 0.002629 loss: 0.6184 (1.0573)
loss_classifier: 0.1425 (0.2887) loss_box_reg: 0.2219 (0.2575)
loss_mask: 0.2009 (0.4885) loss_objectness: 0.0079 (0.0166)
loss_rpn_box_reg: 0.0073 (0.0060) time: 0.5860 data: 0.0089 max
mem: 3041
Epoch: [0] [40/60] eta: 0:00:16 lr: 0.003476 loss: 0.4260 (0.8925)
loss_classifier: 0.0628 (0.2297) loss_box_reg: 0.1934 (0.2366)
loss_mask: 0.1704 (0.4070) loss_objectness: 0.0042 (0.0136)
loss_rpn_box_reg: 0.0039 (0.0056) time: 0.5971 data: 0.0095 max
mem: 3041
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.3985 (0.7969)
loss_classifier: 0.0396 (0.1942) loss_box_reg: 0.1551 (0.2237)
loss_mask: 0.1441 (0.3614) loss_objectness: 0.0022 (0.0119)
loss_rpn_box_reg: 0.0034 (0.0057) time: 0.5937 data: 0.0086 max
mem: 3041
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2984 (0.7316)

```



```
loss_classifier: 0.0380 (0.1735) loss_box_reg: 0.1196 (0.2106)
loss_mask: 0.1490 (0.3309) loss_objectness: 0.0013 (0.0106)
loss_rpn_box_reg: 0.0044 (0.0059) time: 0.6078 data: 0.0084 max
mem: 3041
Epoch: [0] Total time: 0:00:44 (0.7489 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.2112 (0.2112)
evaluator_time: 0.0035 (0.0035) time: 0.5069 data: 0.2905 max mem:
3041
Test: [49/50] eta: 0:00:00 model_time: 0.1151 (0.1320)
evaluator_time: 0.0060 (0.0107) time: 0.1344 data: 0.0054 max mem:
3041
Test: Total time: 0:00:08 (0.1603 s / it)
Averaged stats: model_time: 0.1151 (0.1320) evaluator_time: 0.0060
(0.0107)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.712
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.887
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.418
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.445
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.724
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.321
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.800
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.774
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.737
Average Precision (AP) @[ IoU=0.50 | area= all |
```

```
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.912
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.362
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.340
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.319
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.774
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Epoch: [1] [ 0/60] eta: 0:00:56 lr: 0.005000 loss: 0.2784 (0.2784)
loss_classifier: 0.0530 (0.0530) loss_box_reg: 0.0800 (0.0800)
loss_mask: 0.1416 (0.1416) loss_objectness: 0.0007 (0.0007)
loss_rpn_box_reg: 0.0032 (0.0032) time: 0.9449 data: 0.3936 max
mem: 3041
Epoch: [1] [10/60] eta: 0:00:31 lr: 0.005000 loss: 0.2784 (0.2844)
loss_classifier: 0.0315 (0.0372) loss_box_reg: 0.0800 (0.0917)
loss_mask: 0.1416 (0.1491) loss_objectness: 0.0007 (0.0011)
loss_rpn_box_reg: 0.0037 (0.0053) time: 0.6239 data: 0.0419 max
mem: 3041
Epoch: [1] [20/60] eta: 0:00:24 lr: 0.005000 loss: 0.2673 (0.2773)
loss_classifier: 0.0337 (0.0367) loss_box_reg: 0.0789 (0.0854)
loss_mask: 0.1440 (0.1488) loss_objectness: 0.0007 (0.0015)
loss_rpn_box_reg: 0.0037 (0.0049) time: 0.5967 data: 0.0098 max
mem: 3041
Epoch: [1] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2713 (0.2893)
loss_classifier: 0.0357 (0.0387) loss_box_reg: 0.0829 (0.0895)
loss_mask: 0.1530 (0.1546) loss_objectness: 0.0008 (0.0015)
loss_rpn_box_reg: 0.0041 (0.0049) time: 0.6009 data: 0.0110 max
mem: 3041
Epoch: [1] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2660 (0.2844)
loss_classifier: 0.0383 (0.0390) loss_box_reg: 0.0762 (0.0884)
loss_mask: 0.1448 (0.1502) loss_objectness: 0.0011 (0.0019)
loss_rpn_box_reg: 0.0041 (0.0048) time: 0.5973 data: 0.0103 max
mem: 3041
Epoch: [1] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2605 (0.2799)
loss_classifier: 0.0405 (0.0392) loss_box_reg: 0.0762 (0.0858)
loss_mask: 0.1319 (0.1476) loss_objectness: 0.0021 (0.0020)
```

```
loss_rpn_box_reg: 0.0040 (0.0053)  time: 0.5930  data: 0.0102  max
mem: 3132
Epoch: [1]  [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.2641 (0.2789)
loss_classifier: 0.0406 (0.0397)  loss_box_reg: 0.0738 (0.0848)
loss_mask: 0.1278 (0.1472)  loss_objectness: 0.0005 (0.0019)
loss_rpn_box_reg: 0.0037 (0.0054)  time: 0.5962  data: 0.0086  max
mem: 3132
Epoch: [1] Total time: 0:00:36 (0.6078 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:25  model_time: 0.1575 (0.1575)
evaluator_time: 0.0035 (0.0035)  time: 0.5199  data: 0.3574  max mem:
3132
Test:  [49/50]  eta: 0:00:00  model_time: 0.1032 (0.1119)
evaluator_time: 0.0036 (0.0054)  time: 0.1195  data: 0.0038  max mem:
3132
Test: Total time: 0:00:06 (0.1344 s / it)
Averaged stats: model_time: 0.1032 (0.1119)  evaluator_time: 0.0036
(0.0054)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.816
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.987
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.960
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.440
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.489
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.829
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.362
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.852
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.852
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.867
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
IoU metric: segm
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.755
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.989
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.921
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.429
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.421
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.766
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.331
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.790
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.790
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.797
Epoch: [2] [ 0/60] eta: 0:01:04 lr: 0.005000 loss: 0.2961 (0.2961)
loss_classifier: 0.0745 (0.0745) loss_box_reg: 0.0721 (0.0721)
loss_mask: 0.1432 (0.1432) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0062 (0.0062) time: 1.0802 data: 0.4941 max
mem: 3132
Epoch: [2] [10/60] eta: 0:00:32 lr: 0.005000 loss: 0.2604 (0.2647)
loss_classifier: 0.0426 (0.0437) loss_box_reg: 0.0721 (0.0735)
loss_mask: 0.1367 (0.1407) loss_objectness: 0.0006 (0.0016)
loss_rpn_box_reg: 0.0046 (0.0051) time: 0.6582 data: 0.0519 max
mem: 3132
Epoch: [2] [20/60] eta: 0:00:25 lr: 0.005000 loss: 0.2255 (0.2370)
loss_classifier: 0.0294 (0.0360) loss_box_reg: 0.0615 (0.0648)
loss_mask: 0.1238 (0.1304) loss_objectness: 0.0006 (0.0013)
loss_rpn_box_reg: 0.0029 (0.0044) time: 0.6032 data: 0.0079 max
mem: 3132
Epoch: [2] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2249 (0.2365)
loss_classifier: 0.0272 (0.0351) loss_box_reg: 0.0517 (0.0633)
loss_mask: 0.1255 (0.1322) loss_objectness: 0.0008 (0.0014)
loss_rpn_box_reg: 0.0029 (0.0044) time: 0.5744 data: 0.0084 max
mem: 3132
Epoch: [2] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2193 (0.2306)
loss_classifier: 0.0292 (0.0344) loss_box_reg: 0.0504 (0.0618)
loss_mask: 0.1233 (0.1292) loss_objectness: 0.0006 (0.0012)
loss_rpn_box_reg: 0.0031 (0.0040) time: 0.5783 data: 0.0091 max
mem: 3132

```

```
Epoch: [2] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2193 (0.2335)
loss_classifier: 0.0298 (0.0348) loss_box_reg: 0.0609 (0.0630)
loss_mask: 0.1224 (0.1302) loss_objectness: 0.0005 (0.0013)
loss_rpn_box_reg: 0.0031 (0.0042) time: 0.5997 data: 0.0085 max
mem: 3132
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2122 (0.2300)
loss_classifier: 0.0267 (0.0341) loss_box_reg: 0.0508 (0.0611)
loss_mask: 0.1272 (0.1294) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0033 (0.0041) time: 0.5890 data: 0.0075 max
mem: 3132
Epoch: [2] Total time: 0:00:36 (0.6033 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1435 (0.1435)
evaluator_time: 0.0036 (0.0036) time: 0.5156 data: 0.3669 max mem:
3132
Test: [49/50] eta: 0:00:00 model_time: 0.1125 (0.1128)
evaluator_time: 0.0052 (0.0057) time: 0.1273 data: 0.0051 max mem:
3132
Test: Total time: 0:00:06 (0.1388 s / it)
Averaged stats: model_time: 0.1125 (0.1128) evaluator_time: 0.0052
(0.0057)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.801
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.985
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.942
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.465
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.614
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.813
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.353
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.834
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.834
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.867
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.842
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.771
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.993
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.946
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.534
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.335
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.785
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.335
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.633
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.814
Epoch: [3] [ 0/60] eta: 0:01:26 lr: 0.000500 loss: 0.1888 (0.1888)
loss_classifier: 0.0211 (0.0211) loss_box_reg: 0.0529 (0.0529)
loss_mask: 0.1116 (0.1116) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0029 (0.0029) time: 1.4483 data: 0.7804 max
mem: 3132
Epoch: [3] [10/60] eta: 0:00:34 lr: 0.000500 loss: 0.1989 (0.2184)
loss_classifier: 0.0280 (0.0328) loss_box_reg: 0.0579 (0.0566)
loss_mask: 0.1126 (0.1228) loss_objectness: 0.0004 (0.0013)
loss_rpn_box_reg: 0.0040 (0.0050) time: 0.6991 data: 0.0769 max
mem: 3409
Epoch: [3] [20/60] eta: 0:00:25 lr: 0.000500 loss: 0.1989 (0.2140)
loss_classifier: 0.0270 (0.0307) loss_box_reg: 0.0579 (0.0554)
loss_mask: 0.1194 (0.1229) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0027 (0.0040) time: 0.6098 data: 0.0080 max
mem: 3409
Epoch: [3] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1789 (0.2006)
loss_classifier: 0.0226 (0.0274) loss_box_reg: 0.0390 (0.0468)
loss_mask: 0.1171 (0.1218) loss_objectness: 0.0002 (0.0010)
loss_rpn_box_reg: 0.0023 (0.0035) time: 0.5718 data: 0.0096 max
mem: 3409
Epoch: [3] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1722 (0.1987)
loss_classifier: 0.0218 (0.0269) loss_box_reg: 0.0320 (0.0467)

```

```

loss_mask: 0.1109 (0.1207) loss_objectness: 0.0002 (0.0010)
loss_rpn_box_reg: 0.0024 (0.0033) time: 0.5643 data: 0.0095 max
mem: 3409
Epoch: [3] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1764 (0.1961)
loss_classifier: 0.0250 (0.0269) loss_box_reg: 0.0332 (0.0465)
loss_mask: 0.1058 (0.1185) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0018 (0.0031) time: 0.5900 data: 0.0113 max
mem: 3409
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1764 (0.1959)
loss_classifier: 0.0283 (0.0270) loss_box_reg: 0.0430 (0.0465)
loss_mask: 0.1093 (0.1183) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0017 (0.0031) time: 0.5978 data: 0.0102 max
mem: 3409
Epoch: [3] Total time: 0:00:36 (0.6087 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:37 model_time: 0.2348 (0.2348)
evaluator_time: 0.0051 (0.0051) time: 0.7442 data: 0.5027 max mem:
3409
Test: [49/50] eta: 0:00:00 model_time: 0.1037 (0.1161)
evaluator_time: 0.0034 (0.0056) time: 0.1183 data: 0.0035 max mem:
3409
Test: Total time: 0:00:07 (0.1430 s / it)
Averaged stats: model_time: 0.1037 (0.1161) evaluator_time: 0.0034
(0.0056)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.829
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.993
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.955
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.549
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.842
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.365
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |

```

```

maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.867
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.869
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.780
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.993
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.947
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.490
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.330
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.791
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.810
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.810
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.667
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.820
Epoch: [4] [ 0/60] eta: 0:01:05 lr: 0.000500 loss: 0.1847 (0.1847)
loss_classifier: 0.0303 (0.0303) loss_box_reg: 0.0470 (0.0470)
loss_mask: 0.1053 (0.1053) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0018 (0.0018) time: 1.0919 data: 0.4482 max
mem: 3409
Epoch: [4] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.1847 (0.1761)
loss_classifier: 0.0267 (0.0258) loss_box_reg: 0.0378 (0.0385)
loss_mask: 0.1053 (0.1082) loss_objectness: 0.0006 (0.0013)
loss_rpn_box_reg: 0.0023 (0.0023) time: 0.6351 data: 0.0488 max
mem: 3409
Epoch: [4] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1661 (0.1715)
loss_classifier: 0.0193 (0.0226) loss_box_reg: 0.0287 (0.0339)
loss_mask: 0.1105 (0.1118) loss_objectness: 0.0006 (0.0013)
loss_rpn_box_reg: 0.0015 (0.0020) time: 0.5681 data: 0.0099 max
mem: 3409
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1761 (0.1765)
loss_classifier: 0.0203 (0.0227) loss_box_reg: 0.0295 (0.0343)
loss_mask: 0.1157 (0.1161) loss_objectness: 0.0005 (0.0010)
loss_rpn_box_reg: 0.0015 (0.0024) time: 0.5754 data: 0.0097 max

```



```
mem: 3409
Epoch: [4] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1798 (0.1763)
loss_classifier: 0.0241 (0.0237) loss_box_reg: 0.0340 (0.0355)
loss_mask: 0.1105 (0.1136) loss_objectness: 0.0004 (0.0010)
loss_rpn_box_reg: 0.0018 (0.0024) time: 0.6095 data: 0.0096 max
mem: 3409
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1686 (0.1776)
loss_classifier: 0.0243 (0.0249) loss_box_reg: 0.0307 (0.0363)
loss_mask: 0.1059 (0.1129) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0021 (0.0026) time: 0.5992 data: 0.0092 max
mem: 3409
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1707 (0.1810)
loss_classifier: 0.0243 (0.0256) loss_box_reg: 0.0365 (0.0379)
loss_mask: 0.1098 (0.1140) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0020 (0.0026) time: 0.5792 data: 0.0075 max
mem: 3409
Epoch: [4] Total time: 0:00:36 (0.6021 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.1503 (0.1503)
evaluator_time: 0.0039 (0.0039) time: 0.5395 data: 0.3835 max mem:
3409
Test: [49/50] eta: 0:00:00 model_time: 0.1037 (0.1115)
evaluator_time: 0.0036 (0.0049) time: 0.1197 data: 0.0040 max mem:
3409
Test: Total time: 0:00:06 (0.1344 s / it)
Averaged stats: model_time: 0.1037 (0.1115) evaluator_time: 0.0036
(0.0049)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.848
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.993
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.955
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.549
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.373
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.874
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.874
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.867
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.883
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.783
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.993
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.947
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.512
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.353
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.796
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.339
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.815
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.815
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.600
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.667
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.824
That's it!

```

So after one epoch of training, we obtain a COCO-style mAP > 50, and a mask mAP of 65.

But what do the predictions look like? Let's take one image in the dataset and verify

```

import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

image =
read_image("drive/MyDrive/_static/img/tv_tutorial/tv_image05.png")
eval_transform = get_transform(train=False)

```

```

model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

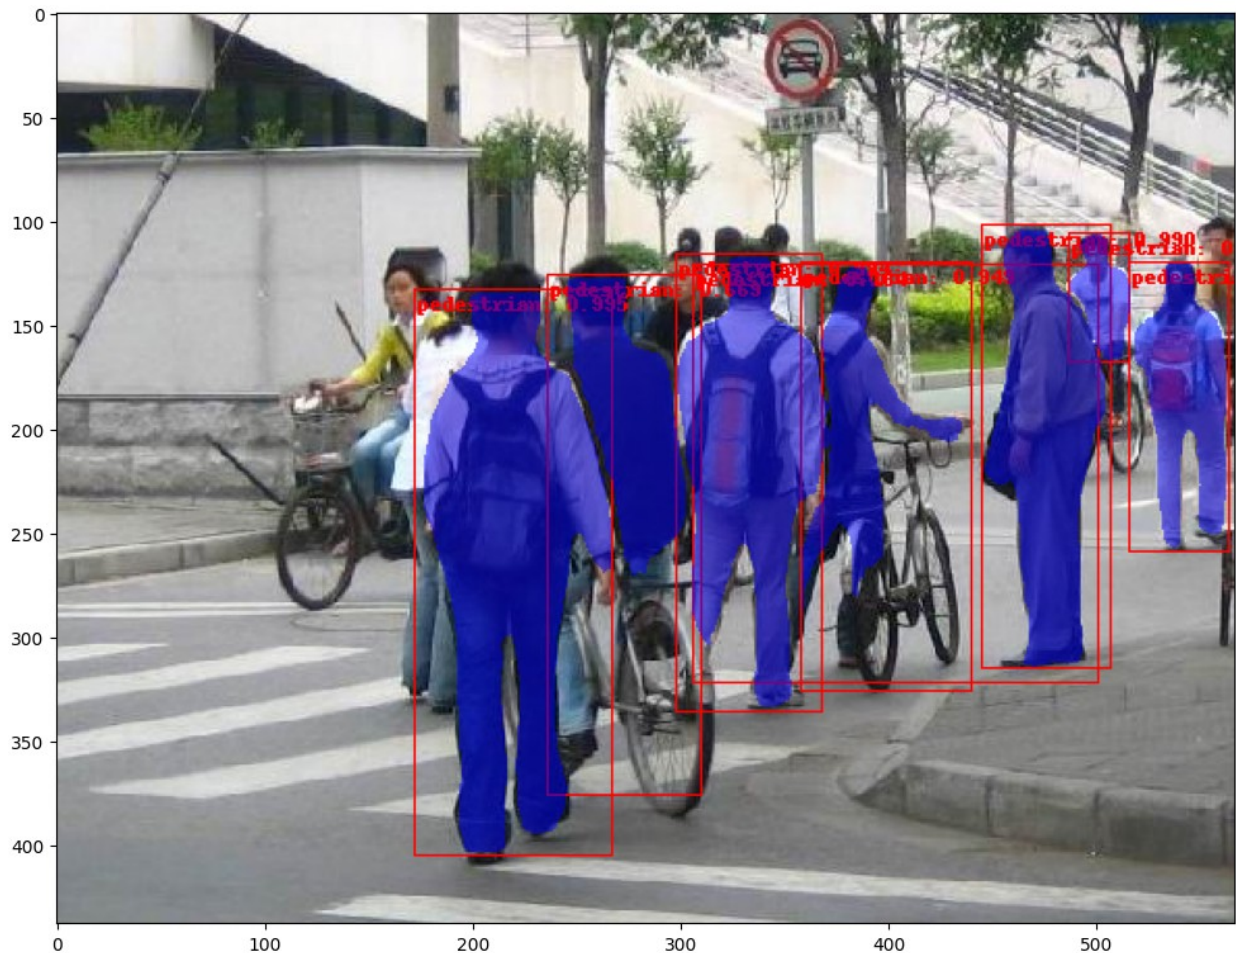
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

<matplotlib.image.AxesImage at 0x7fe26fd7dcf0>

```



The results look good!

Wrapping up

In this tutorial, you have learned how to create your own training pipeline for object detection models on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes and segmentation masks. You also leveraged a Mask R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-GPU training, check `references/detection/train.py`, which is present in the torchvision repository.

You can download a full source file for this tutorial [here](#).

This is for Question 5 Part B and C using the finetuning option

```
# For tips on running notebooks in Google Colab, see
# https://pytorch.org/tutorials/beginner/colab
%matplotlib inline

from google.colab import drive
drive.mount('/content/drive') #Updated the spots where they needed the
locations to change, also downloaded the PennFudan dataset and
tutorial source for the test photo at the end
# Dataset: https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
# Test Photo:
https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425fa
a00c06bcf502/_static/img/tv_tutorial/tv_image05.png
#Source Code:
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

Mounted at /content/drive
```

TorchVision Object Detection Finetuning Tutorial

.. tip::

To get the most of this tutorial, we suggest using this [Colab Version](https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_fineting_instance_segmentation.ipynb). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model on the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection and instance segmentation model on a custom dataset.

.. note ::

This tutorial works only with torchvision version `>=0.16` or `nightly`. If you're using `torchvision<=0.15`, please follow [this tutorial instead](https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425faa00c06bcf502/intermediate_source/torchvision_tutorial.rst).

Defining the Dataset

The reference scripts for training object detection, instance segmentation and person keypoint detection allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return a tuple:

- `image`: `:class:torchvision.tv_tensors.Image` of shape `[3, H, W]`, a pure tensor, or a PIL Image of size `(H, W)`
- `target`: a dict containing the following fields
 - `boxes`, `:class:torchvision.tv_tensors.BoundingBoxes` of shape `[N, 4]`: the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels`, integer `:class:torch.Tensor` of shape `[N]`: the label for each bounding box. `0` represents always the background class.
 - `image_id`, int: an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area`, float `:class:torch.Tensor` of shape `[N]`: the area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd`, uint8 `:class:torch.Tensor` of shape `[N]`: instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks`, `:class:torchvision.tv_tensors.Mask` of shape `[N, H, W]`: the segmentation masks for each one of the objects

If your dataset is compliant with above requirements then it will work for both training and evaluation codes from the reference script. Evaluation code will use scripts from `pycocotools` which can be installed with `pip install pycocotools`.

.. note :: For Windows, please install `pycocotools` from [gautamchitnis_](https://github.com/gautamchitnis/cocoapi) with command

```
pip install
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-
master#subdirectory=PythonAPI
```

One note on the `labels`. The model considers class `0` as background. If your dataset does not contain the background class, you should not have `0` in your `labels`. For example, assuming you have just two classes, *cat* and *dog*, you can define `1` (not `0`) to represent *cats* and `2` to represent *dogs*. So, for instance, if one of the images has both classes, your `labels` tensor should look like `[1, 2]`.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratios), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

Writing a custom dataset for PennFudan

Let's write a dataset for the PennFudan dataset. After [downloading and extracting the zip file](#), we have the following folder structure:

::

```
PennFudanPed/ PedMasks/ FudanPed00001_mask.png FudanPed00002_mask.png
FudanPed00003_mask.png FudanPed00004_mask.png ... PNGImages/ FudanPed00001.png
FudanPed00002.png FudanPed00003.png FudanPed00004.png
```

Here is one example of a pair of images and segmentation masks

So each image has a corresponding segmentation mask, where each color correspond to a different instance. Let's write a :class:torch.utils.data.Dataset class for this dataset. In the code below, we are wrapping images, bounding boxes and masks into torchvision.TVTensor classes so that we will be able to apply torchvision built-in transformations ([new Transforms API](#)) for the given object detection and segmentation task. Namely, image tensors will be wrapped by :class:torchvision.tv_tensors.Image, bounding boxes into :class:torchvision.tv_tensors.BoundingBoxes and masks into :class:torchvision.tv_tensors.Mask. As torchvision.TVTensor are :class:torch.Tensor subclasses, wrapped objects are also tensors and inherit the plain :class:torch.Tensor API. For more information about torchvision tv_tensors see [this documentation](#).

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
```



```

        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

That's all for the dataset. Now let's define a model that can perform predictions on this dataset.

Defining your model

In this tutorial, we will be using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.

Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

There are two common situations where one might want to modify one of the available models in TorchVision Model Zoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

1 - Finetuning from a pretrained model

Let's suppose that you want to start from a model pre-trained on COCO and want to finetune it for your particular classes. Here is a possible way of doing it:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

# load a model pre-trained on COCO
model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)

# replace the classifier with a new one, that has
# num_classes which is user-defined
num_classes = 2 # 1 class (person) + background
# get number of input features for the classifier
in_features = model.roi_heads.box_predictor.cls_score.in_features
# replace the pre-trained head with a new one
model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

Downloading:
"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth" to
/root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth
100%|██████████| 160M/160M [00:01<00:00, 157MB/s]
```

Object detection and instance segmentation model for PennFudan Dataset

In our case, we want to finetune from a pre-trained model, given that our dataset is very small, so we will be following approach number 1.

Here we want to also compute the instance segmentation masks, so we will be using Mask R-CNN:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
    torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
    num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
    model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256
    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(
        in_features_mask,
        hidden_layer,
        num_classes
    )

    return model
```

That's it, this will make `model` be ready to be trained and evaluated on your custom dataset.

Putting everything together

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py` and `references/detection/utils.py`. Just download everything under `references/detection` to your folder and use them here. On Linux if you have `wget`, you can download them using below commands:

```

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API
<https://pytorch.org/vision/stable/transforms.html>`_
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#
# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
        transforms.append(T.ToDtype(torch.float, scale=True))
        transforms.append(T.ToPureTensor())
    return T.Compose(transforms)

# Testing ``forward()`` method (Optional)
# -----
#
# Before iterating over the dataset, it's good to see what the model
# expects during training and inference time on sample data.
import utils

model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
data_loader = torch.utils.data.DataLoader(

```

```

        dataset,
        batch_size=2,
        shuffle=True,
        num_workers=4,
        collate_fn=utils.collate_fn
    )

# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
print(output)

# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions
print(predictions[0])

/usr/local/lib/python3.10/dist-packages/torch/utils/data/
dataloader.py:557: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

{'loss_classifier': tensor(0.0989, grad_fn=<NllLossBackward0>),
'loss_box_reg': tensor(0.0605, grad_fn=<DivBackward0>),
'loss_objectness': tensor(0.0056,
grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), 'loss_rpn_box_reg':
tensor(0.0073, grad_fn=<DivBackward0>)}
{'boxes': tensor([], size=(0, 4), grad_fn=<StackBackward0>), 'labels':
tensor([], dtype=torch.int64), 'scores': tensor([],
grad_fn=<IndexBackward0>)}

```

Let's now write the main function which performs the training and the validation:

```

from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',

```

```

get_transform(train=True))
dataset_test = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 10

```

```

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

Downloading:

```

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth" to
/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth
100%|██████████| 170M/170M [00:01<00:00, 92.5MB/s]

```

```

Epoch: [0] [ 0/60] eta: 0:09:22 lr: 0.000090 loss: 4.0232 (4.0232)
loss_classifier: 0.6608 (0.6608) loss_box_reg: 0.1815 (0.1815)
loss_mask: 3.1586 (3.1586) loss_objectness: 0.0194 (0.0194)
loss_rpn_box_reg: 0.0029 (0.0029) time: 9.3771 data: 1.6565 max
mem: 2596
Epoch: [0] [10/60] eta: 0:01:07 lr: 0.000936 loss: 1.7460 (2.3481)
loss_classifier: 0.4496 (0.4475) loss_box_reg: 0.3494 (0.3553)
loss_mask: 1.0930 (1.5102) loss_objectness: 0.0304 (0.0277)
loss_rpn_box_reg: 0.0061 (0.0074) time: 1.3599 data: 0.1567 max
mem: 2978
Epoch: [0] [20/60] eta: 0:00:39 lr: 0.001783 loss: 1.0728 (1.5656)
loss_classifier: 0.2308 (0.3128) loss_box_reg: 0.2986 (0.3006)
loss_mask: 0.4108 (0.9232) loss_objectness: 0.0153 (0.0225)
loss_rpn_box_reg: 0.0058 (0.0066) time: 0.5580 data: 0.0099 max
mem: 2978
Epoch: [0] [30/60] eta: 0:00:25 lr: 0.002629 loss: 0.5781 (1.2421)
loss_classifier: 0.1208 (0.2424) loss_box_reg: 0.1701 (0.2828)
loss_mask: 0.2195 (0.6921) loss_objectness: 0.0083 (0.0175)
loss_rpn_box_reg: 0.0041 (0.0072) time: 0.5665 data: 0.0111 max
mem: 3066
Epoch: [0] [40/60] eta: 0:00:15 lr: 0.003476 loss: 0.5432 (1.0732)
loss_classifier: 0.0734 (0.2017) loss_box_reg: 0.1987 (0.2805)
loss_mask: 0.2021 (0.5695) loss_objectness: 0.0044 (0.0144)
loss_rpn_box_reg: 0.0059 (0.0072) time: 0.5562 data: 0.0098 max
mem: 3066
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.4499 (0.9406)
loss_classifier: 0.0452 (0.1706) loss_box_reg: 0.1744 (0.2575)
loss_mask: 0.1770 (0.4932) loss_objectness: 0.0026 (0.0123)
loss_rpn_box_reg: 0.0052 (0.0070) time: 0.5357 data: 0.0101 max
mem: 3066
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3485 (0.8478)

```

```

loss_classifier: 0.0368 (0.1514) loss_box_reg: 0.1203 (0.2363)
loss_mask: 0.1702 (0.4425) loss_objectness: 0.0015 (0.0108)
loss_rpn_box_reg: 0.0044 (0.0069) time: 0.5412 data: 0.0090 max
mem: 3066
Epoch: [0] Total time: 0:00:42 (0.7033 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.2204 (0.2204)
evaluator_time: 0.0211 (0.0211) time: 0.5590 data: 0.3109 max mem:
3066
Test: [49/50] eta: 0:00:00 model_time: 0.1112 (0.1196)
evaluator_time: 0.0061 (0.0107) time: 0.1321 data: 0.0052 max mem:
3066
Test: Total time: 0:00:07 (0.1491 s / it)
Averaged stats: model_time: 0.1112 (0.1196) evaluator_time: 0.0061
(0.0107)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.662
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.849
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.357
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.559
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.680
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.338
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.727
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.727
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.736
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.693
Average Precision (AP) @[ IoU=0.50 | area= all |

```

```

maxDets=100 ] = 0.960
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.868
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.369
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.459
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.353
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.751
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.754
Epoch: [1] [ 0/60] eta: 0:01:02 lr: 0.005000 loss: 0.3659 (0.3659)
loss_classifier: 0.0570 (0.0570) loss_box_reg: 0.1407 (0.1407)
loss_mask: 0.1630 (0.1630) loss_objectness: 0.0016 (0.0016)
loss_rpn_box_reg: 0.0036 (0.0036) time: 1.0337 data: 0.4756 max
mem: 3066
Epoch: [1] [10/60] eta: 0:00:30 lr: 0.005000 loss: 0.3129 (0.3190)
loss_classifier: 0.0385 (0.0399) loss_box_reg: 0.1143 (0.1172)
loss_mask: 0.1380 (0.1558) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0049 (0.0052) time: 0.6191 data: 0.0496 max
mem: 3066
Epoch: [1] [20/60] eta: 0:00:24 lr: 0.005000 loss: 0.2991 (0.3156)
loss_classifier: 0.0334 (0.0396) loss_box_reg: 0.0931 (0.1061)
loss_mask: 0.1543 (0.1634) loss_objectness: 0.0007 (0.0016)
loss_rpn_box_reg: 0.0048 (0.0050) time: 0.5884 data: 0.0081 max
mem: 3066
Epoch: [1] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2847 (0.3090)
loss_classifier: 0.0421 (0.0408) loss_box_reg: 0.0901 (0.1063)
loss_mask: 0.1517 (0.1552) loss_objectness: 0.0015 (0.0017)
loss_rpn_box_reg: 0.0043 (0.0049) time: 0.5998 data: 0.0097 max
mem: 3464
Epoch: [1] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2822 (0.3030)
loss_classifier: 0.0421 (0.0400) loss_box_reg: 0.0887 (0.1028)
loss_mask: 0.1373 (0.1540) loss_objectness: 0.0007 (0.0014)
loss_rpn_box_reg: 0.0038 (0.0048) time: 0.6153 data: 0.0090 max
mem: 3464
Epoch: [1] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2812 (0.3023)
loss_classifier: 0.0341 (0.0400) loss_box_reg: 0.0887 (0.1020)
loss_mask: 0.1392 (0.1538) loss_objectness: 0.0005 (0.0014)

```



```
loss_rpn_box_reg: 0.0038 (0.0051)  time: 0.6257  data: 0.0098  max
mem: 3464
Epoch: [1]  [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.2764 (0.2969)
loss_classifier: 0.0402 (0.0396)  loss_box_reg: 0.0861 (0.0994)
loss_mask: 0.1343 (0.1512)  loss_objectness: 0.0006 (0.0014)
loss_rpn_box_reg: 0.0057 (0.0052)  time: 0.6134  data: 0.0094  max
mem: 3464
Epoch: [1] Total time: 0:00:36 (0.6161 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:39  model_time: 0.2075 (0.2075)
evaluator_time: 0.0281 (0.0281)  time: 0.7880  data: 0.5507  max mem:
3464
Test:  [49/50]  eta: 0:00:00  model_time: 0.0969 (0.1079)
evaluator_time: 0.0032 (0.0060)  time: 0.1103  data: 0.0038  max mem:
3464
Test: Total time: 0:00:06 (0.1350 s / it)
Averaged stats: model_time: 0.0969 (0.1079)  evaluator_time: 0.0032
(0.0060)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.777
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.966
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.916
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.367
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.581
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.801
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.396
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.827
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.827
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.725
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.846
IoU metric: segm
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.743
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.971
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.921
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.354
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.532
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.761
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.377
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.798
Epoch: [2] [ 0/60] eta: 0:00:57 lr: 0.005000 loss: 0.2624 (0.2624)
loss_classifier: 0.0328 (0.0328) loss_box_reg: 0.0768 (0.0768)
loss_mask: 0.1485 (0.1485) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0043 (0.0043) time: 0.9614 data: 0.3912 max
mem: 3464
Epoch: [2] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2395 (0.2361)
loss_classifier: 0.0326 (0.0314) loss_box_reg: 0.0622 (0.0653)
loss_mask: 0.1212 (0.1340) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0043 (0.0047) time: 0.5913 data: 0.0433 max
mem: 3464
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2164 (0.2381)
loss_classifier: 0.0306 (0.0326) loss_box_reg: 0.0592 (0.0659)
loss_mask: 0.1212 (0.1343) loss_objectness: 0.0005 (0.0006)
loss_rpn_box_reg: 0.0042 (0.0047) time: 0.5783 data: 0.0097 max
mem: 3464
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2076 (0.2267)
loss_classifier: 0.0270 (0.0301) loss_box_reg: 0.0574 (0.0631)
loss_mask: 0.1179 (0.1286) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0042) time: 0.6047 data: 0.0093 max
mem: 3464
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2118 (0.2347)
loss_classifier: 0.0270 (0.0311) loss_box_reg: 0.0595 (0.0677)
loss_mask: 0.1246 (0.1308) loss_objectness: 0.0002 (0.0005)
loss_rpn_box_reg: 0.0037 (0.0046) time: 0.5985 data: 0.0094 max
mem: 3464
```

```
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2745 (0.2435)
loss_classifier: 0.0311 (0.0325) loss_box_reg: 0.0837 (0.0720)
loss_mask: 0.1364 (0.1328) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0054 (0.0054) time: 0.5857 data: 0.0099 max
mem: 3464
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2684 (0.2462)
loss_classifier: 0.0331 (0.0332) loss_box_reg: 0.0837 (0.0742)
loss_mask: 0.1263 (0.1328) loss_objectness: 0.0006 (0.0008)
loss_rpn_box_reg: 0.0035 (0.0052) time: 0.5700 data: 0.0085 max
mem: 3464
Epoch: [2] Total time: 0:00:35 (0.5971 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1593 (0.1593)
evaluator_time: 0.0136 (0.0136) time: 0.5565 data: 0.3816 max mem:
3464
Test: [49/50] eta: 0:00:00 model_time: 0.0989 (0.1055)
evaluator_time: 0.0030 (0.0047) time: 0.1097 data: 0.0037 max mem:
3464
Test: Total time: 0:00:06 (0.1273 s / it)
Averaged stats: model_time: 0.0989 (0.1055) evaluator_time: 0.0030
(0.0047)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.746
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.352
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.579
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.760
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.385
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.798
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.739
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.972
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.890
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.394
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.495
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.751
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.376
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.781
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.783
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.667
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.725
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.792
Epoch: [3] [ 0/60] eta: 0:00:59 lr: 0.000500 loss: 0.2168 (0.2168)
loss_classifier: 0.0278 (0.0278) loss_box_reg: 0.0707 (0.0707)
loss_mask: 0.1120 (0.1120) loss_objectness: 0.0021 (0.0021)
loss_rpn_box_reg: 0.0043 (0.0043) time: 0.9965 data: 0.3444 max
mem: 3464
Epoch: [3] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.2418 (0.2369)
loss_classifier: 0.0279 (0.0308) loss_box_reg: 0.0656 (0.0665)
loss_mask: 0.1386 (0.1345) loss_objectness: 0.0007 (0.0010)
loss_rpn_box_reg: 0.0036 (0.0041) time: 0.6190 data: 0.0402 max
mem: 3464
Epoch: [3] [20/60] eta: 0:00:25 lr: 0.000500 loss: 0.2288 (0.2314)
loss_classifier: 0.0299 (0.0303) loss_box_reg: 0.0602 (0.0632)
loss_mask: 0.1326 (0.1333) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0032 (0.0037) time: 0.6124 data: 0.0094 max
mem: 3464
Epoch: [3] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.2084 (0.2180)
loss_classifier: 0.0284 (0.0304) loss_box_reg: 0.0457 (0.0558)
loss_mask: 0.1106 (0.1270) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0030 (0.0038) time: 0.6178 data: 0.0095 max
mem: 3464
Epoch: [3] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1870 (0.2101)
loss_classifier: 0.0254 (0.0296) loss_box_reg: 0.0372 (0.0535)

```

```
loss_mask: 0.1054 (0.1226) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0022 (0.0035) time: 0.5934 data: 0.0098 max
mem: 3464
Epoch: [3] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1886 (0.2074)
loss_classifier: 0.0208 (0.0285) loss_box_reg: 0.0435 (0.0524)
loss_mask: 0.1120 (0.1220) loss_objectness: 0.0004 (0.0010)
loss_rpn_box_reg: 0.0020 (0.0035) time: 0.6067 data: 0.0106 max
mem: 3464
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1886 (0.2092)
loss_classifier: 0.0252 (0.0288) loss_box_reg: 0.0435 (0.0524)
loss_mask: 0.1151 (0.1234) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0031 (0.0036) time: 0.6042 data: 0.0099 max
mem: 3464
Epoch: [3] Total time: 0:00:36 (0.6153 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1925 (0.1925)
evaluator_time: 0.0142 (0.0142) time: 0.5605 data: 0.3520 max mem:
3464
Test: [49/50] eta: 0:00:00 model_time: 0.1048 (0.1132)
evaluator_time: 0.0047 (0.0065) time: 0.1237 data: 0.0090 max mem:
3464
Test: Total time: 0:00:07 (0.1403 s / it)
Averaged stats: model_time: 0.1048 (0.1132) evaluator_time: 0.0047
(0.0065)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.828
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.368
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.650
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.846
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.419
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.866
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.866
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
```

```

maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.880
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.758
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.378
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.517
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.772
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.382
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.797
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
Epoch: [4] [ 0/60] eta: 0:01:03 lr: 0.000500 loss: 0.1764 (0.1764)
loss_classifier: 0.0218 (0.0218) loss_box_reg: 0.0406 (0.0406)
loss_mask: 0.1091 (0.1091) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0045 (0.0045) time: 1.0649 data: 0.3697 max
mem: 3464
Epoch: [4] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.2057 (0.2016)
loss_classifier: 0.0229 (0.0288) loss_box_reg: 0.0416 (0.0444)
loss_mask: 0.1205 (0.1236) loss_objectness: 0.0006 (0.0006)
loss_rpn_box_reg: 0.0044 (0.0043) time: 0.6200 data: 0.0414 max
mem: 3464
Epoch: [4] [20/60] eta: 0:00:25 lr: 0.000500 loss: 0.2057 (0.2106)
loss_classifier: 0.0295 (0.0301) loss_box_reg: 0.0496 (0.0495)
loss_mask: 0.1205 (0.1261) loss_objectness: 0.0007 (0.0008)
loss_rpn_box_reg: 0.0037 (0.0041) time: 0.6055 data: 0.0097 max
mem: 3700
Epoch: [4] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1907 (0.2052)
loss_classifier: 0.0274 (0.0292) loss_box_reg: 0.0480 (0.0478)
loss_mask: 0.1135 (0.1238) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0029 (0.0036) time: 0.6023 data: 0.0097 max

```

```
mem: 3700
Epoch: [4] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1797 (0.2047)
loss_classifier: 0.0246 (0.0302) loss_box_reg: 0.0437 (0.0480)
loss_mask: 0.1135 (0.1223) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0027 (0.0034) time: 0.5845 data: 0.0109 max
mem: 3700
Epoch: [4] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1739 (0.1996)
loss_classifier: 0.0253 (0.0290) loss_box_reg: 0.0385 (0.0468)
loss_mask: 0.1050 (0.1199) loss_objectness: 0.0005 (0.0008)
loss_rpn_box_reg: 0.0023 (0.0032) time: 0.5931 data: 0.0120 max
mem: 3700
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1693 (0.1997)
loss_classifier: 0.0275 (0.0292) loss_box_reg: 0.0385 (0.0467)
loss_mask: 0.1055 (0.1197) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0023 (0.0032) time: 0.6095 data: 0.0095 max
mem: 3700
Epoch: [4] Total time: 0:00:36 (0.6121 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.1775 (0.1775)
evaluator_time: 0.0206 (0.0206) time: 0.5346 data: 0.3344 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.0987 (0.1066)
evaluator_time: 0.0026 (0.0048) time: 0.1097 data: 0.0037 max mem:
3700
Test: Total time: 0:00:06 (0.1277 s / it)
Averaged stats: model_time: 0.0987 (0.1066) evaluator_time: 0.0026
(0.0048)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.840
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.368
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.674
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.426
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.881
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.881
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.897
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.757
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.970
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.913
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.378
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.530
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.770
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.385
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.799
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.801
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.810
Epoch: [5] [ 0/60] eta: 0:00:59  lr: 0.000500  loss: 0.2363 (0.2363)
loss_classifier: 0.0346 (0.0346)  loss_box_reg: 0.0638 (0.0638)
loss_mask: 0.1342 (0.1342)  loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0035 (0.0035)  time: 0.9918  data: 0.4285  max
mem: 3700
Epoch: [5] [10/60] eta: 0:00:30  lr: 0.000500  loss: 0.1874 (0.1939)
loss_classifier: 0.0242 (0.0280)  loss_box_reg: 0.0449 (0.0424)
loss_mask: 0.1151 (0.1199)  loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0021 (0.0025)  time: 0.6078  data: 0.0489  max
mem: 3700
Epoch: [5] [20/60] eta: 0:00:24  lr: 0.000500  loss: 0.1931 (0.2105)
loss_classifier: 0.0356 (0.0326)  loss_box_reg: 0.0477 (0.0501)
loss_mask: 0.1145 (0.1234)  loss_objectness: 0.0004 (0.0014)
loss_rpn_box_reg: 0.0021 (0.0030)  time: 0.5918  data: 0.0103  max
mem: 3700
Epoch: [5] [30/60] eta: 0:00:17  lr: 0.000500  loss: 0.1931 (0.2008)
loss_classifier: 0.0284 (0.0289)  loss_box_reg: 0.0477 (0.0472)

```



```
loss_mask: 0.1142 (0.1206) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5921 data: 0.0098 max
mem: 3700
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1741 (0.1982)
loss_classifier: 0.0214 (0.0286) loss_box_reg: 0.0373 (0.0459)
loss_mask: 0.1122 (0.1197) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0026 (0.0030) time: 0.5814 data: 0.0102 max
mem: 3700
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1935 (0.1971)
loss_classifier: 0.0248 (0.0283) loss_box_reg: 0.0424 (0.0465)
loss_mask: 0.1114 (0.1183) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0034 (0.0031) time: 0.6012 data: 0.0099 max
mem: 3700
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1741 (0.1930)
loss_classifier: 0.0218 (0.0271) loss_box_reg: 0.0325 (0.0444)
loss_mask: 0.1061 (0.1175) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0026 (0.0031) time: 0.6125 data: 0.0085 max
mem: 3700
Epoch: [5] Total time: 0:00:36 (0.6051 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1591 (0.1591)
evaluator_time: 0.0116 (0.0116) time: 0.5638 data: 0.3914 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.1025 (0.1087)
evaluator_time: 0.0039 (0.0052) time: 0.1167 data: 0.0041 max mem:
3700
Test: Total time: 0:00:06 (0.1351 s / it)
Averaged stats: model_time: 0.1025 (0.1087) evaluator_time: 0.0039
(0.0052)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.837
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.930
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.858
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.423
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.875
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.875
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.812
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.892
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.762
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.969
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.932
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.544
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.385
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.804
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.806
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.813
Epoch: [6] [ 0/60] eta: 0:01:24 lr: 0.000050 loss: 0.2115 (0.2115)
loss_classifier: 0.0332 (0.0332) loss_box_reg: 0.0507 (0.0507)
loss_mask: 0.1236 (0.1236) loss_objectness: 0.0005 (0.0005)
loss_rpn_box_reg: 0.0036 (0.0036) time: 1.4092 data: 0.6883 max
mem: 3700
Epoch: [6] [10/60] eta: 0:00:33 lr: 0.000050 loss: 0.2112 (0.1921)
loss_classifier: 0.0329 (0.0284) loss_box_reg: 0.0381 (0.0406)
loss_mask: 0.1210 (0.1200) loss_objectness: 0.0002 (0.0004)
loss_rpn_box_reg: 0.0027 (0.0027) time: 0.6647 data: 0.0694 max
mem: 3700
Epoch: [6] [20/60] eta: 0:00:25 lr: 0.000050 loss: 0.1861 (0.1938)
loss_classifier: 0.0253 (0.0281) loss_box_reg: 0.0354 (0.0426)
loss_mask: 0.1169 (0.1195) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0027 (0.0032) time: 0.6113 data: 0.0086 max

```

```
mem: 3700
Epoch: [6] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1861 (0.1884)
loss_classifier: 0.0226 (0.0266) loss_box_reg: 0.0354 (0.0416)
loss_mask: 0.1085 (0.1164) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0024 (0.0031) time: 0.6113 data: 0.0095 max
mem: 3700
Epoch: [6] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1644 (0.1829)
loss_classifier: 0.0184 (0.0256) loss_box_reg: 0.0317 (0.0396)
loss_mask: 0.1084 (0.1143) loss_objectness: 0.0002 (0.0005)
loss_rpn_box_reg: 0.0022 (0.0029) time: 0.5842 data: 0.0091 max
mem: 3700
Epoch: [6] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.1649 (0.1882)
loss_classifier: 0.0205 (0.0266) loss_box_reg: 0.0322 (0.0410)
loss_mask: 0.1118 (0.1170) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0029) time: 0.5804 data: 0.0091 max
mem: 3700
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1811 (0.1904)
loss_classifier: 0.0275 (0.0268) loss_box_reg: 0.0370 (0.0415)
loss_mask: 0.1207 (0.1183) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0030) time: 0.5921 data: 0.0083 max
mem: 3700
Epoch: [6] Total time: 0:00:36 (0.6128 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:40 model_time: 0.1714 (0.1714)
evaluator_time: 0.0194 (0.0194) time: 0.8045 data: 0.6114 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.1001 (0.1105)
evaluator_time: 0.0025 (0.0057) time: 0.1101 data: 0.0037 max mem:
3700
Test: Total time: 0:00:07 (0.1411 s / it)
Averaged stats: model_time: 0.1001 (0.1105) evaluator_time: 0.0025
(0.0057)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.836
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.941
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.655
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
```

```

maxDets=100 ] = 0.858
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.420
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.874
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.874
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.812
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.892
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.763
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.970
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.917
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.540
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.386
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.805
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [7] [ 0/60] eta: 0:00:58 lr: 0.000050 loss: 0.1546 (0.1546)
loss_classifier: 0.0144 (0.0144) loss_box_reg: 0.0233 (0.0233)
loss_mask: 0.1152 (0.1152) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0015 (0.0015) time: 0.9699 data: 0.3689 max
mem: 3700
Epoch: [7] [10/60] eta: 0:00:31 lr: 0.000050 loss: 0.1793 (0.1782)
loss_classifier: 0.0260 (0.0260) loss_box_reg: 0.0386 (0.0390)
loss_mask: 0.1070 (0.1103) loss_objectness: 0.0002 (0.0003)
loss_rpn_box_reg: 0.0030 (0.0026) time: 0.6248 data: 0.0409 max
mem: 3700
Epoch: [7] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1793 (0.1823)

```

```
loss_classifier: 0.0254 (0.0261) loss_box_reg: 0.0397 (0.0411)
loss_mask: 0.1049 (0.1116) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0030 (0.0029) time: 0.6018 data: 0.0095 max
mem: 3700
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1716 (0.1806)
loss_classifier: 0.0230 (0.0257) loss_box_reg: 0.0397 (0.0396)
loss_mask: 0.1068 (0.1118) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0027 (0.0029) time: 0.5863 data: 0.0099 max
mem: 3700
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1733 (0.1817)
loss_classifier: 0.0228 (0.0256) loss_box_reg: 0.0364 (0.0394)
loss_mask: 0.1077 (0.1133) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5703 data: 0.0087 max
mem: 3700
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1733 (0.1853)
loss_classifier: 0.0205 (0.0257) loss_box_reg: 0.0364 (0.0405)
loss_mask: 0.1171 (0.1155) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0030) time: 0.5830 data: 0.0088 max
mem: 3700
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1799 (0.1869)
loss_classifier: 0.0235 (0.0261) loss_box_reg: 0.0418 (0.0411)
loss_mask: 0.1171 (0.1161) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5975 data: 0.0084 max
mem: 3700
Epoch: [7] Total time: 0:00:36 (0.6016 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1995 (0.1995)
evaluator_time: 0.0102 (0.0102) time: 0.5437 data: 0.3323 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.0981 (0.1083)
evaluator_time: 0.0026 (0.0047) time: 0.1100 data: 0.0036 max mem:
3700
Test: Total time: 0:00:06 (0.1294 s / it)
Averaged stats: model_time: 0.0981 (0.1083) evaluator_time: 0.0026
(0.0047)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.839
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.941
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.424
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.879
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.879
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.896
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.763
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.534
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.386
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [8] [ 0/60] eta: 0:01:05 lr: 0.000050 loss: 0.1549 (0.1549)
loss_classifier: 0.0242 (0.0242) loss_box_reg: 0.0308 (0.0308)
loss_mask: 0.0977 (0.0977) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0020 (0.0020) time: 1.0998 data: 0.4827 max
mem: 3700
Epoch: [8] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1596 (0.1569)
loss_classifier: 0.0237 (0.0219) loss_box_reg: 0.0263 (0.0260)
loss_mask: 0.1095 (0.1062) loss_objectness: 0.0003 (0.0005)

```

```
loss_rpn_box_reg: 0.0020 (0.0022)  time: 0.5801  data: 0.0547  max
mem: 3700
Epoch: [8]  [20/60]  eta: 0:00:23  lr: 0.000050  loss: 0.1797 (0.1906)
loss_classifier: 0.0262 (0.0277)  loss_box_reg: 0.0315 (0.0408)
loss_mask: 0.1151 (0.1184)  loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5693  data: 0.0113  max
mem: 3700
Epoch: [8]  [30/60]  eta: 0:00:18  lr: 0.000050  loss: 0.1824 (0.1826)
loss_classifier: 0.0262 (0.0254)  loss_box_reg: 0.0394 (0.0391)
loss_mask: 0.1151 (0.1145)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0028 (0.0029)  time: 0.6157  data: 0.0108  max
mem: 3700
Epoch: [8]  [40/60]  eta: 0:00:12  lr: 0.000050  loss: 0.1793 (0.1847)
loss_classifier: 0.0236 (0.0260)  loss_box_reg: 0.0394 (0.0401)
loss_mask: 0.1075 (0.1150)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.6154  data: 0.0103  max
mem: 3700
Epoch: [8]  [50/60]  eta: 0:00:05  lr: 0.000050  loss: 0.1794 (0.1860)
loss_classifier: 0.0266 (0.0254)  loss_box_reg: 0.0393 (0.0395)
loss_mask: 0.1112 (0.1176)  loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5890  data: 0.0097  max
mem: 3700
Epoch: [8]  [59/60]  eta: 0:00:00  lr: 0.000050  loss: 0.1794 (0.1867)
loss_classifier: 0.0276 (0.0260)  loss_box_reg: 0.0365 (0.0408)
loss_mask: 0.1053 (0.1166)  loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0028)  time: 0.5954  data: 0.0087  max
mem: 3700
Epoch: [8] Total time: 0:00:36 (0.6055 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:28  model_time: 0.1800 (0.1800)
evaluator_time: 0.0126 (0.0126)  time: 0.5767  data: 0.3823  max mem:
3700
Test:  [49/50]  eta: 0:00:00  model_time: 0.1195 (0.1192)
evaluator_time: 0.0047 (0.0068)  time: 0.1389  data: 0.0058  max mem:
3700
Test: Total time: 0:00:07 (0.1499 s / it)
Averaged stats: model_time: 0.1195 (0.1192)  evaluator_time: 0.0047
(0.0068)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.840
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.979
Average Precision  (AP) @[ IoU=0.75      | area=  all |
```

```

maxDets=100 ] = 0.931
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.861
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.425
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.878
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.878
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.895
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.765
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.535
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.777
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.385
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [9] [ 0/60] eta: 0:01:36 lr: 0.000005 loss: 0.1179 (0.1179)
loss_classifier: 0.0076 (0.0076) loss_box_reg: 0.0123 (0.0123)
loss_mask: 0.0972 (0.0972) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0006 (0.0006) time: 1.6021 data: 0.8491 max
mem: 3700

```



```
Epoch: [9] [10/60] eta: 0:00:34 lr: 0.000005 loss: 0.1728 (0.1847)
loss_classifier: 0.0242 (0.0255) loss_box_reg: 0.0357 (0.0386)
loss_mask: 0.1076 (0.1176) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0020 (0.0023) time: 0.6896 data: 0.0844 max
mem: 3700
Epoch: [9] [20/60] eta: 0:00:26 lr: 0.000005 loss: 0.1840 (0.1879)
loss_classifier: 0.0292 (0.0284) loss_box_reg: 0.0365 (0.0414)
loss_mask: 0.1125 (0.1150) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0025) time: 0.6100 data: 0.0099 max
mem: 3700
Epoch: [9] [30/60] eta: 0:00:19 lr: 0.000005 loss: 0.1834 (0.1829)
loss_classifier: 0.0286 (0.0269) loss_box_reg: 0.0356 (0.0394)
loss_mask: 0.1061 (0.1132) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0031 (0.0026) time: 0.6119 data: 0.0100 max
mem: 3778
Epoch: [9] [40/60] eta: 0:00:12 lr: 0.000005 loss: 0.1670 (0.1854)
loss_classifier: 0.0234 (0.0268) loss_box_reg: 0.0321 (0.0401)
loss_mask: 0.1090 (0.1150) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0032 (0.0028) time: 0.5948 data: 0.0085 max
mem: 3778
Epoch: [9] [50/60] eta: 0:00:06 lr: 0.000005 loss: 0.1898 (0.1883)
loss_classifier: 0.0244 (0.0267) loss_box_reg: 0.0388 (0.0407)
loss_mask: 0.1180 (0.1174) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0032 (0.0029) time: 0.5872 data: 0.0101 max
mem: 3778
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1793 (0.1891)
loss_classifier: 0.0253 (0.0269) loss_box_reg: 0.0378 (0.0417)
loss_mask: 0.1127 (0.1169) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0029 (0.0030) time: 0.5977 data: 0.0096 max
mem: 3778
Epoch: [9] Total time: 0:00:37 (0.6213 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:42 model_time: 0.2411 (0.2411)
evaluator_time: 0.0180 (0.0180) time: 0.8438 data: 0.5829 max mem:
3778
Test: [49/50] eta: 0:00:00 model_time: 0.0998 (0.1116)
evaluator_time: 0.0029 (0.0051) time: 0.1125 data: 0.0038 max mem:
3778
Test: Total time: 0:00:06 (0.1391 s / it)
Averaged stats: model_time: 0.0998 (0.1116) evaluator_time: 0.0029
(0.0051)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.839
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.931
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.424
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.877
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.877
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.894
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.767
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.535
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.779
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.386
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.804
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.806
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
That's it!

```

So after one epoch of training, we obtain a COCO-style mAP > 50, and a mask mAP of 65.

But what do the predictions look like? Let's take one image in the dataset and verify

```
import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

image =
read_image("drive/MyDrive/_static/img/tv_tutorial/tv_image05.png")
eval_transform = get_transform(train=False)

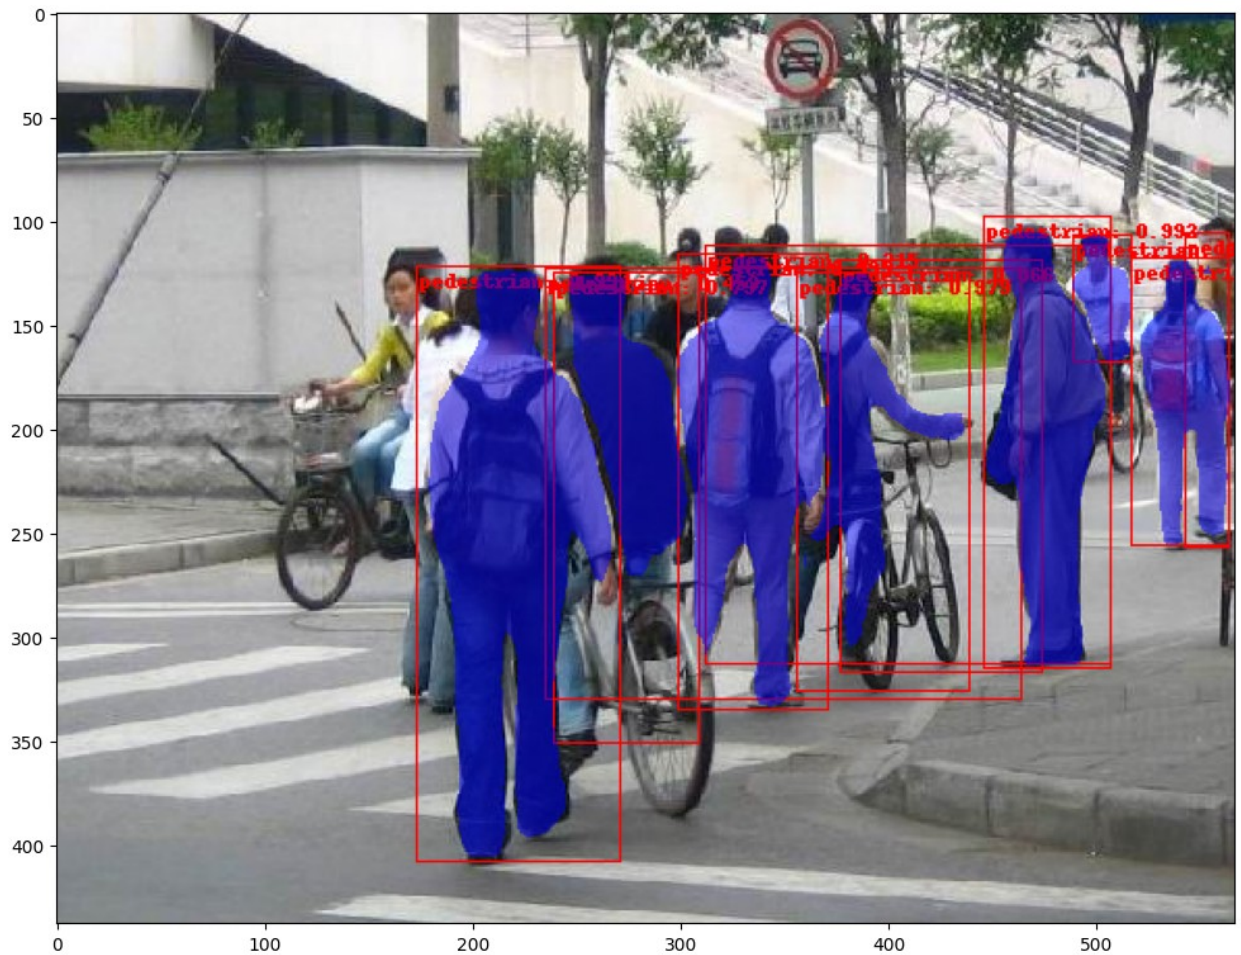
model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

<matplotlib.image.AxesImage at 0x7ef574e209a0>
```



pred_labels

```
['pedestrian: 0.995',  
'pedestrian: 0.994',  
'pedestrian: 0.992',  
'pedestrian: 0.988',  
'pedestrian: 0.979',  
'pedestrian: 0.797',  
'pedestrian: 0.453',  
'pedestrian: 0.215',  
'pedestrian: 0.089',  
'pedestrian: 0.066',  
'pedestrian: 0.055']
```

pred_boxes

```
tensor([[299, 116, 371, 335],  
        [173, 122, 271, 408],  
        [446, 98, 507, 315],  
        [517, 118, 564, 256],  
        [356, 125, 439, 326],
```

```
[239, 125, 309, 351],  
[235, 123, 464, 330],  
[312, 112, 507, 313],  
[489, 107, 517, 168],  
[377, 119, 474, 317],  
[543, 106, 564, 257]], device='cuda:0')
```

The results look good!

Wrapping up


In this tutorial, you have learned how to create your own training pipeline for object detection models on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes and segmentation masks. You also leveraged a Mask R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-GPU training, check `references/detection/train.py`, which is present in the torchvision repository.

You can download a full source file for this tutorial [here](#).

```
image = read_image("drive/MyDrive/Beatles_-_Abbey_Road.jpg")  
eval_transform = get_transform(train=False)  
  
model.eval()  
with torch.no_grad():  
    x = eval_transform(image)  
    # convert RGBA -> RGB and move to device  
    x = x[:3, ...].to(device)  
    predictions = model([x, ])  
    pred = predictions[0]  
  
image = (255.0 * (image - image.min()) / (image.max() -  
image.min())).to(torch.uint8)  
image = image[:3, ...]  
pred_labels = [f"pedestrian: {score:.3f}" for label, score in  
zip(pred["labels"], pred["scores"])]  
pred_boxes = pred["boxes"].long()  
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,  
colors="red")  
  
masks = (pred["masks"] > 0.7).squeeze(1)  
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,  
colors="blue")  
  
plt.figure(figsize=(12, 12))  
plt.imshow(output_image.permute(1, 2, 0))
```


<matplotlib.image.AxesImage at 0x7ef4a2a48cd0>



pred_labels

```
['pedestrian: 0.987',  
'pedestrian: 0.985',  
'pedestrian: 0.957',  
'pedestrian: 0.939',  
'pedestrian: 0.103',  
'pedestrian: 0.093',  
'pedestrian: 0.089',
```

```
'pedestrian: 0.060',  
'pedestrian: 0.059']
```

pred_boxes

```
tensor([[ 21, 161,  84, 273],  
        [ 92, 164, 144, 272],  
        [148, 168, 219, 296],  
        [223, 163, 290, 293],  
        [ 41, 162, 237, 277],  
        [239, 147, 245, 174],  
        [114, 160, 277, 288],  
        [ 57, 159, 136, 269],  
        [104, 162, 190, 273]], device='cuda:0')
```

This is for Question 5 Part B and C using the backbone option

```
# For tips on running notebooks in Google Colab, see
# https://pytorch.org/tutorials/beginner/colab
%matplotlib inline

from google.colab import drive
drive.mount('/content/drive') #Updated the spots where they needed the
locations to change, also downloaded the PennFudan dataset and
tutorial source for the test photo at the end
# Dataset: https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
# Test Photo:
https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425fa
a00c06bcf502/_static/img/tv_tutorial/tv_image05.png
#Source Code:
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

Mounted at /content/drive
```

TorchVision Object Detection Finetuning Tutorial

.. tip::

To get the most of this tutorial, we suggest using this [Colab Version](https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_finetuning_instance_segmentation.ipynb). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model on the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection and instance segmentation model on a custom dataset.

.. note ::

This tutorial works only with torchvision version `>=0.16` or `nightly`. If you're using `torchvision<=0.15`, please follow [this tutorial instead](https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425faa00c06bcf502/intermediate_source/torchvision_tutorial.rst).

Defining the Dataset

The reference scripts for training object detection, instance segmentation and person keypoint detection allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return a tuple:

- `image`: `:class:torchvision.tv_tensors.Image` of shape `[3, H, W]`, a pure tensor, or a PIL Image of size `(H, W)`
- `target`: a dict containing the following fields
 - `boxes`, `:class:torchvision.tv_tensors.BoundingBoxes` of shape `[N, 4]`: the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels`, integer `:class:torch.Tensor` of shape `[N]`: the label for each bounding box. `0` represents always the background class.
 - `image_id`, int: an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area`, float `:class:torch.Tensor` of shape `[N]`: the area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd`, uint8 `:class:torch.Tensor` of shape `[N]`: instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks`, `:class:torchvision.tv_tensors.Mask` of shape `[N, H, W]`: the segmentation masks for each one of the objects

If your dataset is compliant with above requirements then it will work for both training and evaluation codes from the reference script. Evaluation code will use scripts from `pycocotools` which can be installed with `pip install pycocotools`.

.. note :: For Windows, please install `pycocotools` from [gautamchitnis_](https://github.com/gautamchitnis/cocoapi) with command

```
pip install
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-
master#subdirectory=PythonAPI
```

One note on the `labels`. The model considers class `0` as background. If your dataset does not contain the background class, you should not have `0` in your `labels`. For example, assuming you have just two classes, *cat* and *dog*, you can define `1` (not `0`) to represent *cats* and `2` to represent *dogs*. So, for instance, if one of the images has both classes, your `labels` tensor should look like `[1, 2]`.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratios), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

Writing a custom dataset for PennFudan

Let's write a dataset for the PennFudan dataset. After [downloading and extracting the zip file](#), we have the following folder structure:

::

```
PennFudanPed/ PedMasks/ FudanPed00001_mask.png FudanPed00002_mask.png
FudanPed00003_mask.png FudanPed00004_mask.png ... PNGImages/ FudanPed00001.png
FudanPed00002.png FudanPed00003.png FudanPed00004.png
```

Here is one example of a pair of images and segmentation masks

So each image has a corresponding segmentation mask, where each color correspond to a different instance. Let's write a `:class:torch.utils.data.Dataset` class for this dataset. In the code below, we are wrapping images, bounding boxes and masks into `torchvision.TVTensor` classes so that we will be able to apply torchvision built-in transformations ([new Transforms API](#)) for the given object detection and segmentation task. Namely, image tensors will be wrapped by `:class:torchvision.tv_tensors.Image`, bounding boxes into `:class:torchvision.tv_tensors.BoundingBoxes` and masks into `:class:torchvision.tv_tensors.Mask`. As `torchvision.TVTensor` are `:class:torch.Tensor` subclasses, wrapped objects are also tensors and inherit the plain `:class:torch.Tensor` API. For more information about torchvision `tv_tensors` see [this documentation](#).

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
```

```

        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

That's all for the dataset. Now let's define a model that can perform predictions on this dataset.

Defining your model

In this tutorial, we will be using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.

Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

There are two common situations where one might want to modify one of the available models in TorchVision Model Zoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

2 - Modifying the model to add a different backbone

```
import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator

# load a pre-trained model for classification and return
# only the features
backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features
# `FasterRCNN` needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(
    sizes=((32, 64, 128, 256, 512),),
    aspect_ratios=((0.5, 1.0, 2.0),)
)

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# `OrderedDict[Tensor]`, and in `featmap_names` you can choose
# which
```

```

# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIALign(
    featmap_names=['0'],
    output_size=7,
    sampling_ratio=2
)

# put the pieces together inside a Faster-RCNN model
model = FasterRCNN(
    backbone,
    num_classes=2,
    rpn_anchor_generator=anchor_generator,
    box_roi_pool=roi_pooler
)

Downloading: "https://download.pytorch.org/models/mobilenet_v2-7ebf99e0.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-7ebf99e0.pth
100%|██████████| 13.6M/13.6M [00:00<00:00, 80.2MB/s]

```

Object detection and instance segmentation model for PennFudan Dataset

In our case, we want to finetune from a pre-trained model, given that our dataset is very small, so we will be following approach number 1.

Here we want to also compute the instance segmentation masks, so we will be using Mask R-CNN:

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
    torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
    num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
    model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256

```

```

# and replace the mask predictor with a new one
model.roi_heads.mask_predictor = MaskRCNNPredictor(
    in_features_mask,
    hidden_layer,
    num_classes
)

return model

```

That's it, this will make `model` be ready to be trained and evaluated on your custom dataset.

Putting everything together

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py` and `references/detection/utils.py`. Just download everything under `references/detection` to your folder and use them here. On Linux if you have `wget`, you can download them using below commands:

```

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API
<https://pytorch.org/vision/stable/transforms.html>`
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#
# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:

```

```

        transforms.append(T.RandomHorizontalFlip(0.5))
        transforms.append(T.ToDtype(torch.float, scale=True))
        transforms.append(T.ToPureTensor())
        return T.Compose(transforms)

# Testing ``forward()`` method (Optional)
# -----
#
# Before iterating over the dataset, it's good to see what the model
# expects during training and inference time on sample data.
import utils

model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
print(output)

# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions
print(predictions[0])

```

Downloading:

"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth" to
 /root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth

100%|██████████| 160M/160M [00:00<00:00, 171MB/s]

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py
 :557: UserWarning: This DataLoader will create 4 worker processes in
 total. Our suggested max number of worker in current system is 2,
 which is smaller than what this DataLoader is going to create. Please

be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn(_create_warning_msg(

{'loss_classifier': tensor(0.2291, grad_fn=<NllLossBackward0>),
'loss_box_reg': tensor(0.1100, grad_fn=<DivBackward0>),
'loss_objectness': tensor(0.0078,
grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), 'loss_rpn_box_reg':
tensor(0.0105, grad_fn=<DivBackward0>)}
{'boxes': tensor([], size=(0, 4), grad_fn=<StackBackward0>), 'labels':
tensor([], dtype=torch.int64), 'scores': tensor([],
grad_fn=<IndexBackward0>)})
```

Let's now write the main function which performs the training and the validation:

```
from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
dataset_test = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
```



```

)

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
    print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

Downloading:

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-bf2d0c1e.pth" to
/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-bf2d0c1e.pth

100%|██████████| 170M/170M [00:01<00:00, 154MB/s]

Epoch: [0] [0/60] eta: 0:09:12 lr: 0.000090 loss: 5.0986 (5.0986)
loss_classifier: 0.5976 (0.5976) loss_box_reg: 0.1450 (0.1450)
loss_mask: 4.3445 (4.3445) loss_objectness: 0.0105 (0.0105)
loss_rpn_box_reg: 0.0011 (0.0011) time: 9.2085 data: 1.7603 max
mem: 2065

```
Epoch: [0] [10/60] eta: 0:01:05 lr: 0.000936 loss: 1.7777 (2.5739)
loss_classifier: 0.4427 (0.4398) loss_box_reg: 0.2589 (0.2873)
loss_mask: 1.1316 (1.8147) loss_objectness: 0.0235 (0.0256)
loss_rpn_box_reg: 0.0042 (0.0064) time: 1.3122 data: 0.1694 max
mem: 2764
Epoch: [0] [20/60] eta: 0:00:37 lr: 0.001783 loss: 0.9253 (1.6445)
loss_classifier: 0.1694 (0.2943) loss_box_reg: 0.2363 (0.2486)
loss_mask: 0.4028 (1.0708) loss_objectness: 0.0174 (0.0243)
loss_rpn_box_reg: 0.0060 (0.0064) time: 0.5337 data: 0.0097 max
mem: 3211
Epoch: [0] [30/60] eta: 0:00:24 lr: 0.002629 loss: 0.5503 (1.2777)
loss_classifier: 0.0942 (0.2279) loss_box_reg: 0.2240 (0.2401)
loss_mask: 0.1915 (0.7836) loss_objectness: 0.0093 (0.0189)
loss_rpn_box_reg: 0.0062 (0.0071) time: 0.5428 data: 0.0086 max
mem: 3211
Epoch: [0] [40/60] eta: 0:00:15 lr: 0.003476 loss: 0.4641 (1.0915)
loss_classifier: 0.0733 (0.1892) loss_box_reg: 0.2205 (0.2392)
loss_mask: 0.1874 (0.6412) loss_objectness: 0.0037 (0.0150)
loss_rpn_box_reg: 0.0075 (0.0069) time: 0.5498 data: 0.0110 max
mem: 3211
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.4495 (0.9692)
loss_classifier: 0.0539 (0.1633) loss_box_reg: 0.2102 (0.2342)
loss_mask: 0.1874 (0.5518) loss_objectness: 0.0021 (0.0130)
loss_rpn_box_reg: 0.0058 (0.0069) time: 0.5793 data: 0.0148 max
mem: 3222
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3703 (0.8745)
loss_classifier: 0.0415 (0.1442) loss_box_reg: 0.1577 (0.2203)
loss_mask: 0.1652 (0.4921) loss_objectness: 0.0017 (0.0112)
loss_rpn_box_reg: 0.0058 (0.0067) time: 0.5651 data: 0.0123 max
mem: 3222
Epoch: [0] Total time: 0:00:41 (0.6973 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.2362 (0.2362)
evaluator_time: 0.0045 (0.0045) time: 0.5697 data: 0.3278 max mem:
3222
Test: [49/50] eta: 0:00:00 model_time: 0.0982 (0.1290)
evaluator_time: 0.0053 (0.0121) time: 0.1231 data: 0.0049 max mem:
3222
Test: Total time: 0:00:07 (0.1594 s / it)
Averaged stats: model_time: 0.0982 (0.1290) evaluator_time: 0.0053
(0.0121)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.584
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.676
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.267
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.596
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.596
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.266
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.667
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.667
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.668
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.664
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.961
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.798
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.294
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.432
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.688
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.300
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.722
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.723
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.692
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.732
Epoch: [1] [ 0/60] eta: 0:00:52  lr: 0.005000  loss: 0.2154 (0.2154)
loss_classifier: 0.0134 (0.0134)  loss_box_reg: 0.0351 (0.0351)

```

```
loss_mask: 0.1655 (0.1655) loss_objectness: 0.0000 (0.0000)
loss_rpn_box_reg: 0.0014 (0.0014) time: 0.8740 data: 0.3722 max
mem: 3222
Epoch: [1] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2632 (0.3034)
loss_classifier: 0.0393 (0.0371) loss_box_reg: 0.0935 (0.1025)
loss_mask: 0.1525 (0.1554) loss_objectness: 0.0010 (0.0015)
loss_rpn_box_reg: 0.0036 (0.0068) time: 0.5839 data: 0.0398 max
mem: 3222
Epoch: [1] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2812 (0.3210)
loss_classifier: 0.0393 (0.0425) loss_box_reg: 0.1028 (0.1159)
loss_mask: 0.1403 (0.1537) loss_objectness: 0.0010 (0.0017)
loss_rpn_box_reg: 0.0061 (0.0072) time: 0.5770 data: 0.0092 max
mem: 3222
Epoch: [1] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2619 (0.2951)
loss_classifier: 0.0371 (0.0390) loss_box_reg: 0.1001 (0.1037)
loss_mask: 0.1330 (0.1445) loss_objectness: 0.0008 (0.0017)
loss_rpn_box_reg: 0.0057 (0.0061) time: 0.5975 data: 0.0097 max
mem: 3499
Epoch: [1] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2293 (0.2814)
loss_classifier: 0.0254 (0.0362) loss_box_reg: 0.0697 (0.0951)
loss_mask: 0.1265 (0.1430) loss_objectness: 0.0004 (0.0014)
loss_rpn_box_reg: 0.0037 (0.0057) time: 0.5854 data: 0.0087 max
mem: 3499
Epoch: [1] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2533 (0.2769)
loss_classifier: 0.0271 (0.0350) loss_box_reg: 0.0671 (0.0901)
loss_mask: 0.1385 (0.1447) loss_objectness: 0.0004 (0.0017)
loss_rpn_box_reg: 0.0040 (0.0055) time: 0.5632 data: 0.0091 max
mem: 3499
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2533 (0.2748)
loss_classifier: 0.0292 (0.0352) loss_box_reg: 0.0677 (0.0869)
loss_mask: 0.1429 (0.1455) loss_objectness: 0.0004 (0.0015)
loss_rpn_box_reg: 0.0044 (0.0058) time: 0.5583 data: 0.0081 max
mem: 3499
Epoch: [1] Total time: 0:00:35 (0.5834 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:36 model_time: 0.2662 (0.2662)
evaluator_time: 0.0041 (0.0041) time: 0.7247 data: 0.4532 max mem:
3499
Test: [49/50] eta: 0:00:00 model_time: 0.0967 (0.1095)
evaluator_time: 0.0038 (0.0059) time: 0.1158 data: 0.0042 max mem:
3499
Test: Total time: 0:00:06 (0.1344 s / it)
Averaged stats: model_time: 0.0967 (0.1095) evaluator_time: 0.0038
(0.0059)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
```

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.738

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.969

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.932

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.381

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.688

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.752

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.330

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.796

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.796

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.500

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.775

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.806

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.692

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.974

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.826

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.329

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.619

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.706

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.311

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.748

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.748

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.533

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.750

Average Recall (AR) @[IoU=0.50:0.95 | area= large |

```
maxDets=100 ] = 0.753
Epoch: [2] [ 0/60] eta: 0:00:53 lr: 0.005000 loss: 0.1752 (0.1752)
loss_classifier: 0.0203 (0.0203) loss_box_reg: 0.0485 (0.0485)
loss_mask: 0.1022 (0.1022) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0038 (0.0038) time: 0.8881 data: 0.3680 max
mem: 3499
Epoch: [2] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2351 (0.2578)
loss_classifier: 0.0322 (0.0362) loss_box_reg: 0.0631 (0.0793)
loss_mask: 0.1351 (0.1367) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0038 (0.0048) time: 0.5899 data: 0.0413 max
mem: 3499
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2351 (0.2564)
loss_classifier: 0.0303 (0.0325) loss_box_reg: 0.0725 (0.0793)
loss_mask: 0.1351 (0.1394) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0036 (0.0046) time: 0.5669 data: 0.0089 max
mem: 3499
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2122 (0.2387)
loss_classifier: 0.0246 (0.0306) loss_box_reg: 0.0517 (0.0711)
loss_mask: 0.1209 (0.1320) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0043) time: 0.5666 data: 0.0091 max
mem: 3499
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2139 (0.2401)
loss_classifier: 0.0334 (0.0315) loss_box_reg: 0.0517 (0.0713)
loss_mask: 0.1146 (0.1319) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0031 (0.0046) time: 0.5879 data: 0.0099 max
mem: 3506
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2241 (0.2373)
loss_classifier: 0.0339 (0.0311) loss_box_reg: 0.0579 (0.0691)
loss_mask: 0.1267 (0.1320) loss_objectness: 0.0007 (0.0007)
loss_rpn_box_reg: 0.0034 (0.0044) time: 0.5892 data: 0.0091 max
mem: 3506
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2103 (0.2388)
loss_classifier: 0.0257 (0.0310) loss_box_reg: 0.0546 (0.0685)
loss_mask: 0.1331 (0.1340) loss_objectness: 0.0005 (0.0008)
loss_rpn_box_reg: 0.0029 (0.0045) time: 0.5629 data: 0.0074 max
mem: 3506
Epoch: [2] Total time: 0:00:34 (0.5829 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1925 (0.1925)
evaluator_time: 0.0037 (0.0037) time: 0.5428 data: 0.3454 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0969 (0.1068)
evaluator_time: 0.0032 (0.0049) time: 0.1137 data: 0.0037 max mem:
3506
Test: Total time: 0:00:06 (0.1284 s / it)
Averaged stats: model_time: 0.0969 (0.1068) evaluator_time: 0.0032
(0.0049)
Accumulating evaluation results...
```

```

DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.764
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.905
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.292
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.696
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.786
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.339
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.820
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.820
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.433
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.758
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.837
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.748
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.973
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.938
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.315
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.567
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.326
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.433

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.758
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.805
Epoch: [3] [ 0/60] eta: 0:00:55 lr: 0.000500 loss: 0.1499 (0.1499)
loss_classifier: 0.0188 (0.0188) loss_box_reg: 0.0275 (0.0275)
loss_mask: 0.0995 (0.0995) loss_objectness: 0.0013 (0.0013)
loss_rpn_box_reg: 0.0028 (0.0028) time: 0.9178 data: 0.3975 max
mem: 3506
Epoch: [3] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1933 (0.1984)
loss_classifier: 0.0250 (0.0240) loss_box_reg: 0.0381 (0.0457)
loss_mask: 0.1201 (0.1236) loss_objectness: 0.0010 (0.0020)
loss_rpn_box_reg: 0.0023 (0.0030) time: 0.6011 data: 0.0473 max
mem: 3506
Epoch: [3] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1933 (0.1941)
loss_classifier: 0.0217 (0.0244) loss_box_reg: 0.0381 (0.0418)
loss_mask: 0.1204 (0.1229) loss_objectness: 0.0007 (0.0018)
loss_rpn_box_reg: 0.0023 (0.0033) time: 0.5586 data: 0.0100 max
mem: 3506
Epoch: [3] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1934 (0.1968)
loss_classifier: 0.0248 (0.0255) loss_box_reg: 0.0394 (0.0434)
loss_mask: 0.1204 (0.1232) loss_objectness: 0.0004 (0.0016)
loss_rpn_box_reg: 0.0032 (0.0032) time: 0.5505 data: 0.0101 max
mem: 3506
Epoch: [3] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1934 (0.1993)
loss_classifier: 0.0257 (0.0261) loss_box_reg: 0.0436 (0.0445)
loss_mask: 0.1213 (0.1236) loss_objectness: 0.0006 (0.0017)
loss_rpn_box_reg: 0.0034 (0.0034) time: 0.5638 data: 0.0108 max
mem: 3506
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1763 (0.1969)
loss_classifier: 0.0241 (0.0258) loss_box_reg: 0.0361 (0.0436)
loss_mask: 0.1137 (0.1223) loss_objectness: 0.0007 (0.0019)
loss_rpn_box_reg: 0.0025 (0.0032) time: 0.5697 data: 0.0088 max
mem: 3506
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1844 (0.1977)
loss_classifier: 0.0217 (0.0261) loss_box_reg: 0.0374 (0.0444)
loss_mask: 0.1123 (0.1220) loss_objectness: 0.0004 (0.0017)
loss_rpn_box_reg: 0.0027 (0.0034) time: 0.5720 data: 0.0080 max
mem: 3506
Epoch: [3] Total time: 0:00:34 (0.5757 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.2132 (0.2132)
evaluator_time: 0.0038 (0.0038) time: 0.5280 data: 0.3096 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.1036 (0.1098)
evaluator_time: 0.0050 (0.0060) time: 0.1220 data: 0.0047 max mem:
3506
Test: Total time: 0:00:06 (0.1367 s / it)
```


Averaged stats: model_time: 0.1036 (0.1098) evaluator_time: 0.0050 (0.0060)

Accumulating evaluation results...

DONE (t=0.03s).

Accumulating evaluation results...

DONE (t=0.03s).

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.791

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.977

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.923

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.372

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.703

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.812

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.348

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.850

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.850

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.567

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.775

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.866

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.754

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.977

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.937

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.368

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.523

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.773

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.331

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.803

Average Recall (AR) @[IoU=0.50:0.95 | area= all |

```

maxDets=100 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.600
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
Epoch: [4] [ 0/60] eta: 0:01:20 lr: 0.000500 loss: 0.2263 (0.2263)
loss_classifier: 0.0418 (0.0418) loss_box_reg: 0.0667 (0.0667)
loss_mask: 0.1081 (0.1081) loss_objectness: 0.0019 (0.0019)
loss_rpn_box_reg: 0.0078 (0.0078) time: 1.3454 data: 0.6900 max
mem: 3506
Epoch: [4] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.1923 (0.1878)
loss_classifier: 0.0294 (0.0274) loss_box_reg: 0.0394 (0.0393)
loss_mask: 0.1081 (0.1172) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0016 (0.0031) time: 0.6336 data: 0.0698 max
mem: 3506
Epoch: [4] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.1878 (0.1919)
loss_classifier: 0.0271 (0.0274) loss_box_reg: 0.0394 (0.0438)
loss_mask: 0.1124 (0.1166) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0022 (0.0034) time: 0.5651 data: 0.0086 max
mem: 3506
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1777 (0.1855)
loss_classifier: 0.0248 (0.0265) loss_box_reg: 0.0356 (0.0394)
loss_mask: 0.1126 (0.1156) loss_objectness: 0.0002 (0.0009)
loss_rpn_box_reg: 0.0017 (0.0030) time: 0.5693 data: 0.0113 max
mem: 3506
Epoch: [4] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1566 (0.1846)
loss_classifier: 0.0217 (0.0255) loss_box_reg: 0.0259 (0.0393)
loss_mask: 0.1126 (0.1161) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0014 (0.0028) time: 0.5630 data: 0.0128 max
mem: 3506
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1653 (0.1845)
loss_classifier: 0.0224 (0.0254) loss_box_reg: 0.0278 (0.0396)
loss_mask: 0.1163 (0.1159) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5632 data: 0.0133 max
mem: 3506
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1707 (0.1862)
loss_classifier: 0.0241 (0.0255) loss_box_reg: 0.0333 (0.0407)
loss_mask: 0.1126 (0.1164) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0029 (0.0030) time: 0.5827 data: 0.0109 max
mem: 3506
Epoch: [4] Total time: 0:00:35 (0.5868 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.1612 (0.1612)
evaluator_time: 0.0042 (0.0042) time: 0.5203 data: 0.3537 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0973 (0.1117)

```

evaluator_time: 0.0033 (0.0072) time: 0.1171 data: 0.0053 max mem: 3506

Test: Total time: 0:00:06 (0.1387 s / it)

Averaged stats: model_time: 0.0973 (0.1117) evaluator_time: 0.0033 (0.0072)

Accumulating evaluation results...

DONE (t=0.02s).

Accumulating evaluation results...

DONE (t=0.02s).

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.802

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.975

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.913

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.335

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.700

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.825

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.354

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.849

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.849

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.467

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.767

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.868

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.758

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.975

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.945

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.313

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.527

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.780

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.334

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=

```
10 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [5] [ 0/60] eta: 0:01:01  lr: 0.000500  loss: 0.1311 (0.1311)
loss_classifier: 0.0145 (0.0145)  loss_box_reg: 0.0197 (0.0197)
loss_mask: 0.0942 (0.0942)  loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0024 (0.0024)  time: 1.0227  data: 0.4592  max
mem: 3506
Epoch: [5] [10/60] eta: 0:00:27  lr: 0.000500  loss: 0.1704 (0.1697)
loss_classifier: 0.0172 (0.0197)  loss_box_reg: 0.0356 (0.0336)
loss_mask: 0.1118 (0.1133)  loss_objectness: 0.0006 (0.0010)
loss_rpn_box_reg: 0.0024 (0.0021)  time: 0.5558  data: 0.0470  max
mem: 3506
Epoch: [5] [20/60] eta: 0:00:22  lr: 0.000500  loss: 0.1704 (0.1813)
loss_classifier: 0.0216 (0.0224)  loss_box_reg: 0.0365 (0.0369)
loss_mask: 0.1150 (0.1184)  loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0026 (0.0028)  time: 0.5431  data: 0.0072  max
mem: 3506
Epoch: [5] [30/60] eta: 0:00:17  lr: 0.000500  loss: 0.1792 (0.1873)
loss_classifier: 0.0254 (0.0234)  loss_box_reg: 0.0371 (0.0392)
loss_mask: 0.1242 (0.1211)  loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0028)  time: 0.5843  data: 0.0089  max
mem: 3506
Epoch: [5] [40/60] eta: 0:00:11  lr: 0.000500  loss: 0.1743 (0.1817)
loss_classifier: 0.0224 (0.0225)  loss_box_reg: 0.0323 (0.0373)
loss_mask: 0.1098 (0.1183)  loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0017 (0.0027)  time: 0.5693  data: 0.0092  max
mem: 3506
Epoch: [5] [50/60] eta: 0:00:05  lr: 0.000500  loss: 0.1669 (0.1854)
loss_classifier: 0.0195 (0.0239)  loss_box_reg: 0.0334 (0.0392)
loss_mask: 0.1054 (0.1183)  loss_objectness: 0.0006 (0.0011)
loss_rpn_box_reg: 0.0024 (0.0029)  time: 0.5706  data: 0.0094  max
mem: 3506
Epoch: [5] [59/60] eta: 0:00:00  lr: 0.000500  loss: 0.1669 (0.1838)
loss_classifier: 0.0234 (0.0236)  loss_box_reg: 0.0390 (0.0390)
loss_mask: 0.1064 (0.1172)  loss_objectness: 0.0003 (0.0011)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5905  data: 0.0086  max
mem: 3506
Epoch: [5] Total time: 0:00:34 (0.5787 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:32  model_time: 0.1621 (0.1621)
evaluator_time: 0.0040 (0.0040)  time: 0.6579  data: 0.4907  max mem:
```

```
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0967 (0.1065)
evaluator_time: 0.0041 (0.0050) time: 0.1136 data: 0.0037 max mem:
3506
Test: Total time: 0:00:06 (0.1310 s / it)
Averaged stats: model_time: 0.0967 (0.1065) evaluator_time: 0.0041
(0.0050)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.809
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.973
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.913
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.706
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.831
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.854
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.854
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.760
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.313
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.562
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.781
```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.334
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [6] [ 0/60] eta: 0:00:59  lr: 0.000050  loss: 0.1643 (0.1643)
loss_classifier: 0.0163 (0.0163)  loss_box_reg: 0.0271 (0.0271)
loss_mask: 0.1181 (0.1181)  loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0027 (0.0027)  time: 0.9919  data: 0.4327  max
mem: 3506
Epoch: [6] [10/60] eta: 0:00:29  lr: 0.000050  loss: 0.1693 (0.1709)
loss_classifier: 0.0187 (0.0200)  loss_box_reg: 0.0339 (0.0332)
loss_mask: 0.1108 (0.1150)  loss_objectness: 0.0003 (0.0004)
loss_rpn_box_reg: 0.0023 (0.0024)  time: 0.5922  data: 0.0479  max
mem: 3506
Epoch: [6] [20/60] eta: 0:00:23  lr: 0.000050  loss: 0.1793 (0.1807)
loss_classifier: 0.0242 (0.0241)  loss_box_reg: 0.0339 (0.0363)
loss_mask: 0.1108 (0.1167)  loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0023 (0.0028)  time: 0.5668  data: 0.0090  max
mem: 3506
Epoch: [6] [30/60] eta: 0:00:17  lr: 0.000050  loss: 0.1793 (0.1808)
loss_classifier: 0.0242 (0.0245)  loss_box_reg: 0.0327 (0.0368)
loss_mask: 0.1136 (0.1161)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0025 (0.0028)  time: 0.5922  data: 0.0084  max
mem: 3506
Epoch: [6] [40/60] eta: 0:00:11  lr: 0.000050  loss: 0.1654 (0.1841)
loss_classifier: 0.0206 (0.0245)  loss_box_reg: 0.0300 (0.0384)
loss_mask: 0.1133 (0.1173)  loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0025 (0.0030)  time: 0.5777  data: 0.0092  max
mem: 3506
Epoch: [6] [50/60] eta: 0:00:05  lr: 0.000050  loss: 0.1654 (0.1815)
loss_classifier: 0.0191 (0.0243)  loss_box_reg: 0.0300 (0.0377)
loss_mask: 0.1091 (0.1157)  loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0022 (0.0029)  time: 0.5662  data: 0.0090  max
mem: 3780
Epoch: [6] [59/60] eta: 0:00:00  lr: 0.000050  loss: 0.1691 (0.1797)
loss_classifier: 0.0206 (0.0238)  loss_box_reg: 0.0285 (0.0370)
loss_mask: 0.1148 (0.1152)  loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0020 (0.0028)  time: 0.5693  data: 0.0076  max
mem: 3780
Epoch: [6] Total time: 0:00:35 (0.5842 s / it)
creating index...
```

```
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1869 (0.1869)
evaluator_time: 0.0060 (0.0060) time: 0.5776 data: 0.3833 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.0973 (0.1077)
evaluator_time: 0.0033 (0.0052) time: 0.1150 data: 0.0036 max mem:
3780
Test: Total time: 0:00:06 (0.1305 s / it)
Averaged stats: model_time: 0.0973 (0.1077) evaluator_time: 0.0033
(0.0052)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.808
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.704
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.831
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.355
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.873
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.759
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
```

```
maxDets=100 ] = 0.528
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.334
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.814
Epoch: [7] [ 0/60] eta: 0:01:28 lr: 0.000050 loss: 0.1616 (0.1616)
loss_classifier: 0.0205 (0.0205) loss_box_reg: 0.0288 (0.0288)
loss_mask: 0.1102 (0.1102) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0021 (0.0021) time: 1.4667 data: 0.6278 max
mem: 3780
Epoch: [7] [10/60] eta: 0:00:32 lr: 0.000050 loss: 0.1826 (0.1937)
loss_classifier: 0.0241 (0.0265) loss_box_reg: 0.0363 (0.0411)
loss_mask: 0.1177 (0.1229) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0021 (0.0025) time: 0.6578 data: 0.0641 max
mem: 3780
Epoch: [7] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1834 (0.1907)
loss_classifier: 0.0241 (0.0258) loss_box_reg: 0.0402 (0.0413)
loss_mask: 0.1172 (0.1201) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0026 (0.0029) time: 0.5638 data: 0.0084 max
mem: 3780
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1867 (0.1909)
loss_classifier: 0.0240 (0.0260) loss_box_reg: 0.0443 (0.0416)
loss_mask: 0.1058 (0.1197) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0030) time: 0.5671 data: 0.0102 max
mem: 3780
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1575 (0.1866)
loss_classifier: 0.0200 (0.0247) loss_box_reg: 0.0340 (0.0401)
loss_mask: 0.1044 (0.1184) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5652 data: 0.0095 max
mem: 3780
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1498 (0.1816)
loss_classifier: 0.0164 (0.0240) loss_box_reg: 0.0250 (0.0378)
loss_mask: 0.1044 (0.1164) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0017 (0.0027) time: 0.5425 data: 0.0088 max
mem: 3780
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1522 (0.1790)
loss_classifier: 0.0212 (0.0239) loss_box_reg: 0.0261 (0.0367)
loss_mask: 0.1036 (0.1150) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0024 (0.0028) time: 0.5749 data: 0.0091 max
```



```
mem: 3780
Epoch: [7] Total time: 0:00:35 (0.5869 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1866 (0.1866)
evaluator_time: 0.0039 (0.0039) time: 0.5075 data: 0.3155 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.1186 (0.1148)
evaluator_time: 0.0055 (0.0068) time: 0.1294 data: 0.0048 max mem:
3780
Test: Total time: 0:00:06 (0.1398 s / it)
Averaged stats: model_time: 0.1186 (0.1148) evaluator_time: 0.0055
(0.0068)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.810
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.834
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.856
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.856
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.876
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.531
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [8] [ 0/60] eta: 0:01:06 lr: 0.000050 loss: 0.1370 (0.1370)
loss_classifier: 0.0163 (0.0163) loss_box_reg: 0.0187 (0.0187)
loss_mask: 0.0959 (0.0959) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0057 (0.0057) time: 1.1106 data: 0.4501 max
mem: 3780
Epoch: [8] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1865 (0.2064)
loss_classifier: 0.0221 (0.0261) loss_box_reg: 0.0353 (0.0489)
loss_mask: 0.1295 (0.1263) loss_objectness: 0.0002 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0044) time: 0.5999 data: 0.0469 max
mem: 3780
Epoch: [8] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1865 (0.1967)
loss_classifier: 0.0244 (0.0263) loss_box_reg: 0.0408 (0.0450)
loss_mask: 0.1164 (0.1205) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0030 (0.0043) time: 0.5597 data: 0.0091 max
mem: 3780
Epoch: [8] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1610 (0.1818)
loss_classifier: 0.0210 (0.0243) loss_box_reg: 0.0292 (0.0389)
loss_mask: 0.0997 (0.1145) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0026 (0.0036) time: 0.5602 data: 0.0101 max
mem: 3780
Epoch: [8] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1455 (0.1775)
loss_classifier: 0.0183 (0.0239) loss_box_reg: 0.0238 (0.0373)
loss_mask: 0.0973 (0.1124) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0021 (0.0032) time: 0.5544 data: 0.0081 max
mem: 3780
Epoch: [8] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1653 (0.1799)
loss_classifier: 0.0229 (0.0248) loss_box_reg: 0.0294 (0.0371)
loss_mask: 0.1098 (0.1143) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0030) time: 0.5777 data: 0.0088 max
mem: 3780
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1663 (0.1786)

```

```
loss_classifier: 0.0229 (0.0245) loss_box_reg: 0.0291 (0.0367)
loss_mask: 0.1133 (0.1139) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0029) time: 0.5682 data: 0.0086 max
mem: 3780
Epoch: [8] Total time: 0:00:34 (0.5768 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:42 model_time: 0.3319 (0.3319)
evaluator_time: 0.0050 (0.0050) time: 0.8421 data: 0.5039 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.0978 (0.1148)
evaluator_time: 0.0033 (0.0058) time: 0.1142 data: 0.0036 max mem:
3780
Test: Total time: 0:00:07 (0.1426 s / it)
Averaged stats: model_time: 0.0978 (0.1148) evaluator_time: 0.0033
(0.0058)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.806
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.830
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.354
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.853
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.853
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
```

```

maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.529
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.816
Epoch: [9] [ 0/60] eta: 0:00:53 lr: 0.000005 loss: 0.1116 (0.1116)
loss_classifier: 0.0113 (0.0113) loss_box_reg: 0.0150 (0.0150)
loss_mask: 0.0852 (0.0852) loss_objectness: 0.0000 (0.0000)
loss_rpn_box_reg: 0.0002 (0.0002) time: 0.8908 data: 0.3519 max
mem: 3780
Epoch: [9] [10/60] eta: 0:00:29 lr: 0.000005 loss: 0.1478 (0.1777)
loss_classifier: 0.0200 (0.0243) loss_box_reg: 0.0245 (0.0380)
loss_mask: 0.0980 (0.1113) loss_objectness: 0.0003 (0.0013)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5810 data: 0.0399 max
mem: 3780
Epoch: [9] [20/60] eta: 0:00:23 lr: 0.000005 loss: 0.1664 (0.1761)
loss_classifier: 0.0231 (0.0243) loss_box_reg: 0.0284 (0.0357)
loss_mask: 0.1047 (0.1119) loss_objectness: 0.0003 (0.0017)
loss_rpn_box_reg: 0.0020 (0.0025) time: 0.5605 data: 0.0095 max
mem: 3780
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1671 (0.1812)
loss_classifier: 0.0244 (0.0245) loss_box_reg: 0.0301 (0.0375)
loss_mask: 0.1119 (0.1150) loss_objectness: 0.0005 (0.0014)
loss_rpn_box_reg: 0.0025 (0.0029) time: 0.5652 data: 0.0098 max
mem: 3780
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1601 (0.1795)
loss_classifier: 0.0228 (0.0243) loss_box_reg: 0.0293 (0.0375)
loss_mask: 0.1107 (0.1135) loss_objectness: 0.0005 (0.0013)
loss_rpn_box_reg: 0.0029 (0.0029) time: 0.5756 data: 0.0094 max
mem: 3780
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1601 (0.1808)
loss_classifier: 0.0220 (0.0245) loss_box_reg: 0.0309 (0.0381)
loss_mask: 0.1089 (0.1141) loss_objectness: 0.0003 (0.0011)

```

```
loss_rpn_box_reg: 0.0029 (0.0029)  time: 0.5913  data: 0.0095  max
mem: 3780
Epoch: [9]  [59/60]  eta: 0:00:00  lr: 0.000005  loss: 0.1621 (0.1808)
loss_classifier: 0.0219 (0.0247)  loss_box_reg: 0.0326 (0.0377)
loss_mask: 0.1089 (0.1143)  loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0033 (0.0029)  time: 0.5792  data: 0.0088  max
mem: 3780
Epoch: [9] Total time: 0:00:34 (0.5811 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:40  model_time: 0.3076 (0.3076)
evaluator_time: 0.0055 (0.0055)  time: 0.8024  data: 0.4880  max mem:
3780
Test:  [49/50]  eta: 0:00:00  model_time: 0.1000 (0.1120)
evaluator_time: 0.0034 (0.0052)  time: 0.1157  data: 0.0037  max mem:
3780
Test: Total time: 0:00:06 (0.1375 s / it)
Averaged stats: model_time: 0.1000 (0.1120)  evaluator_time: 0.0034
(0.0052)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.806
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.975
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.919
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.830
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.354
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.853
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.853
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.313
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.529
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.816
That's it!

```

So after one epoch of training, we obtain a COCO-style mAP > 50, and a mask mAP of 65.

But what do the predictions look like? Let's take one image in the dataset and verify

```

import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

image =
read_image("drive/MyDrive/_static/img/tv_tutorial/tv_image05.png")
eval_transform = get_transform(train=False)

model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

```


```

image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))
<matplotlib.image.AxesImage at 0x7c62d689c3d0>

```



```
pred_labels
```

```
['pedestrian: 0.996',  
'pedestrian: 0.996',  
'pedestrian: 0.994',  
'pedestrian: 0.987',  
'pedestrian: 0.976',  
'pedestrian: 0.937',  
'pedestrian: 0.203',  
'pedestrian: 0.181',  
'pedestrian: 0.119',  
'pedestrian: 0.057']
```

```
pred_boxes
```

```
tensor([[170, 129, 268, 410],  
        [297, 117, 369, 340],  
        [445, 101, 508, 314],  
        [515, 118, 565, 255],  
        [361, 120, 440, 325],  
        [239, 125, 308, 360],  
        [268, 126, 507, 329],  
        [486, 105, 517, 172],  
        [541, 110, 564, 250],  
        [284, 109, 317, 169]] , device='cuda:0')
```

The results look good!

Wrapping up

In this tutorial, you have learned how to create your own training pipeline for object detection models on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes and segmentation masks. You also leveraged a Mask R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-GPU training, check `references/detection/train.py`, which is present in the torchvision repository.

You can download a full source file for this tutorial [here](#).

```
image = read_image("drive/MyDrive/Beatles_-_Abbey_Road.jpg")  
eval_transform = get_transform(train=False)  
  
model.eval()  
with torch.no_grad():  
    x = eval_transform(image)  
    # convert RGBA -> RGB and move to device  
    x = x[:3, ...].to(device)  
    predictions = model([x, ])
```



```
pred = predictions[0]

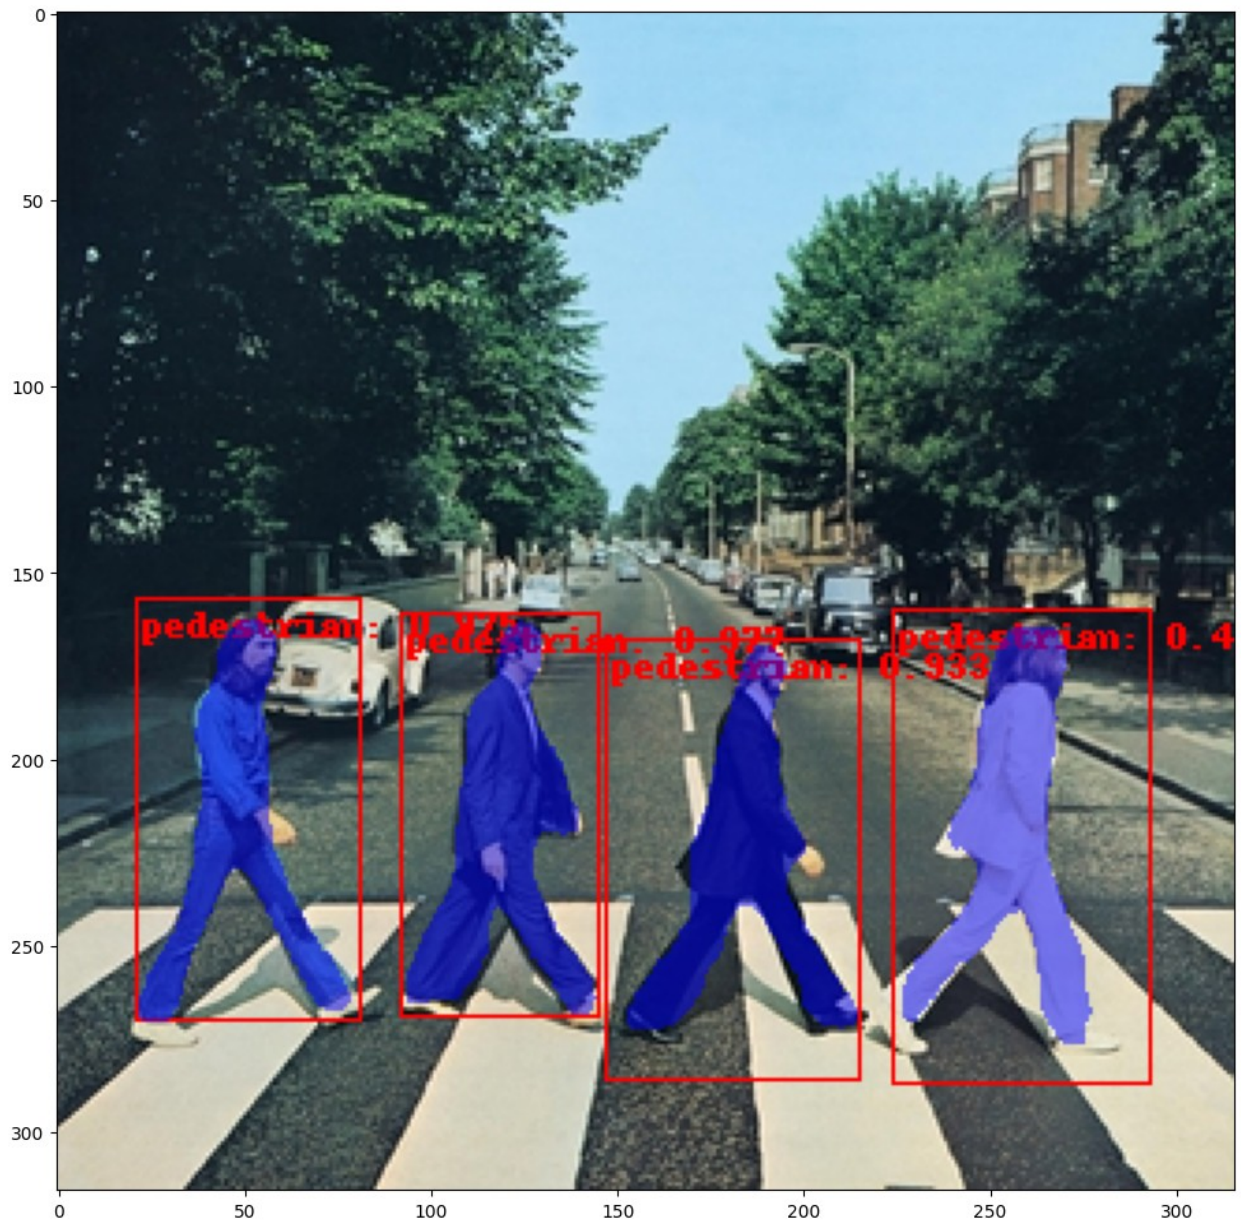
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red", font_size=2)

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

/usr/local/lib/python3.10/dist-packages/torchvision/utils.py:223:
UserWarning: Argument 'font_size' will be ignored since 'font' is not
set.
  warnings.warn("Argument 'font_size' will be ignored since 'font' is
not set.")

<matplotlib.image.AxesImage at 0x7c62d68430a0>
```



pred_labels

```
['pedestrian: 0.977',  
'pedestrian: 0.975',  
'pedestrian: 0.933',  
'pedestrian: 0.496']
```

pred_boxes

```
tensor([[ 92, 161, 145, 269],  
        [ 21, 157,  81, 270],  
        [147, 168, 215, 286],  
        [224, 160, 293, 287]], device='cuda:0')
```

