

This is for Question 5 Part B and C using the backbone option

```
# For tips on running notebooks in Google Colab, see
# https://pytorch.org/tutorials/beginner/colab
%matplotlib inline

from google.colab import drive
drive.mount('/content/drive') #Updated the spots where they needed the
locations to change, also downloaded the PennFudan dataset and
tutorial source for the test photo at the end
# Dataset: https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
# Test Photo:
https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425fa
a00c06bcf502/_static/img/tv_tutorial/tv_image05.png
#Source Code:
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

Mounted at /content/drive
```

TorchVision Object Detection Finetuning Tutorial

.. tip::

To get the most of this tutorial, we suggest using this [Colab Version](https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_finetuning_instance_segmentation.ipynb). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model on the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection and instance segmentation model on a custom dataset.

.. note ::

This tutorial works only with torchvision version `>=0.16` or `nightly`. If you're using `torchvision<=0.15`, please follow [this tutorial instead](https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425faa00c06bcf502/intermediate_source/torchvision_tutorial.rst).

Defining the Dataset

The reference scripts for training object detection, instance segmentation and person keypoint detection allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return a tuple:

- `image`: :class:torchvision.tv_tensors.Image of shape `[3, H, W]`, a pure tensor, or a PIL Image of size `(H, W)`
- `target`: a dict containing the following fields
 - `boxes`, :class:torchvision.tv_tensors.BoundingBoxes of shape `[N, 4]`: the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels`, integer :class:torch.Tensor of shape `[N]`: the label for each bounding box. `0` represents always the background class.
 - `image_id`, int: an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area`, float :class:torch.Tensor of shape `[N]`: the area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd`, uint8 :class:torch.Tensor of shape `[N]`: instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks`, :class:torchvision.tv_tensors.Mask of shape `[N, H, W]`: the segmentation masks for each one of the objects

If your dataset is compliant with above requirements then it will work for both training and evaluation codes from the reference script. Evaluation code will use scripts from `pycocotools` which can be installed with `pip install pycocotools`.

.. note :: For Windows, please install `pycocotools` from [gautamchitnis_](https://github.com/gautamchitnis/cocoapi) with command

```
pip install
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-
master#subdirectory=PythonAPI
```

One note on the `labels`. The model considers class `0` as background. If your dataset does not contain the background class, you should not have `0` in your `labels`. For example, assuming you have just two classes, *cat* and *dog*, you can define `1` (not `0`) to represent *cats* and `2` to represent *dogs*. So, for instance, if one of the images has both classes, your `labels` tensor should look like `[1, 2]`.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratios), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

Writing a custom dataset for PennFudan

Let's write a dataset for the PennFudan dataset. After [downloading and extracting the zip file](#), we have the following folder structure:

::

```
PennFudanPed/ PedMasks/ FudanPed00001_mask.png FudanPed00002_mask.png
FudanPed00003_mask.png FudanPed00004_mask.png ... PNGImages/ FudanPed00001.png
FudanPed00002.png FudanPed00003.png FudanPed00004.png
```

Here is one example of a pair of images and segmentation masks

So each image has a corresponding segmentation mask, where each color correspond to a different instance. Let's write a :class:torch.utils.data.Dataset class for this dataset. In the code below, we are wrapping images, bounding boxes and masks into torchvision.TVTensor classes so that we will be able to apply torchvision built-in transformations ([new Transforms API](#)) for the given object detection and segmentation task. Namely, image tensors will be wrapped by :class:torchvision.tv_tensors.Image, bounding boxes into :class:torchvision.tv_tensors.BoundingBoxes and masks into :class:torchvision.tv_tensors.Mask. As torchvision.TVTensor are :class:torch.Tensor subclasses, wrapped objects are also tensors and inherit the plain :class:torch.Tensor API. For more information about torchvision tv_tensors see [this documentation](#).

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
```

```

        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,)), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,)), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

That's all for the dataset. Now let's define a model that can perform predictions on this dataset.

Defining your model

In this tutorial, we will be using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.

Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

There are two common situations where one might want to modify one of the available models in TorchVision Model Zoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

2 - Modifying the model to add a different backbone

```
import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator

# load a pre-trained model for classification and return
# only the features
backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features
# `FasterRCNN` needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(
    sizes=((32, 64, 128, 256, 512),),
    aspect_ratios=((0.5, 1.0, 2.0),)
)

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# `OrderedDict[Tensor]`, and in `featmap_names` you can choose
# which
```

```

# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(
    featmap_names=['0'],
    output_size=7,
    sampling_ratio=2
)

# put the pieces together inside a Faster-RCNN model
model = FasterRCNN(
    backbone,
    num_classes=2,
    rpn_anchor_generator=anchor_generator,
    box_roi_pool=roi_pooler
)

Downloading: "https://download.pytorch.org/models/mobilenet_v2-7ebf99e0.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-7ebf99e0.pth
100%|██████████| 13.6M/13.6M [00:00<00:00, 80.2MB/s]

```

Object detection and instance segmentation model for PennFudan Dataset

In our case, we want to finetune from a pre-trained model, given that our dataset is very small, so we will be following approach number 1.

Here we want to also compute the instance segmentation masks, so we will be using Mask R-CNN:

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
    torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
    num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
    model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256

```

```

# and replace the mask predictor with a new one
model.roi_heads.mask_predictor = MaskRCNNPredictor(
    in_features_mask,
    hidden_layer,
    num_classes
)

return model

```

That's it, this will make `model` be ready to be trained and evaluated on your custom dataset.

Putting everything together

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py` and `references/detection/utils.py`. Just download everything under `references/detection` to your folder and use them here. On Linux if you have `wget`, you can download them using below commands:

```

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API`
<https://pytorch.org/vision/stable/transforms.html>`
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#
# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:

```

```

        transforms.append(T.RandomHorizontalFlip(0.5))
        transforms.append(T.ToDtype(torch.float, scale=True))
        transforms.append(T.ToPureTensor())
        return T.Compose(transforms)

# Testing ``forward()`` method (Optional)
# -----
#
# Before iterating over the dataset, it's good to see what the model
# expects during training and inference time on sample data.
import utils

model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
print(output)

# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions
print(predictions[0])

```

Downloading:

"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth" to

/root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth

100%|██████████| 160M/160M [00:00<00:00, 171MB/s]

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please

be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn(_create_warning_msg(
```

```
{'loss_classifier': tensor(0.2291, grad_fn=<NllLossBackward0>),  
'loss_box_reg': tensor(0.1100, grad_fn=<DivBackward0>),  
'loss_objectness': tensor(0.0078,  
grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), 'loss_rpn_box_reg':  
tensor(0.0105, grad_fn=<DivBackward0>)}  
{'boxes': tensor([], size=(0, 4), grad_fn=<StackBackward0>), 'labels':  
tensor([], dtype=torch.int64), 'scores': tensor([],  
grad_fn=<IndexBackward0>)}
```

Let's now write the main function which performs the training and the validation:

```
from engine import train_one_epoch, evaluate  
  
# train on the GPU or on the CPU, if a GPU is not available  
device = torch.device('cuda') if torch.cuda.is_available() else  
torch.device('cpu')  
  
# our dataset has two classes only - background and person  
num_classes = 2  
# use our dataset and defined transformations  
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',  
get_transform(train=True))  
dataset_test = PennFudanDataset('drive/MyDrive/data/PennFudanPed',  
get_transform(train=False))  
  
# split the dataset in train and test set  
indices = torch.randperm(len(dataset)).tolist()  
dataset = torch.utils.data.Subset(dataset, indices[:-50])  
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])  
  
# define training and validation data loaders  
data_loader = torch.utils.data.DataLoader(  
    dataset,  
    batch_size=2,  
    shuffle=True,  
    num_workers=4,  
    collate_fn=utils.collate_fn  
)  
  
data_loader_test = torch.utils.data.DataLoader(  
    dataset_test,  
    batch_size=1,  
    shuffle=False,  
    num_workers=4,  
    collate_fn=utils.collate_fn
```

```

)

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
    print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

Downloading:

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-bf2d0c1e.pth" to
/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-bf2d0c1e.pth

100%|██████████| 170M/170M [00:01<00:00, 154MB/s]

Epoch: [0] [0/60] eta: 0:09:12 lr: 0.000090 loss: 5.0986 (5.0986)
loss_classifier: 0.5976 (0.5976) loss_box_reg: 0.1450 (0.1450)
loss_mask: 4.3445 (4.3445) loss_objectness: 0.0105 (0.0105)
loss_rpn_box_reg: 0.0011 (0.0011) time: 9.2085 data: 1.7603 max
mem: 2065

```
Epoch: [0] [10/60] eta: 0:01:05 lr: 0.000936 loss: 1.7777 (2.5739)
loss_classifier: 0.4427 (0.4398) loss_box_reg: 0.2589 (0.2873)
loss_mask: 1.1316 (1.8147) loss_objectness: 0.0235 (0.0256)
loss_rpn_box_reg: 0.0042 (0.0064) time: 1.3122 data: 0.1694 max
mem: 2764
Epoch: [0] [20/60] eta: 0:00:37 lr: 0.001783 loss: 0.9253 (1.6445)
loss_classifier: 0.1694 (0.2943) loss_box_reg: 0.2363 (0.2486)
loss_mask: 0.4028 (1.0708) loss_objectness: 0.0174 (0.0243)
loss_rpn_box_reg: 0.0060 (0.0064) time: 0.5337 data: 0.0097 max
mem: 3211
Epoch: [0] [30/60] eta: 0:00:24 lr: 0.002629 loss: 0.5503 (1.2777)
loss_classifier: 0.0942 (0.2279) loss_box_reg: 0.2240 (0.2401)
loss_mask: 0.1915 (0.7836) loss_objectness: 0.0093 (0.0189)
loss_rpn_box_reg: 0.0062 (0.0071) time: 0.5428 data: 0.0086 max
mem: 3211
Epoch: [0] [40/60] eta: 0:00:15 lr: 0.003476 loss: 0.4641 (1.0915)
loss_classifier: 0.0733 (0.1892) loss_box_reg: 0.2205 (0.2392)
loss_mask: 0.1874 (0.6412) loss_objectness: 0.0037 (0.0150)
loss_rpn_box_reg: 0.0075 (0.0069) time: 0.5498 data: 0.0110 max
mem: 3211
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.4495 (0.9692)
loss_classifier: 0.0539 (0.1633) loss_box_reg: 0.2102 (0.2342)
loss_mask: 0.1874 (0.5518) loss_objectness: 0.0021 (0.0130)
loss_rpn_box_reg: 0.0058 (0.0069) time: 0.5793 data: 0.0148 max
mem: 3222
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3703 (0.8745)
loss_classifier: 0.0415 (0.1442) loss_box_reg: 0.1577 (0.2203)
loss_mask: 0.1652 (0.4921) loss_objectness: 0.0017 (0.0112)
loss_rpn_box_reg: 0.0058 (0.0067) time: 0.5651 data: 0.0123 max
mem: 3222
Epoch: [0] Total time: 0:00:41 (0.6973 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.2362 (0.2362)
evaluator_time: 0.0045 (0.0045) time: 0.5697 data: 0.3278 max mem:
3222
Test: [49/50] eta: 0:00:00 model_time: 0.0982 (0.1290)
evaluator_time: 0.0053 (0.0121) time: 0.1231 data: 0.0049 max mem:
3222
Test: Total time: 0:00:07 (0.1594 s / it)
Averaged stats: model_time: 0.0982 (0.1290) evaluator_time: 0.0053
(0.0121)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.584
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.676
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.267
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.596
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.596
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.266
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.667
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.667
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.668
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.664
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.961
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.798
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.294
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.432
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.688
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.300
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.722
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.723
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.533
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.692
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.732
Epoch: [1] [ 0/60] eta: 0:00:52  lr: 0.005000  loss: 0.2154 (0.2154)
loss_classifier: 0.0134 (0.0134)  loss_box_reg: 0.0351 (0.0351)

```

```
loss_mask: 0.1655 (0.1655) loss_objectness: 0.0000 (0.0000)
loss_rpn_box_reg: 0.0014 (0.0014) time: 0.8740 data: 0.3722 max
mem: 3222
Epoch: [1] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2632 (0.3034)
loss_classifier: 0.0393 (0.0371) loss_box_reg: 0.0935 (0.1025)
loss_mask: 0.1525 (0.1554) loss_objectness: 0.0010 (0.0015)
loss_rpn_box_reg: 0.0036 (0.0068) time: 0.5839 data: 0.0398 max
mem: 3222
Epoch: [1] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2812 (0.3210)
loss_classifier: 0.0393 (0.0425) loss_box_reg: 0.1028 (0.1159)
loss_mask: 0.1403 (0.1537) loss_objectness: 0.0010 (0.0017)
loss_rpn_box_reg: 0.0061 (0.0072) time: 0.5770 data: 0.0092 max
mem: 3222
Epoch: [1] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2619 (0.2951)
loss_classifier: 0.0371 (0.0390) loss_box_reg: 0.1001 (0.1037)
loss_mask: 0.1330 (0.1445) loss_objectness: 0.0008 (0.0017)
loss_rpn_box_reg: 0.0057 (0.0061) time: 0.5975 data: 0.0097 max
mem: 3499
Epoch: [1] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2293 (0.2814)
loss_classifier: 0.0254 (0.0362) loss_box_reg: 0.0697 (0.0951)
loss_mask: 0.1265 (0.1430) loss_objectness: 0.0004 (0.0014)
loss_rpn_box_reg: 0.0037 (0.0057) time: 0.5854 data: 0.0087 max
mem: 3499
Epoch: [1] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2533 (0.2769)
loss_classifier: 0.0271 (0.0350) loss_box_reg: 0.0671 (0.0901)
loss_mask: 0.1385 (0.1447) loss_objectness: 0.0004 (0.0017)
loss_rpn_box_reg: 0.0040 (0.0055) time: 0.5632 data: 0.0091 max
mem: 3499
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2533 (0.2748)
loss_classifier: 0.0292 (0.0352) loss_box_reg: 0.0677 (0.0869)
loss_mask: 0.1429 (0.1455) loss_objectness: 0.0004 (0.0015)
loss_rpn_box_reg: 0.0044 (0.0058) time: 0.5583 data: 0.0081 max
mem: 3499
Epoch: [1] Total time: 0:00:35 (0.5834 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:36 model_time: 0.2662 (0.2662)
evaluator_time: 0.0041 (0.0041) time: 0.7247 data: 0.4532 max mem:
3499
Test: [49/50] eta: 0:00:00 model_time: 0.0967 (0.1095)
evaluator_time: 0.0038 (0.0059) time: 0.1158 data: 0.0042 max mem:
3499
Test: Total time: 0:00:06 (0.1344 s / it)
Averaged stats: model_time: 0.0967 (0.1095) evaluator_time: 0.0038
(0.0059)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
```

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.738

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.969

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.932

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.381

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.688

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.752

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.330

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.796

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.796

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.500

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.775

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.806

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.692

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.974

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.826

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.329

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.619

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.706

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.311

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.748

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.748

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.533

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.750

Average Recall (AR) @[IoU=0.50:0.95 | area= large |

```
maxDets=100 ] = 0.753
Epoch: [2] [ 0/60] eta: 0:00:53 lr: 0.005000 loss: 0.1752 (0.1752)
loss_classifier: 0.0203 (0.0203) loss_box_reg: 0.0485 (0.0485)
loss_mask: 0.1022 (0.1022) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0038 (0.0038) time: 0.8881 data: 0.3680 max
mem: 3499
Epoch: [2] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2351 (0.2578)
loss_classifier: 0.0322 (0.0362) loss_box_reg: 0.0631 (0.0793)
loss_mask: 0.1351 (0.1367) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0038 (0.0048) time: 0.5899 data: 0.0413 max
mem: 3499
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2351 (0.2564)
loss_classifier: 0.0303 (0.0325) loss_box_reg: 0.0725 (0.0793)
loss_mask: 0.1351 (0.1394) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0036 (0.0046) time: 0.5669 data: 0.0089 max
mem: 3499
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2122 (0.2387)
loss_classifier: 0.0246 (0.0306) loss_box_reg: 0.0517 (0.0711)
loss_mask: 0.1209 (0.1320) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0043) time: 0.5666 data: 0.0091 max
mem: 3499
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2139 (0.2401)
loss_classifier: 0.0334 (0.0315) loss_box_reg: 0.0517 (0.0713)
loss_mask: 0.1146 (0.1319) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0031 (0.0046) time: 0.5879 data: 0.0099 max
mem: 3506
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2241 (0.2373)
loss_classifier: 0.0339 (0.0311) loss_box_reg: 0.0579 (0.0691)
loss_mask: 0.1267 (0.1320) loss_objectness: 0.0007 (0.0007)
loss_rpn_box_reg: 0.0034 (0.0044) time: 0.5892 data: 0.0091 max
mem: 3506
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2103 (0.2388)
loss_classifier: 0.0257 (0.0310) loss_box_reg: 0.0546 (0.0685)
loss_mask: 0.1331 (0.1340) loss_objectness: 0.0005 (0.0008)
loss_rpn_box_reg: 0.0029 (0.0045) time: 0.5629 data: 0.0074 max
mem: 3506
Epoch: [2] Total time: 0:00:34 (0.5829 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1925 (0.1925)
evaluator_time: 0.0037 (0.0037) time: 0.5428 data: 0.3454 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0969 (0.1068)
evaluator_time: 0.0032 (0.0049) time: 0.1137 data: 0.0037 max mem:
3506
Test: Total time: 0:00:06 (0.1284 s / it)
Averaged stats: model_time: 0.0969 (0.1068) evaluator_time: 0.0032
(0.0049)
Accumulating evaluation results...
```

```

DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.764
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.905
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.292
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.696
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.786
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.339
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.820
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.820
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.433
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.758
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.837
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.748
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.973
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.938
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.315
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.567
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.326
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.433

```



```
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.758
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.805
Epoch: [3] [ 0/60] eta: 0:00:55 lr: 0.000500 loss: 0.1499 (0.1499)
loss_classifier: 0.0188 (0.0188) loss_box_reg: 0.0275 (0.0275)
loss_mask: 0.0995 (0.0995) loss_objectness: 0.0013 (0.0013)
loss_rpn_box_reg: 0.0028 (0.0028) time: 0.9178 data: 0.3975 max
mem: 3506
Epoch: [3] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1933 (0.1984)
loss_classifier: 0.0250 (0.0240) loss_box_reg: 0.0381 (0.0457)
loss_mask: 0.1201 (0.1236) loss_objectness: 0.0010 (0.0020)
loss_rpn_box_reg: 0.0023 (0.0030) time: 0.6011 data: 0.0473 max
mem: 3506
Epoch: [3] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1933 (0.1941)
loss_classifier: 0.0217 (0.0244) loss_box_reg: 0.0381 (0.0418)
loss_mask: 0.1204 (0.1229) loss_objectness: 0.0007 (0.0018)
loss_rpn_box_reg: 0.0023 (0.0033) time: 0.5586 data: 0.0100 max
mem: 3506
Epoch: [3] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1934 (0.1968)
loss_classifier: 0.0248 (0.0255) loss_box_reg: 0.0394 (0.0434)
loss_mask: 0.1204 (0.1232) loss_objectness: 0.0004 (0.0016)
loss_rpn_box_reg: 0.0032 (0.0032) time: 0.5505 data: 0.0101 max
mem: 3506
Epoch: [3] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1934 (0.1993)
loss_classifier: 0.0257 (0.0261) loss_box_reg: 0.0436 (0.0445)
loss_mask: 0.1213 (0.1236) loss_objectness: 0.0006 (0.0017)
loss_rpn_box_reg: 0.0034 (0.0034) time: 0.5638 data: 0.0108 max
mem: 3506
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1763 (0.1969)
loss_classifier: 0.0241 (0.0258) loss_box_reg: 0.0361 (0.0436)
loss_mask: 0.1137 (0.1223) loss_objectness: 0.0007 (0.0019)
loss_rpn_box_reg: 0.0025 (0.0032) time: 0.5697 data: 0.0088 max
mem: 3506
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1844 (0.1977)
loss_classifier: 0.0217 (0.0261) loss_box_reg: 0.0374 (0.0444)
loss_mask: 0.1123 (0.1220) loss_objectness: 0.0004 (0.0017)
loss_rpn_box_reg: 0.0027 (0.0034) time: 0.5720 data: 0.0080 max
mem: 3506
Epoch: [3] Total time: 0:00:34 (0.5757 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.2132 (0.2132)
evaluator_time: 0.0038 (0.0038) time: 0.5280 data: 0.3096 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.1036 (0.1098)
evaluator_time: 0.0050 (0.0060) time: 0.1220 data: 0.0047 max mem:
3506
Test: Total time: 0:00:06 (0.1367 s / it)
```

Averaged stats: model_time: 0.1036 (0.1098) evaluator_time: 0.0050 (0.0060)

Accumulating evaluation results...

DONE (t=0.03s).

Accumulating evaluation results...

DONE (t=0.03s).

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.791

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.977

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.923

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.372

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.703

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.812

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.348

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.850

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.850

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.567

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.775

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.866

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.754

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.977

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.937

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.368

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.523

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.773

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.331

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.803

Average Recall (AR) @[IoU=0.50:0.95 | area= all |

```
maxDets=100 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.600
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
Epoch: [4] [ 0/60] eta: 0:01:20 lr: 0.000500 loss: 0.2263 (0.2263)
loss_classifier: 0.0418 (0.0418) loss_box_reg: 0.0667 (0.0667)
loss_mask: 0.1081 (0.1081) loss_objectness: 0.0019 (0.0019)
loss_rpn_box_reg: 0.0078 (0.0078) time: 1.3454 data: 0.6900 max
mem: 3506
Epoch: [4] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.1923 (0.1878)
loss_classifier: 0.0294 (0.0274) loss_box_reg: 0.0394 (0.0393)
loss_mask: 0.1081 (0.1172) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0016 (0.0031) time: 0.6336 data: 0.0698 max
mem: 3506
Epoch: [4] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.1878 (0.1919)
loss_classifier: 0.0271 (0.0274) loss_box_reg: 0.0394 (0.0438)
loss_mask: 0.1124 (0.1166) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0022 (0.0034) time: 0.5651 data: 0.0086 max
mem: 3506
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1777 (0.1855)
loss_classifier: 0.0248 (0.0265) loss_box_reg: 0.0356 (0.0394)
loss_mask: 0.1126 (0.1156) loss_objectness: 0.0002 (0.0009)
loss_rpn_box_reg: 0.0017 (0.0030) time: 0.5693 data: 0.0113 max
mem: 3506
Epoch: [4] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1566 (0.1846)
loss_classifier: 0.0217 (0.0255) loss_box_reg: 0.0259 (0.0393)
loss_mask: 0.1126 (0.1161) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0014 (0.0028) time: 0.5630 data: 0.0128 max
mem: 3506
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1653 (0.1845)
loss_classifier: 0.0224 (0.0254) loss_box_reg: 0.0278 (0.0396)
loss_mask: 0.1163 (0.1159) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5632 data: 0.0133 max
mem: 3506
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1707 (0.1862)
loss_classifier: 0.0241 (0.0255) loss_box_reg: 0.0333 (0.0407)
loss_mask: 0.1126 (0.1164) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0029 (0.0030) time: 0.5827 data: 0.0109 max
mem: 3506
Epoch: [4] Total time: 0:00:35 (0.5868 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.1612 (0.1612)
evaluator_time: 0.0042 (0.0042) time: 0.5203 data: 0.3537 max mem:
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0973 (0.1117)
```

evaluator_time: 0.0033 (0.0072) time: 0.1171 data: 0.0053 max mem: 3506

Test: Total time: 0:00:06 (0.1387 s / it)

Averaged stats: model_time: 0.0973 (0.1117) evaluator_time: 0.0033 (0.0072)

Accumulating evaluation results...

DONE (t=0.02s).

Accumulating evaluation results...

DONE (t=0.02s).

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.802

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.975

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.913

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.335

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.700

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.825

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.354

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=10] = 0.849

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.849

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.467

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.767

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.868

IoU metric: segm

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.758

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.975

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.945

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.313

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.527

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.780

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=1] = 0.334

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=

```
10 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [5] [ 0/60] eta: 0:01:01 lr: 0.000500 loss: 0.1311 (0.1311)
loss_classifier: 0.0145 (0.0145) loss_box_reg: 0.0197 (0.0197)
loss_mask: 0.0942 (0.0942) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0024 (0.0024) time: 1.0227 data: 0.4592 max
mem: 3506
Epoch: [5] [10/60] eta: 0:00:27 lr: 0.000500 loss: 0.1704 (0.1697)
loss_classifier: 0.0172 (0.0197) loss_box_reg: 0.0356 (0.0336)
loss_mask: 0.1118 (0.1133) loss_objectness: 0.0006 (0.0010)
loss_rpn_box_reg: 0.0024 (0.0021) time: 0.5558 data: 0.0470 max
mem: 3506
Epoch: [5] [20/60] eta: 0:00:22 lr: 0.000500 loss: 0.1704 (0.1813)
loss_classifier: 0.0216 (0.0224) loss_box_reg: 0.0365 (0.0369)
loss_mask: 0.1150 (0.1184) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0026 (0.0028) time: 0.5431 data: 0.0072 max
mem: 3506
Epoch: [5] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1792 (0.1873)
loss_classifier: 0.0254 (0.0234) loss_box_reg: 0.0371 (0.0392)
loss_mask: 0.1242 (0.1211) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0028) time: 0.5843 data: 0.0089 max
mem: 3506
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1743 (0.1817)
loss_classifier: 0.0224 (0.0225) loss_box_reg: 0.0323 (0.0373)
loss_mask: 0.1098 (0.1183) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0017 (0.0027) time: 0.5693 data: 0.0092 max
mem: 3506
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1669 (0.1854)
loss_classifier: 0.0195 (0.0239) loss_box_reg: 0.0334 (0.0392)
loss_mask: 0.1054 (0.1183) loss_objectness: 0.0006 (0.0011)
loss_rpn_box_reg: 0.0024 (0.0029) time: 0.5706 data: 0.0094 max
mem: 3506
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1669 (0.1838)
loss_classifier: 0.0234 (0.0236) loss_box_reg: 0.0390 (0.0390)
loss_mask: 0.1064 (0.1172) loss_objectness: 0.0003 (0.0011)
loss_rpn_box_reg: 0.0025 (0.0029) time: 0.5905 data: 0.0086 max
mem: 3506
Epoch: [5] Total time: 0:00:34 (0.5787 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:32 model_time: 0.1621 (0.1621)
evaluator_time: 0.0040 (0.0040) time: 0.6579 data: 0.4907 max mem:
```

```
3506
Test: [49/50] eta: 0:00:00 model_time: 0.0967 (0.1065)
evaluator_time: 0.0041 (0.0050) time: 0.1136 data: 0.0037 max mem:
3506
Test: Total time: 0:00:06 (0.1310 s / it)
Averaged stats: model_time: 0.0967 (0.1065) evaluator_time: 0.0041
(0.0050)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.809
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.973
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.913
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.706
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.831
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.854
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.854
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.760
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.313
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.562
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.781
```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.334
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.802
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [6] [ 0/60] eta: 0:00:59  lr: 0.000050  loss: 0.1643 (0.1643)
loss_classifier: 0.0163 (0.0163)  loss_box_reg: 0.0271 (0.0271)
loss_mask: 0.1181 (0.1181)  loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0027 (0.0027)  time: 0.9919  data: 0.4327  max
mem: 3506
Epoch: [6] [10/60] eta: 0:00:29  lr: 0.000050  loss: 0.1693 (0.1709)
loss_classifier: 0.0187 (0.0200)  loss_box_reg: 0.0339 (0.0332)
loss_mask: 0.1108 (0.1150)  loss_objectness: 0.0003 (0.0004)
loss_rpn_box_reg: 0.0023 (0.0024)  time: 0.5922  data: 0.0479  max
mem: 3506
Epoch: [6] [20/60] eta: 0:00:23  lr: 0.000050  loss: 0.1793 (0.1807)
loss_classifier: 0.0242 (0.0241)  loss_box_reg: 0.0339 (0.0363)
loss_mask: 0.1108 (0.1167)  loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0023 (0.0028)  time: 0.5668  data: 0.0090  max
mem: 3506
Epoch: [6] [30/60] eta: 0:00:17  lr: 0.000050  loss: 0.1793 (0.1808)
loss_classifier: 0.0242 (0.0245)  loss_box_reg: 0.0327 (0.0368)
loss_mask: 0.1136 (0.1161)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0025 (0.0028)  time: 0.5922  data: 0.0084  max
mem: 3506
Epoch: [6] [40/60] eta: 0:00:11  lr: 0.000050  loss: 0.1654 (0.1841)
loss_classifier: 0.0206 (0.0245)  loss_box_reg: 0.0300 (0.0384)
loss_mask: 0.1133 (0.1173)  loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0025 (0.0030)  time: 0.5777  data: 0.0092  max
mem: 3506
Epoch: [6] [50/60] eta: 0:00:05  lr: 0.000050  loss: 0.1654 (0.1815)
loss_classifier: 0.0191 (0.0243)  loss_box_reg: 0.0300 (0.0377)
loss_mask: 0.1091 (0.1157)  loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0022 (0.0029)  time: 0.5662  data: 0.0090  max
mem: 3780
Epoch: [6] [59/60] eta: 0:00:00  lr: 0.000050  loss: 0.1691 (0.1797)
loss_classifier: 0.0206 (0.0238)  loss_box_reg: 0.0285 (0.0370)
loss_mask: 0.1148 (0.1152)  loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0020 (0.0028)  time: 0.5693  data: 0.0076  max
mem: 3780
Epoch: [6] Total time: 0:00:35 (0.5842 s / it)
creating index...
```

```
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1869 (0.1869)
evaluator_time: 0.0060 (0.0060) time: 0.5776 data: 0.3833 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.0973 (0.1077)
evaluator_time: 0.0033 (0.0052) time: 0.1150 data: 0.0036 max mem:
3780
Test: Total time: 0:00:06 (0.1305 s / it)
Averaged stats: model_time: 0.0973 (0.1077) evaluator_time: 0.0033
(0.0052)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.808
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.704
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.831
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.355
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.855
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.873
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.759
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
```



```
maxDets=100 ] = 0.528
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.334
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.814
Epoch: [7] [ 0/60] eta: 0:01:28 lr: 0.000050 loss: 0.1616 (0.1616)
loss_classifier: 0.0205 (0.0205) loss_box_reg: 0.0288 (0.0288)
loss_mask: 0.1102 (0.1102) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0021 (0.0021) time: 1.4667 data: 0.6278 max
mem: 3780
Epoch: [7] [10/60] eta: 0:00:32 lr: 0.000050 loss: 0.1826 (0.1937)
loss_classifier: 0.0241 (0.0265) loss_box_reg: 0.0363 (0.0411)
loss_mask: 0.1177 (0.1229) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0021 (0.0025) time: 0.6578 data: 0.0641 max
mem: 3780
Epoch: [7] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1834 (0.1907)
loss_classifier: 0.0241 (0.0258) loss_box_reg: 0.0402 (0.0413)
loss_mask: 0.1172 (0.1201) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0026 (0.0029) time: 0.5638 data: 0.0084 max
mem: 3780
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1867 (0.1909)
loss_classifier: 0.0240 (0.0260) loss_box_reg: 0.0443 (0.0416)
loss_mask: 0.1058 (0.1197) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0030) time: 0.5671 data: 0.0102 max
mem: 3780
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1575 (0.1866)
loss_classifier: 0.0200 (0.0247) loss_box_reg: 0.0340 (0.0401)
loss_mask: 0.1044 (0.1184) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5652 data: 0.0095 max
mem: 3780
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1498 (0.1816)
loss_classifier: 0.0164 (0.0240) loss_box_reg: 0.0250 (0.0378)
loss_mask: 0.1044 (0.1164) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0017 (0.0027) time: 0.5425 data: 0.0088 max
mem: 3780
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1522 (0.1790)
loss_classifier: 0.0212 (0.0239) loss_box_reg: 0.0261 (0.0367)
loss_mask: 0.1036 (0.1150) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0024 (0.0028) time: 0.5749 data: 0.0091 max
```

```
mem: 3780
Epoch: [7] Total time: 0:00:35 (0.5869 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1866 (0.1866)
evaluator_time: 0.0039 (0.0039) time: 0.5075 data: 0.3155 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.1186 (0.1148)
evaluator_time: 0.0055 (0.0068) time: 0.1294 data: 0.0048 max mem:
3780
Test: Total time: 0:00:06 (0.1398 s / it)
Averaged stats: model_time: 0.1186 (0.1148) evaluator_time: 0.0055
(0.0068)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.810
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.834
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.856
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.856
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.876
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.531
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.802
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.815
Epoch: [8] [ 0/60] eta: 0:01:06 lr: 0.000050 loss: 0.1370 (0.1370)
loss_classifier: 0.0163 (0.0163) loss_box_reg: 0.0187 (0.0187)
loss_mask: 0.0959 (0.0959) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0057 (0.0057) time: 1.1106 data: 0.4501 max
mem: 3780
Epoch: [8] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1865 (0.2064)
loss_classifier: 0.0221 (0.0261) loss_box_reg: 0.0353 (0.0489)
loss_mask: 0.1295 (0.1263) loss_objectness: 0.0002 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0044) time: 0.5999 data: 0.0469 max
mem: 3780
Epoch: [8] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1865 (0.1967)
loss_classifier: 0.0244 (0.0263) loss_box_reg: 0.0408 (0.0450)
loss_mask: 0.1164 (0.1205) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0030 (0.0043) time: 0.5597 data: 0.0091 max
mem: 3780
Epoch: [8] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1610 (0.1818)
loss_classifier: 0.0210 (0.0243) loss_box_reg: 0.0292 (0.0389)
loss_mask: 0.0997 (0.1145) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0026 (0.0036) time: 0.5602 data: 0.0101 max
mem: 3780
Epoch: [8] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1455 (0.1775)
loss_classifier: 0.0183 (0.0239) loss_box_reg: 0.0238 (0.0373)
loss_mask: 0.0973 (0.1124) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0021 (0.0032) time: 0.5544 data: 0.0081 max
mem: 3780
Epoch: [8] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1653 (0.1799)
loss_classifier: 0.0229 (0.0248) loss_box_reg: 0.0294 (0.0371)
loss_mask: 0.1098 (0.1143) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0030) time: 0.5777 data: 0.0088 max
mem: 3780
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1663 (0.1786)
```

```
loss_classifier: 0.0229 (0.0245) loss_box_reg: 0.0291 (0.0367)
loss_mask: 0.1133 (0.1139) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0029) time: 0.5682 data: 0.0086 max
mem: 3780
Epoch: [8] Total time: 0:00:34 (0.5768 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:42 model_time: 0.3319 (0.3319)
evaluator_time: 0.0050 (0.0050) time: 0.8421 data: 0.5039 max mem:
3780
Test: [49/50] eta: 0:00:00 model_time: 0.0978 (0.1148)
evaluator_time: 0.0033 (0.0058) time: 0.1142 data: 0.0036 max mem:
3780
Test: Total time: 0:00:07 (0.1426 s / it)
Averaged stats: model_time: 0.0978 (0.1148) evaluator_time: 0.0033
(0.0058)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.806
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.830
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.354
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.853
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.853
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
```

```

maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.529
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.816
Epoch: [9] [ 0/60] eta: 0:00:53 lr: 0.000005 loss: 0.1116 (0.1116)
loss_classifier: 0.0113 (0.0113) loss_box_reg: 0.0150 (0.0150)
loss_mask: 0.0852 (0.0852) loss_objectness: 0.0000 (0.0000)
loss_rpn_box_reg: 0.0002 (0.0002) time: 0.8908 data: 0.3519 max
mem: 3780
Epoch: [9] [10/60] eta: 0:00:29 lr: 0.000005 loss: 0.1478 (0.1777)
loss_classifier: 0.0200 (0.0243) loss_box_reg: 0.0245 (0.0380)
loss_mask: 0.0980 (0.1113) loss_objectness: 0.0003 (0.0013)
loss_rpn_box_reg: 0.0020 (0.0028) time: 0.5810 data: 0.0399 max
mem: 3780
Epoch: [9] [20/60] eta: 0:00:23 lr: 0.000005 loss: 0.1664 (0.1761)
loss_classifier: 0.0231 (0.0243) loss_box_reg: 0.0284 (0.0357)
loss_mask: 0.1047 (0.1119) loss_objectness: 0.0003 (0.0017)
loss_rpn_box_reg: 0.0020 (0.0025) time: 0.5605 data: 0.0095 max
mem: 3780
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1671 (0.1812)
loss_classifier: 0.0244 (0.0245) loss_box_reg: 0.0301 (0.0375)
loss_mask: 0.1119 (0.1150) loss_objectness: 0.0005 (0.0014)
loss_rpn_box_reg: 0.0025 (0.0029) time: 0.5652 data: 0.0098 max
mem: 3780
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1601 (0.1795)
loss_classifier: 0.0228 (0.0243) loss_box_reg: 0.0293 (0.0375)
loss_mask: 0.1107 (0.1135) loss_objectness: 0.0005 (0.0013)
loss_rpn_box_reg: 0.0029 (0.0029) time: 0.5756 data: 0.0094 max
mem: 3780
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1601 (0.1808)
loss_classifier: 0.0220 (0.0245) loss_box_reg: 0.0309 (0.0381)
loss_mask: 0.1089 (0.1141) loss_objectness: 0.0003 (0.0011)

```

```
loss_rpn_box_reg: 0.0029 (0.0029)  time: 0.5913  data: 0.0095  max
mem: 3780
Epoch: [9]  [59/60]  eta: 0:00:00  lr: 0.000005  loss: 0.1621 (0.1808)
loss_classifier: 0.0219 (0.0247)  loss_box_reg: 0.0326 (0.0377)
loss_mask: 0.1089 (0.1143)  loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0033 (0.0029)  time: 0.5792  data: 0.0088  max
mem: 3780
Epoch: [9] Total time: 0:00:34 (0.5811 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:40  model_time: 0.3076 (0.3076)
evaluator_time: 0.0055 (0.0055)  time: 0.8024  data: 0.4880  max mem:
3780
Test:  [49/50]  eta: 0:00:00  model_time: 0.1000 (0.1120)
evaluator_time: 0.0034 (0.0052)  time: 0.1157  data: 0.0037  max mem:
3780
Test: Total time: 0:00:06 (0.1375 s / it)
Averaged stats: model_time: 0.1000 (0.1120)  evaluator_time: 0.0034
(0.0052)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.806
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.975
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.919
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.348
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.700
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.830
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.354
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.853
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.853
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.872
IoU metric: segm
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.761
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.975
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.313
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.529
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.816
That's it!

```

So after one epoch of training, we obtain a COCO-style mAP > 50, and a mask mAP of 65.

But what do the predictions look like? Let's take one image in the dataset and verify

```

import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

image =
read_image("drive/MyDrive/_static/img/tv_tutorial/tv_image05.png")
eval_transform = get_transform(train=False)

model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

```


```

image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))
<matplotlib.image.AxesImage at 0x7c62d689c3d0>

```




```
pred_labels
```

```
['pedestrian: 0.996',  
'pedestrian: 0.996',  
'pedestrian: 0.994',  
'pedestrian: 0.987',  
'pedestrian: 0.976',  
'pedestrian: 0.937',  
'pedestrian: 0.203',  
'pedestrian: 0.181',  
'pedestrian: 0.119',  
'pedestrian: 0.057']
```

```
pred_boxes
```

```
tensor([[170, 129, 268, 410],  
        [297, 117, 369, 340],  
        [445, 101, 508, 314],  
        [515, 118, 565, 255],  
        [361, 120, 440, 325],  
        [239, 125, 308, 360],  
        [268, 126, 507, 329],  
        [486, 105, 517, 172],  
        [541, 110, 564, 250],  
        [284, 109, 317, 169]] , device='cuda:0')
```

The results look good!

Wrapping up

In this tutorial, you have learned how to create your own training pipeline for object detection models on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes and segmentation masks. You also leveraged a Mask R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-GPU training, check `references/detection/train.py`, which is present in the torchvision repository.

You can download a full source file for this tutorial [here](#).

```
image = read_image("drive/MyDrive/Beatles_-_Abbey_Road.jpg")  
eval_transform = get_transform(train=False)  
  
model.eval()  
with torch.no_grad():  
    x = eval_transform(image)  
    # convert RGBA -> RGB and move to device  
    x = x[:3, ...].to(device)  
    predictions = model([x, ])
```

```
pred = predictions[0]


image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red", font_size=2)

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

/usr/local/lib/python3.10/dist-packages/torchvision/utils.py:223:
UserWarning: Argument 'font_size' will be ignored since 'font' is not
set.
  warnings.warn("Argument 'font_size' will be ignored since 'font' is
not set.")

<matplotlib.image.AxesImage at 0x7c62d68430a0>
```



```
pred_labels
```

```
['pedestrian: 0.977',  
'pedestrian: 0.975',  
'pedestrian: 0.933',  
'pedestrian: 0.496']
```

```
pred_boxes
```

```
tensor([[ 92, 161, 145, 269],  
        [ 21, 157,  81, 270],  
        [147, 168, 215, 286],  
        [224, 160, 293, 287]], device='cuda:0')
```

