

This is for Question 5 Part B and C using the finetuning option

```
# For tips on running notebooks in Google Colab, see
# https://pytorch.org/tutorials/beginner/colab
%matplotlib inline

from google.colab import drive
drive.mount('/content/drive') #Updated the spots where they needed the
locations to change, also downloaded the PennFudan dataset and
tutorial source for the test photo at the end
# Dataset: https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
# Test Photo:
https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425fa
a00c06bcf502/_static/img/tv_tutorial/tv_image05.png
#Source Code:
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

Mounted at /content/drive
```

TorchVision Object Detection Finetuning Tutorial

.. tip::

To get the most of this tutorial, we suggest using this [Colab Version](https://colab.research.google.com/github/pytorch/tutorials/blob/gh-pages/_downloads/torchvision_fineting_instance_segmentation.ipynb). This will allow you to experiment with the information presented below.

For this tutorial, we will be finetuning a pre-trained [Mask R-CNN](#) model on the [Penn-Fudan Database for Pedestrian Detection and Segmentation](#). It contains 170 images with 345 instances of pedestrians, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection and instance segmentation model on a custom dataset.

.. note ::

This tutorial works only with torchvision version `>=0.16` or `nightly`. If you're using `torchvision<=0.15`, please follow [this tutorial instead](https://github.com/pytorch/tutorials/blob/d686b662932a380a58b7683425faa00c06bcf502/intermediate_source/torchvision_tutorial.rst).

Defining the Dataset

The reference scripts for training object detection, instance segmentation and person keypoint detection allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return a tuple:

- `image`: :class:torchvision.tv_tensors.Image of shape `[3, H, W]`, a pure tensor, or a PIL Image of size `(H, W)`
- `target`: a dict containing the following fields
 - `boxes`, :class:torchvision.tv_tensors.BoundingBoxes of shape `[N, 4]`: the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels`, integer :class:torch.Tensor of shape `[N]`: the label for each bounding box. `0` represents always the background class.
 - `image_id`, int: an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area`, float :class:torch.Tensor of shape `[N]`: the area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd`, uint8 :class:torch.Tensor of shape `[N]`: instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks`, :class:torchvision.tv_tensors.Mask of shape `[N, H, W]`: the segmentation masks for each one of the objects

If your dataset is compliant with above requirements then it will work for both training and evaluation codes from the reference script. Evaluation code will use scripts from `pycocotools` which can be installed with `pip install pycocotools`.

.. note :: For Windows, please install `pycocotools` from [gautamchitnis_](https://github.com/gautamchitnis/cocoapi) with command

```
pip install
git+https://github.com/gautamchitnis/cocoapi.git@cocodataset-
master#subdirectory=PythonAPI
```

One note on the `labels`. The model considers class `0` as background. If your dataset does not contain the background class, you should not have `0` in your `labels`. For example, assuming you have just two classes, *cat* and *dog*, you can define `1` (not `0`) to represent *cats* and `2` to represent *dogs*. So, for instance, if one of the images has both classes, your `labels` tensor should look like `[1, 2]`.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratios), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

Writing a custom dataset for PennFudan

Let's write a dataset for the PennFudan dataset. After [downloading and extracting the zip file](#), we have the following folder structure:

::

```
PennFudanPed/ PedMasks/ FudanPed00001_mask.png FudanPed00002_mask.png
FudanPed00003_mask.png FudanPed00004_mask.png ... PNGImages/ FudanPed00001.png
FudanPed00002.png FudanPed00003.png FudanPed00004.png
```

Here is one example of a pair of images and segmentation masks

So each image has a corresponding segmentation mask, where each color correspond to a different instance. Let's write a :class:torch.utils.data.Dataset class for this dataset. In the code below, we are wrapping images, bounding boxes and masks into torchvision.TVTensor classes so that we will be able to apply torchvision built-in transformations ([new Transforms API](#)) for the given object detection and segmentation task. Namely, image tensors will be wrapped by :class:torchvision.tv_tensors.Image, bounding boxes into :class:torchvision.tv_tensors.BoundingBoxes and masks into :class:torchvision.tv_tensors.Mask. As torchvision.TVTensor are :class:torch.Tensor subclasses, wrapped objects are also tensors and inherit the plain :class:torch.Tensor API. For more information about torchvision tv_tensors see [this documentation](#).

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
```

```

        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
        labels = torch.ones((num_objs,), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

That's all for the dataset. Now let's define a model that can perform predictions on this dataset.

Defining your model

In this tutorial, we will be using [Mask R-CNN](#), which is based on top of [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.

Mask R-CNN adds an extra branch into Faster R-CNN, which also predicts segmentation masks for each instance.

There are two common situations where one might want to modify one of the available models in TorchVision Model Zoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

1 - Finetuning from a pretrained model

Let's suppose that you want to start from a model pre-trained on COCO and want to finetune it for your particular classes. Here is a possible way of doing it:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

# load a model pre-trained on COCO
model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)

# replace the classifier with a new one, that has
# num_classes which is user-defined
num_classes = 2 # 1 class (person) + background
# get number of input features for the classifier
in_features = model.roi_heads.box_predictor.cls_score.in_features
# replace the pre-trained head with a new one
model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

Downloading:
"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth" to
/root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-
258fb6c6.pth
100%|██████████| 160M/160M [00:01<00:00, 157MB/s]
```

Object detection and instance segmentation model for PennFudan Dataset

In our case, we want to finetune from a pre-trained model, given that our dataset is very small, so we will be following approach number 1.

Here we want to also compute the instance segmentation masks, so we will be using Mask R-CNN:

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256
    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(
        in_features_mask,
        hidden_layer,
        num_classes
    )

    return model
```

That's it, this will make `model` be ready to be trained and evaluated on your custom dataset.

Putting everything together

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py` and `references/detection/utils.py`. Just download everything under `references/detection` to your folder and use them here. On Linux if you have `wget`, you can download them using below commands:

```

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detection/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API
<https://pytorch.org/vision/stable/transforms.html>`_
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#
# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T

def get_transform(train):
    transforms = []
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
        transforms.append(T.ToDtype(torch.float, scale=True))
        transforms.append(T.ToPureTensor())
    return T.Compose(transforms)

# Testing ``forward()`` method (Optional)
# -----
#
# Before iterating over the dataset, it's good to see what the model
# expects during training and inference time on sample data.
import utils

model =
torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT"
)
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=True))
data_loader = torch.utils.data.DataLoader(

```

```

        dataset,
        batch_size=2,
        shuffle=True,
        num_workers=4,
        collate_fn=utils.collate_fn
    )

# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
print(output)

# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions
print(predictions[0])

/usr/local/lib/python3.10/dist-packages/torch/utils/data/
dataloader.py:557: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(

{'loss_classifier': tensor(0.0989, grad_fn=<NllLossBackward0>),
'loss_box_reg': tensor(0.0605, grad_fn=<DivBackward0>),
'loss_objectness': tensor(0.0056,
grad_fn=<BinaryCrossEntropyWithLogitsBackward0>), 'loss_rpn_box_reg':
tensor(0.0073, grad_fn=<DivBackward0>)}
{'boxes': tensor([], size=(0, 4), grad_fn=<StackBackward0>), 'labels':
tensor([], dtype=torch.int64), 'scores': tensor([],
grad_fn=<IndexBackward0>)}

```

Let's now write the main function which performs the training and the validation:

```

from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('drive/MyDrive/data/PennFudanPed',

```



```

get_transform(train=True))
dataset_test = PennFudanDataset('drive/MyDrive/data/PennFudanPed',
get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 10

```

```

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

```

Downloading:

```

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth" to
/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-
bf2d0c1e.pth
100%|██████████| 170M/170M [00:01<00:00, 92.5MB/s]

```

```

Epoch: [0] [ 0/60] eta: 0:09:22 lr: 0.000090 loss: 4.0232 (4.0232)
loss_classifier: 0.6608 (0.6608) loss_box_reg: 0.1815 (0.1815)
loss_mask: 3.1586 (3.1586) loss_objectness: 0.0194 (0.0194)
loss_rpn_box_reg: 0.0029 (0.0029) time: 9.3771 data: 1.6565 max
mem: 2596
Epoch: [0] [10/60] eta: 0:01:07 lr: 0.000936 loss: 1.7460 (2.3481)
loss_classifier: 0.4496 (0.4475) loss_box_reg: 0.3494 (0.3553)
loss_mask: 1.0930 (1.5102) loss_objectness: 0.0304 (0.0277)
loss_rpn_box_reg: 0.0061 (0.0074) time: 1.3599 data: 0.1567 max
mem: 2978
Epoch: [0] [20/60] eta: 0:00:39 lr: 0.001783 loss: 1.0728 (1.5656)
loss_classifier: 0.2308 (0.3128) loss_box_reg: 0.2986 (0.3006)
loss_mask: 0.4108 (0.9232) loss_objectness: 0.0153 (0.0225)
loss_rpn_box_reg: 0.0058 (0.0066) time: 0.5580 data: 0.0099 max
mem: 2978
Epoch: [0] [30/60] eta: 0:00:25 lr: 0.002629 loss: 0.5781 (1.2421)
loss_classifier: 0.1208 (0.2424) loss_box_reg: 0.1701 (0.2828)
loss_mask: 0.2195 (0.6921) loss_objectness: 0.0083 (0.0175)
loss_rpn_box_reg: 0.0041 (0.0072) time: 0.5665 data: 0.0111 max
mem: 3066
Epoch: [0] [40/60] eta: 0:00:15 lr: 0.003476 loss: 0.5432 (1.0732)
loss_classifier: 0.0734 (0.2017) loss_box_reg: 0.1987 (0.2805)
loss_mask: 0.2021 (0.5695) loss_objectness: 0.0044 (0.0144)
loss_rpn_box_reg: 0.0059 (0.0072) time: 0.5562 data: 0.0098 max
mem: 3066
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.4499 (0.9406)
loss_classifier: 0.0452 (0.1706) loss_box_reg: 0.1744 (0.2575)
loss_mask: 0.1770 (0.4932) loss_objectness: 0.0026 (0.0123)
loss_rpn_box_reg: 0.0052 (0.0070) time: 0.5357 data: 0.0101 max
mem: 3066
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3485 (0.8478)

```

```
loss_classifier: 0.0368 (0.1514) loss_box_reg: 0.1203 (0.2363)
loss_mask: 0.1702 (0.4425) loss_objectness: 0.0015 (0.0108)
loss_rpn_box_reg: 0.0044 (0.0069) time: 0.5412 data: 0.0090 max
mem: 3066
Epoch: [0] Total time: 0:00:42 (0.7033 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.2204 (0.2204)
evaluator_time: 0.0211 (0.0211) time: 0.5590 data: 0.3109 max mem:
3066
Test: [49/50] eta: 0:00:00 model_time: 0.1112 (0.1196)
evaluator_time: 0.0061 (0.0107) time: 0.1321 data: 0.0052 max mem:
3066
Test: Total time: 0:00:07 (0.1491 s / it)
Averaged stats: model_time: 0.1112 (0.1196) evaluator_time: 0.0061
(0.0107)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.662
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.954
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.849
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.357
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.559
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.680
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.338
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.727
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.727
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.736
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.693
Average Precision (AP) @[ IoU=0.50 | area= all |
```

```

maxDets=100 ] = 0.960
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.868
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.369
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.459
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.353
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.751
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.754
Epoch: [1] [ 0/60] eta: 0:01:02 lr: 0.005000 loss: 0.3659 (0.3659)
loss_classifier: 0.0570 (0.0570) loss_box_reg: 0.1407 (0.1407)
loss_mask: 0.1630 (0.1630) loss_objectness: 0.0016 (0.0016)
loss_rpn_box_reg: 0.0036 (0.0036) time: 1.0337 data: 0.4756 max
mem: 3066
Epoch: [1] [10/60] eta: 0:00:30 lr: 0.005000 loss: 0.3129 (0.3190)
loss_classifier: 0.0385 (0.0399) loss_box_reg: 0.1143 (0.1172)
loss_mask: 0.1380 (0.1558) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0049 (0.0052) time: 0.6191 data: 0.0496 max
mem: 3066
Epoch: [1] [20/60] eta: 0:00:24 lr: 0.005000 loss: 0.2991 (0.3156)
loss_classifier: 0.0334 (0.0396) loss_box_reg: 0.0931 (0.1061)
loss_mask: 0.1543 (0.1634) loss_objectness: 0.0007 (0.0016)
loss_rpn_box_reg: 0.0048 (0.0050) time: 0.5884 data: 0.0081 max
mem: 3066
Epoch: [1] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2847 (0.3090)
loss_classifier: 0.0421 (0.0408) loss_box_reg: 0.0901 (0.1063)
loss_mask: 0.1517 (0.1552) loss_objectness: 0.0015 (0.0017)
loss_rpn_box_reg: 0.0043 (0.0049) time: 0.5998 data: 0.0097 max
mem: 3464
Epoch: [1] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2822 (0.3030)
loss_classifier: 0.0421 (0.0400) loss_box_reg: 0.0887 (0.1028)
loss_mask: 0.1373 (0.1540) loss_objectness: 0.0007 (0.0014)
loss_rpn_box_reg: 0.0038 (0.0048) time: 0.6153 data: 0.0090 max
mem: 3464
Epoch: [1] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2812 (0.3023)
loss_classifier: 0.0341 (0.0400) loss_box_reg: 0.0887 (0.1020)
loss_mask: 0.1392 (0.1538) loss_objectness: 0.0005 (0.0014)

```

```
loss_rpn_box_reg: 0.0038 (0.0051)  time: 0.6257  data: 0.0098  max
mem: 3464
Epoch: [1]  [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.2764 (0.2969)
loss_classifier: 0.0402 (0.0396)  loss_box_reg: 0.0861 (0.0994)
loss_mask: 0.1343 (0.1512)  loss_objectness: 0.0006 (0.0014)
loss_rpn_box_reg: 0.0057 (0.0052)  time: 0.6134  data: 0.0094  max
mem: 3464
Epoch: [1] Total time: 0:00:36 (0.6161 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:39  model_time: 0.2075 (0.2075)
evaluator_time: 0.0281 (0.0281)  time: 0.7880  data: 0.5507  max mem:
3464
Test:  [49/50]  eta: 0:00:00  model_time: 0.0969 (0.1079)
evaluator_time: 0.0032 (0.0060)  time: 0.1103  data: 0.0038  max mem:
3464
Test: Total time: 0:00:06 (0.1350 s / it)
Averaged stats: model_time: 0.0969 (0.1079)  evaluator_time: 0.0032
(0.0060)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.777
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.966
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.916
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.367
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.581
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.801
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.396
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.827
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.827
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.725
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.846
IoU metric: segm
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.743
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.971
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.921
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.354
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.532
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.761
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.377
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.798
Epoch: [2] [ 0/60] eta: 0:00:57 lr: 0.005000 loss: 0.2624 (0.2624)
loss_classifier: 0.0328 (0.0328) loss_box_reg: 0.0768 (0.0768)
loss_mask: 0.1485 (0.1485) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0043 (0.0043) time: 0.9614 data: 0.3912 max
mem: 3464
Epoch: [2] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2395 (0.2361)
loss_classifier: 0.0326 (0.0314) loss_box_reg: 0.0622 (0.0653)
loss_mask: 0.1212 (0.1340) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0043 (0.0047) time: 0.5913 data: 0.0433 max
mem: 3464
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2164 (0.2381)
loss_classifier: 0.0306 (0.0326) loss_box_reg: 0.0592 (0.0659)
loss_mask: 0.1212 (0.1343) loss_objectness: 0.0005 (0.0006)
loss_rpn_box_reg: 0.0042 (0.0047) time: 0.5783 data: 0.0097 max
mem: 3464
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2076 (0.2267)
loss_classifier: 0.0270 (0.0301) loss_box_reg: 0.0574 (0.0631)
loss_mask: 0.1179 (0.1286) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0042) time: 0.6047 data: 0.0093 max
mem: 3464
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2118 (0.2347)
loss_classifier: 0.0270 (0.0311) loss_box_reg: 0.0595 (0.0677)
loss_mask: 0.1246 (0.1308) loss_objectness: 0.0002 (0.0005)
loss_rpn_box_reg: 0.0037 (0.0046) time: 0.5985 data: 0.0094 max
mem: 3464

```

```
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2745 (0.2435)
loss_classifier: 0.0311 (0.0325) loss_box_reg: 0.0837 (0.0720)
loss_mask: 0.1364 (0.1328) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0054 (0.0054) time: 0.5857 data: 0.0099 max
mem: 3464
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2684 (0.2462)
loss_classifier: 0.0331 (0.0332) loss_box_reg: 0.0837 (0.0742)
loss_mask: 0.1263 (0.1328) loss_objectness: 0.0006 (0.0008)
loss_rpn_box_reg: 0.0035 (0.0052) time: 0.5700 data: 0.0085 max
mem: 3464
Epoch: [2] Total time: 0:00:35 (0.5971 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1593 (0.1593)
evaluator_time: 0.0136 (0.0136) time: 0.5565 data: 0.3816 max mem:
3464
Test: [49/50] eta: 0:00:00 model_time: 0.0989 (0.1055)
evaluator_time: 0.0030 (0.0047) time: 0.1097 data: 0.0037 max mem:
3464
Test: Total time: 0:00:06 (0.1273 s / it)
Averaged stats: model_time: 0.0989 (0.1055) evaluator_time: 0.0030
(0.0047)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.746
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.352
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.579
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.760
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.385
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.788
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.798
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.739
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.972
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.890
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.394
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.495
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.751
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.376
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.781
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.783
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.667
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.725
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.792
Epoch: [3] [ 0/60] eta: 0:00:59 lr: 0.000500 loss: 0.2168 (0.2168)
loss_classifier: 0.0278 (0.0278) loss_box_reg: 0.0707 (0.0707)
loss_mask: 0.1120 (0.1120) loss_objectness: 0.0021 (0.0021)
loss_rpn_box_reg: 0.0043 (0.0043) time: 0.9965 data: 0.3444 max
mem: 3464
Epoch: [3] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.2418 (0.2369)
loss_classifier: 0.0279 (0.0308) loss_box_reg: 0.0656 (0.0665)
loss_mask: 0.1386 (0.1345) loss_objectness: 0.0007 (0.0010)
loss_rpn_box_reg: 0.0036 (0.0041) time: 0.6190 data: 0.0402 max
mem: 3464
Epoch: [3] [20/60] eta: 0:00:25 lr: 0.000500 loss: 0.2288 (0.2314)
loss_classifier: 0.0299 (0.0303) loss_box_reg: 0.0602 (0.0632)
loss_mask: 0.1326 (0.1333) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0032 (0.0037) time: 0.6124 data: 0.0094 max
mem: 3464
Epoch: [3] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.2084 (0.2180)
loss_classifier: 0.0284 (0.0304) loss_box_reg: 0.0457 (0.0558)
loss_mask: 0.1106 (0.1270) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0030 (0.0038) time: 0.6178 data: 0.0095 max
mem: 3464
Epoch: [3] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1870 (0.2101)
loss_classifier: 0.0254 (0.0296) loss_box_reg: 0.0372 (0.0535)

```



```
loss_mask: 0.1054 (0.1226) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0022 (0.0035) time: 0.5934 data: 0.0098 max
mem: 3464
Epoch: [3] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1886 (0.2074)
loss_classifier: 0.0208 (0.0285) loss_box_reg: 0.0435 (0.0524)
loss_mask: 0.1120 (0.1220) loss_objectness: 0.0004 (0.0010)
loss_rpn_box_reg: 0.0020 (0.0035) time: 0.6067 data: 0.0106 max
mem: 3464
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1886 (0.2092)
loss_classifier: 0.0252 (0.0288) loss_box_reg: 0.0435 (0.0524)
loss_mask: 0.1151 (0.1234) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0031 (0.0036) time: 0.6042 data: 0.0099 max
mem: 3464
Epoch: [3] Total time: 0:00:36 (0.6153 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1925 (0.1925)
evaluator_time: 0.0142 (0.0142) time: 0.5605 data: 0.3520 max mem:
3464
Test: [49/50] eta: 0:00:00 model_time: 0.1048 (0.1132)
evaluator_time: 0.0047 (0.0065) time: 0.1237 data: 0.0090 max mem:
3464
Test: Total time: 0:00:07 (0.1403 s / it)
Averaged stats: model_time: 0.1048 (0.1132) evaluator_time: 0.0047
(0.0065)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.828
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.368
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.650
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.846
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.419
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.866
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.866
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
```

```

maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.880
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.758
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.378
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.517
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.772
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.382
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.797
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
Epoch: [4] [ 0/60] eta: 0:01:03 lr: 0.000500 loss: 0.1764 (0.1764)
loss_classifier: 0.0218 (0.0218) loss_box_reg: 0.0406 (0.0406)
loss_mask: 0.1091 (0.1091) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0045 (0.0045) time: 1.0649 data: 0.3697 max
mem: 3464
Epoch: [4] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.2057 (0.2016)
loss_classifier: 0.0229 (0.0288) loss_box_reg: 0.0416 (0.0444)
loss_mask: 0.1205 (0.1236) loss_objectness: 0.0006 (0.0006)
loss_rpn_box_reg: 0.0044 (0.0043) time: 0.6200 data: 0.0414 max
mem: 3464
Epoch: [4] [20/60] eta: 0:00:25 lr: 0.000500 loss: 0.2057 (0.2106)
loss_classifier: 0.0295 (0.0301) loss_box_reg: 0.0496 (0.0495)
loss_mask: 0.1205 (0.1261) loss_objectness: 0.0007 (0.0008)
loss_rpn_box_reg: 0.0037 (0.0041) time: 0.6055 data: 0.0097 max
mem: 3700
Epoch: [4] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1907 (0.2052)
loss_classifier: 0.0274 (0.0292) loss_box_reg: 0.0480 (0.0478)
loss_mask: 0.1135 (0.1238) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0029 (0.0036) time: 0.6023 data: 0.0097 max

```

```
mem: 3700
Epoch: [4] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1797 (0.2047)
loss_classifier: 0.0246 (0.0302) loss_box_reg: 0.0437 (0.0480)
loss_mask: 0.1135 (0.1223) loss_objectness: 0.0005 (0.0007)
loss_rpn_box_reg: 0.0027 (0.0034) time: 0.5845 data: 0.0109 max
mem: 3700
Epoch: [4] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1739 (0.1996)
loss_classifier: 0.0253 (0.0290) loss_box_reg: 0.0385 (0.0468)
loss_mask: 0.1050 (0.1199) loss_objectness: 0.0005 (0.0008)
loss_rpn_box_reg: 0.0023 (0.0032) time: 0.5931 data: 0.0120 max
mem: 3700
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1693 (0.1997)
loss_classifier: 0.0275 (0.0292) loss_box_reg: 0.0385 (0.0467)
loss_mask: 0.1055 (0.1197) loss_objectness: 0.0005 (0.0009)
loss_rpn_box_reg: 0.0023 (0.0032) time: 0.6095 data: 0.0095 max
mem: 3700
Epoch: [4] Total time: 0:00:36 (0.6121 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:26 model_time: 0.1775 (0.1775)
evaluator_time: 0.0206 (0.0206) time: 0.5346 data: 0.3344 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.0987 (0.1066)
evaluator_time: 0.0026 (0.0048) time: 0.1097 data: 0.0037 max mem:
3700
Test: Total time: 0:00:06 (0.1277 s / it)
Averaged stats: model_time: 0.0987 (0.1066) evaluator_time: 0.0026
(0.0048)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.840
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.981
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.368
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.674
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.426
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.881
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.881
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.897
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.757
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.970
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.913
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.378
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.530
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.770
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.385
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.799
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.801
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.567
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.810
Epoch: [5] [ 0/60] eta: 0:00:59  lr: 0.000500  loss: 0.2363 (0.2363)
loss_classifier: 0.0346 (0.0346)  loss_box_reg: 0.0638 (0.0638)
loss_mask: 0.1342 (0.1342)  loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0035 (0.0035)  time: 0.9918  data: 0.4285  max
mem: 3700
Epoch: [5] [10/60] eta: 0:00:30  lr: 0.000500  loss: 0.1874 (0.1939)
loss_classifier: 0.0242 (0.0280)  loss_box_reg: 0.0449 (0.0424)
loss_mask: 0.1151 (0.1199)  loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0021 (0.0025)  time: 0.6078  data: 0.0489  max
mem: 3700
Epoch: [5] [20/60] eta: 0:00:24  lr: 0.000500  loss: 0.1931 (0.2105)
loss_classifier: 0.0356 (0.0326)  loss_box_reg: 0.0477 (0.0501)
loss_mask: 0.1145 (0.1234)  loss_objectness: 0.0004 (0.0014)
loss_rpn_box_reg: 0.0021 (0.0030)  time: 0.5918  data: 0.0103  max
mem: 3700
Epoch: [5] [30/60] eta: 0:00:17  lr: 0.000500  loss: 0.1931 (0.2008)
loss_classifier: 0.0284 (0.0289)  loss_box_reg: 0.0477 (0.0472)

```

```
loss_mask: 0.1142 (0.1206) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5921 data: 0.0098 max
mem: 3700
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1741 (0.1982)
loss_classifier: 0.0214 (0.0286) loss_box_reg: 0.0373 (0.0459)
loss_mask: 0.1122 (0.1197) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0026 (0.0030) time: 0.5814 data: 0.0102 max
mem: 3700
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1935 (0.1971)
loss_classifier: 0.0248 (0.0283) loss_box_reg: 0.0424 (0.0465)
loss_mask: 0.1114 (0.1183) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0034 (0.0031) time: 0.6012 data: 0.0099 max
mem: 3700
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1741 (0.1930)
loss_classifier: 0.0218 (0.0271) loss_box_reg: 0.0325 (0.0444)
loss_mask: 0.1061 (0.1175) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0026 (0.0031) time: 0.6125 data: 0.0085 max
mem: 3700
Epoch: [5] Total time: 0:00:36 (0.6051 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:28 model_time: 0.1591 (0.1591)
evaluator_time: 0.0116 (0.0116) time: 0.5638 data: 0.3914 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.1025 (0.1087)
evaluator_time: 0.0039 (0.0052) time: 0.1167 data: 0.0041 max mem:
3700
Test: Total time: 0:00:06 (0.1351 s / it)
Averaged stats: model_time: 0.1025 (0.1087) evaluator_time: 0.0039
(0.0052)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.837
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.930
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.858
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.423
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.875
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.875
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.500
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.812
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.892
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.762
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.969
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.932
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.544
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.385
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.804
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.806
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.813
Epoch: [6] [ 0/60] eta: 0:01:24 lr: 0.000050 loss: 0.2115 (0.2115)
loss_classifier: 0.0332 (0.0332) loss_box_reg: 0.0507 (0.0507)
loss_mask: 0.1236 (0.1236) loss_objectness: 0.0005 (0.0005)
loss_rpn_box_reg: 0.0036 (0.0036) time: 1.4092 data: 0.6883 max
mem: 3700
Epoch: [6] [10/60] eta: 0:00:33 lr: 0.000050 loss: 0.2112 (0.1921)
loss_classifier: 0.0329 (0.0284) loss_box_reg: 0.0381 (0.0406)
loss_mask: 0.1210 (0.1200) loss_objectness: 0.0002 (0.0004)
loss_rpn_box_reg: 0.0027 (0.0027) time: 0.6647 data: 0.0694 max
mem: 3700
Epoch: [6] [20/60] eta: 0:00:25 lr: 0.000050 loss: 0.1861 (0.1938)
loss_classifier: 0.0253 (0.0281) loss_box_reg: 0.0354 (0.0426)
loss_mask: 0.1169 (0.1195) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0027 (0.0032) time: 0.6113 data: 0.0086 max

```

```
mem: 3700
Epoch: [6] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1861 (0.1884)
loss_classifier: 0.0226 (0.0266) loss_box_reg: 0.0354 (0.0416)
loss_mask: 0.1085 (0.1164) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0024 (0.0031) time: 0.6113 data: 0.0095 max
mem: 3700
Epoch: [6] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1644 (0.1829)
loss_classifier: 0.0184 (0.0256) loss_box_reg: 0.0317 (0.0396)
loss_mask: 0.1084 (0.1143) loss_objectness: 0.0002 (0.0005)
loss_rpn_box_reg: 0.0022 (0.0029) time: 0.5842 data: 0.0091 max
mem: 3700
Epoch: [6] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.1649 (0.1882)
loss_classifier: 0.0205 (0.0266) loss_box_reg: 0.0322 (0.0410)
loss_mask: 0.1118 (0.1170) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0018 (0.0029) time: 0.5804 data: 0.0091 max
mem: 3700
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1811 (0.1904)
loss_classifier: 0.0275 (0.0268) loss_box_reg: 0.0370 (0.0415)
loss_mask: 0.1207 (0.1183) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0027 (0.0030) time: 0.5921 data: 0.0083 max
mem: 3700
Epoch: [6] Total time: 0:00:36 (0.6128 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:40 model_time: 0.1714 (0.1714)
evaluator_time: 0.0194 (0.0194) time: 0.8045 data: 0.6114 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.1001 (0.1105)
evaluator_time: 0.0025 (0.0057) time: 0.1101 data: 0.0037 max mem:
3700
Test: Total time: 0:00:07 (0.1411 s / it)
Averaged stats: model_time: 0.1001 (0.1105) evaluator_time: 0.0025
(0.0057)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.836
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.941
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.655
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
```

```

maxDets=100 ] = 0.858
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.420
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.874
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.874
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.812
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.892
IoU metric: segm
Average Precision    (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.763
Average Precision    (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.970
Average Precision    (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.917
Average Precision    (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision    (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.540
Average Precision    (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.386
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.803
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.805
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [7] [ 0/60] eta: 0:00:58 lr: 0.000050 loss: 0.1546 (0.1546)
loss_classifier: 0.0144 (0.0144) loss_box_reg: 0.0233 (0.0233)
loss_mask: 0.1152 (0.1152) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0015 (0.0015) time: 0.9699 data: 0.3689 max
mem: 3700
Epoch: [7] [10/60] eta: 0:00:31 lr: 0.000050 loss: 0.1793 (0.1782)
loss_classifier: 0.0260 (0.0260) loss_box_reg: 0.0386 (0.0390)
loss_mask: 0.1070 (0.1103) loss_objectness: 0.0002 (0.0003)
loss_rpn_box_reg: 0.0030 (0.0026) time: 0.6248 data: 0.0409 max
mem: 3700
Epoch: [7] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1793 (0.1823)

```



```
loss_classifier: 0.0254 (0.0261) loss_box_reg: 0.0397 (0.0411)
loss_mask: 0.1049 (0.1116) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0030 (0.0029) time: 0.6018 data: 0.0095 max
mem: 3700
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1716 (0.1806)
loss_classifier: 0.0230 (0.0257) loss_box_reg: 0.0397 (0.0396)
loss_mask: 0.1068 (0.1118) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0027 (0.0029) time: 0.5863 data: 0.0099 max
mem: 3700
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1733 (0.1817)
loss_classifier: 0.0228 (0.0256) loss_box_reg: 0.0364 (0.0394)
loss_mask: 0.1077 (0.1133) loss_objectness: 0.0003 (0.0005)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5703 data: 0.0087 max
mem: 3700
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1733 (0.1853)
loss_classifier: 0.0205 (0.0257) loss_box_reg: 0.0364 (0.0405)
loss_mask: 0.1171 (0.1155) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0031 (0.0030) time: 0.5830 data: 0.0088 max
mem: 3700
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1799 (0.1869)
loss_classifier: 0.0235 (0.0261) loss_box_reg: 0.0418 (0.0411)
loss_mask: 0.1171 (0.1161) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0028 (0.0030) time: 0.5975 data: 0.0084 max
mem: 3700
Epoch: [7] Total time: 0:00:36 (0.6016 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:27 model_time: 0.1995 (0.1995)
evaluator_time: 0.0102 (0.0102) time: 0.5437 data: 0.3323 max mem:
3700
Test: [49/50] eta: 0:00:00 model_time: 0.0981 (0.1083)
evaluator_time: 0.0026 (0.0047) time: 0.1100 data: 0.0036 max mem:
3700
Test: Total time: 0:00:06 (0.1294 s / it)
Averaged stats: model_time: 0.0981 (0.1083) evaluator_time: 0.0026
(0.0047)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.839
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.941
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.424
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.879
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.879
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.896
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.763
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.534
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.386
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [8] [ 0/60] eta: 0:01:05 lr: 0.000050 loss: 0.1549 (0.1549)
loss_classifier: 0.0242 (0.0242) loss_box_reg: 0.0308 (0.0308)
loss_mask: 0.0977 (0.0977) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0020 (0.0020) time: 1.0998 data: 0.4827 max
mem: 3700
Epoch: [8] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1596 (0.1569)
loss_classifier: 0.0237 (0.0219) loss_box_reg: 0.0263 (0.0260)
loss_mask: 0.1095 (0.1062) loss_objectness: 0.0003 (0.0005)

```

```
loss_rpn_box_reg: 0.0020 (0.0022)  time: 0.5801  data: 0.0547  max
mem: 3700
Epoch: [8]  [20/60]  eta: 0:00:23  lr: 0.000050  loss: 0.1797 (0.1906)
loss_classifier: 0.0262 (0.0277)  loss_box_reg: 0.0315 (0.0408)
loss_mask: 0.1151 (0.1184)  loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5693  data: 0.0113  max
mem: 3700
Epoch: [8]  [30/60]  eta: 0:00:18  lr: 0.000050  loss: 0.1824 (0.1826)
loss_classifier: 0.0262 (0.0254)  loss_box_reg: 0.0394 (0.0391)
loss_mask: 0.1151 (0.1145)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0028 (0.0029)  time: 0.6157  data: 0.0108  max
mem: 3700
Epoch: [8]  [40/60]  eta: 0:00:12  lr: 0.000050  loss: 0.1793 (0.1847)
loss_classifier: 0.0236 (0.0260)  loss_box_reg: 0.0394 (0.0401)
loss_mask: 0.1075 (0.1150)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.6154  data: 0.0103  max
mem: 3700
Epoch: [8]  [50/60]  eta: 0:00:05  lr: 0.000050  loss: 0.1794 (0.1860)
loss_classifier: 0.0266 (0.0254)  loss_box_reg: 0.0393 (0.0395)
loss_mask: 0.1112 (0.1176)  loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5890  data: 0.0097  max
mem: 3700
Epoch: [8]  [59/60]  eta: 0:00:00  lr: 0.000050  loss: 0.1794 (0.1867)
loss_classifier: 0.0276 (0.0260)  loss_box_reg: 0.0365 (0.0408)
loss_mask: 0.1053 (0.1166)  loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0028)  time: 0.5954  data: 0.0087  max
mem: 3700
Epoch: [8] Total time: 0:00:36 (0.6055 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:28  model_time: 0.1800 (0.1800)
evaluator_time: 0.0126 (0.0126)  time: 0.5767  data: 0.3823  max mem:
3700
Test:  [49/50]  eta: 0:00:00  model_time: 0.1195 (0.1192)
evaluator_time: 0.0047 (0.0068)  time: 0.1389  data: 0.0058  max mem:
3700
Test: Total time: 0:00:07 (0.1499 s / it)
Averaged stats: model_time: 0.1195 (0.1192)  evaluator_time: 0.0047
(0.0068)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.840
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.979
Average Precision  (AP) @[ IoU=0.75      | area=  all |
```

```

maxDets=100 ] = 0.931
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.861
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.425
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.878
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.878
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.895
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.765
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.535
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.777
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.385
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.811
Epoch: [9] [ 0/60] eta: 0:01:36 lr: 0.000005 loss: 0.1179 (0.1179)
loss_classifier: 0.0076 (0.0076) loss_box_reg: 0.0123 (0.0123)
loss_mask: 0.0972 (0.0972) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0006 (0.0006) time: 1.6021 data: 0.8491 max
mem: 3700

```

```
Epoch: [9] [10/60] eta: 0:00:34 lr: 0.000005 loss: 0.1728 (0.1847)
loss_classifier: 0.0242 (0.0255) loss_box_reg: 0.0357 (0.0386)
loss_mask: 0.1076 (0.1176) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0020 (0.0023) time: 0.6896 data: 0.0844 max
mem: 3700
Epoch: [9] [20/60] eta: 0:00:26 lr: 0.000005 loss: 0.1840 (0.1879)
loss_classifier: 0.0292 (0.0284) loss_box_reg: 0.0365 (0.0414)
loss_mask: 0.1125 (0.1150) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0026 (0.0025) time: 0.6100 data: 0.0099 max
mem: 3700
Epoch: [9] [30/60] eta: 0:00:19 lr: 0.000005 loss: 0.1834 (0.1829)
loss_classifier: 0.0286 (0.0269) loss_box_reg: 0.0356 (0.0394)
loss_mask: 0.1061 (0.1132) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0031 (0.0026) time: 0.6119 data: 0.0100 max
mem: 3778
Epoch: [9] [40/60] eta: 0:00:12 lr: 0.000005 loss: 0.1670 (0.1854)
loss_classifier: 0.0234 (0.0268) loss_box_reg: 0.0321 (0.0401)
loss_mask: 0.1090 (0.1150) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0032 (0.0028) time: 0.5948 data: 0.0085 max
mem: 3778
Epoch: [9] [50/60] eta: 0:00:06 lr: 0.000005 loss: 0.1898 (0.1883)
loss_classifier: 0.0244 (0.0267) loss_box_reg: 0.0388 (0.0407)
loss_mask: 0.1180 (0.1174) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0032 (0.0029) time: 0.5872 data: 0.0101 max
mem: 3778
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1793 (0.1891)
loss_classifier: 0.0253 (0.0269) loss_box_reg: 0.0378 (0.0417)
loss_mask: 0.1127 (0.1169) loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0029 (0.0030) time: 0.5977 data: 0.0096 max
mem: 3778
Epoch: [9] Total time: 0:00:37 (0.6213 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:42 model_time: 0.2411 (0.2411)
evaluator_time: 0.0180 (0.0180) time: 0.8438 data: 0.5829 max mem:
3778
Test: [49/50] eta: 0:00:00 model_time: 0.0998 (0.1116)
evaluator_time: 0.0029 (0.0051) time: 0.1125 data: 0.0038 max mem:
3778
Test: Total time: 0:00:06 (0.1391 s / it)
Averaged stats: model_time: 0.0998 (0.1116) evaluator_time: 0.0029
(0.0051)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.839
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.931
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.424
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.877
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.877
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.467
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.825
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.894
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.767
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.535
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.779
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.386
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.804
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.806
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.633
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.787
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.812
That's it!

```

So after one epoch of training, we obtain a COCO-style mAP > 50, and a mask mAP of 65.

But what do the predictions look like? Let's take one image in the dataset and verify

```
import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

image =
read_image("drive/MyDrive/_static/img/tv_tutorial/tv_image05.png")
eval_transform = get_transform(train=False)


model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

<matplotlib.image.AxesImage at 0x7ef574e209a0>
```



pred_labels

```
['pedestrian: 0.995',  
'pedestrian: 0.994',  
'pedestrian: 0.992',  
'pedestrian: 0.988',  
'pedestrian: 0.979',  
'pedestrian: 0.797',  
'pedestrian: 0.453',  
'pedestrian: 0.215',  
'pedestrian: 0.089',  
'pedestrian: 0.066',  
'pedestrian: 0.055']
```

pred_boxes

```
tensor([[299, 116, 371, 335],  
        [173, 122, 271, 408],  
        [446, 98, 507, 315],  
        [517, 118, 564, 256],  
        [356, 125, 439, 326],
```



```
[239, 125, 309, 351],
[235, 123, 464, 330],
[312, 112, 507, 313],
[489, 107, 517, 168],
[377, 119, 474, 317],
[543, 106, 564, 257]], device='cuda:0')
```

The results look good!

Wrapping up

In this tutorial, you have learned how to create your own training pipeline for object detection models on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes and segmentation masks. You also leveraged a Mask R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-GPU training, check `references/detection/train.py`, which is present in the torchvision repository.

You can download a full source file for this tutorial [here](#).

```
image = read_image("drive/MyDrive/Beatles_-_Abbey_Road.jpg")
eval_transform = get_transform(train=False)


model.eval()
with torch.no_grad():
    x = eval_transform(image)
    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)
    predictions = model([x, ])
    pred = predictions[0]

image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]
pred_labels = [f"pedestrian: {score:.3f}" for label, score in
zip(pred["labels"], pred["scores"])]
pred_boxes = pred["boxes"].long()
output_image = draw_bounding_boxes(image, pred_boxes, pred_labels,
colors="red")

masks = (pred["masks"] > 0.7).squeeze(1)
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))
```

<matplotlib.image.AxesImage at 0x7ef4a2a48cd0>



pred_labels

```
['pedestrian: 0.987',  
'pedestrian: 0.985',  
'pedestrian: 0.957',  
'pedestrian: 0.939',  
'pedestrian: 0.103',  
'pedestrian: 0.093',  
'pedestrian: 0.089',
```

```
'pedestrian: 0.060',  
'pedestrian: 0.059']
```

pred_boxes

```
tensor([[ 21, 161,  84, 273],  
        [ 92, 164, 144, 272],  
        [148, 168, 219, 296],  
        [223, 163, 290, 293],  
        [ 41, 162, 237, 277],  
        [239, 147, 245, 174],  
        [114, 160, 277, 288],  
        [ 57, 159, 136, 269],  
        [104, 162, 190, 273]], device='cuda:0')
```